Observability and Search in the AI Era; Initiate DT (Buy), ESTC (Neutral)

21 April 2026 | 12:01AM EDT

We initiate coverage of Dynatrace (DT, Buy, $45 PT) and Elastic (ESTC, Neutral, $50 PT). Observability is entering a particularly dynamic period, as AI deepens the strategic importance of machine data, and raises the bar for the platforms built to harness it. Compute intensity is scaling rapidly as AI workloads proliferate, workload complexity is expanding the coverage surface area, and the shift from reactive monitoring to autonomous remediation is pulling budget away from fragmented point solutions and toward platforms built to compress signal-to-action. Dynatrace sits squarely in that current. A fresh product cycle in logs, a DPS contract transition that resets the growth algorithm into next year, and strong margin support (FCFM 25%+) provide a compelling setup against a reset valuation (~14x EV/FCF). On Elastic, the platform has structural relevance across search, observability and security, and execution has improved over the past two years. We still need to see clearer evidence that cloud growth can durably inflect – validating both the AI narrative and reinforcing Elastic's positioning in an increasingly competitive backdrop – before getting more constructive. **Why read this report:** Alongside company deep-dives, we map how AI and competition are reshaping observability - and validate the direction with signals from the AI private ecosystem.

订阅研报添加微信 LCD123321Cad 04

Table of Contents

Executive Summary: Initiating Coverage of Dynatrace and Elastic; Buy DT, Neutral ESTC	3
Industry Primer: The Next Phase of Observability	5
Market Overview: Infrastructure, Data and Telemetry Tailwinds	13
AI Private Ecosystem Signals	15
Dynatrace: Initiating Buy, PT $45	19
Elastic: Initiating Neutral, PT $50	30
Disclosure Appendix	44

订阅研报添加微信 LCD123321Cad 04-21

订阅研报添加微信 LCD123321Cad 04-21

订阅研报添加微信 LCD123321Cad 04-21

We initiate coverage of Dynatrace (DT) with a Buy rating and $45 Price Target and Elastic (ESTC) with a Neutral rating and $50 Price Target. Dynatrace operates at the center of the enterprise observability and AIOps market, where value is shifting from telemetry collection and analysis toward automated diagnosis, remediation and AI-driven operations – a transition we believe Dynatrace is well positioned to capitalize on, with a well-established platform and growth trajectory that we think the market is underappreciating. Elastic spans a broader set of markets – search, observability and security – with a common data and retrieval architecture that is becoming more relevant as AI expands the importance of context, relevance and unstructured data handling, though the company still needs to show that strategic breadth can translate into a more durable acceleration in growth. **We highlight the key dynamics driving our views below:**

订阅研报添加微信 LCD12321Cad 04-21

- **Dynatrace has a strong platform position in enterprise observability and a more visible path to growth acceleration than the market currently reflects.** Dynatrace helps large enterprises monitor, troubleshoot, and increasingly automate their most complex digital environments, combining full-stack observability with Davis AI – purpose-built for causation and automated remediation – and extending further into agentic operations and security through Dynatrace Intelligence. The platform is well positioned against the most important shifts in the category – the move toward AI-driven automation, observability and security convergence, and rising log analytics demand – while the financial profile is already strong with near-30% non-GAAP OpM's and 26% FCFM's. Looking ahead, an upcoming DPS renewal cycle and a scaling logs business support a growth outlook that we believe the market is not fully pricing in. Taken together, we think Dynatrace offers a compelling combination of platform durability, improving growth visibility and strong cash generation that provides solid downside support.
- **Elastic provides a credible platform across three strategic markets but converting strategic relevance into durable growth acceleration remains the central debate.** Built around Elasticsearch, the company serves as the data foundation for enterprise search, log analytics, SIEM, and increasingly AI-driven retrieval and agent workflows. Newer capabilities including Agent Builder, Workflows and Attack Discovery are compelling and signal a move up-stack to capture higher-value workflows. GTM execution has also improved over the past two years, showing up in stable NER (112%) and net-new revenue. But Elastic Cloud continues to decelerate even as the AI narrative has broadened and consumption peers accelerate, while recent self-managed outperformance (while encouraging) has not been rewarded by the market. Competition, meanwhile, is intensifying across all three markets as broader platforms aim to own more of the workflow and data layer. We believe a more constructive stance requires greater confidence that cloud growth can stabilize and durably inflect, validating both the AI narrative and Elastic's enduring relevance across its end markets.

Why read this report:

- **An industry primer on how the observability market is being redefined.** Introduces a six-theme framework that explains how changes in architecture, telemetry standards, AI workloads, data economics, and buyer consolidation are reshaping where differentiation, pricing power, and long-term value accrue.
- **An AI company study showing where competitive pressure is forming first.** Draws on signals from the startup ecosystem to highlight how innovation is clustering around telemetry economics, pipeline control, and AI-assisted operations, and what that signals for incumbent platforms. We also apply the GS AI Impact Framework to assess how Dynatrace and Elastic are positioned against AI disruption debates.
- **A scenario-based framework for Dynatrace's path to growth acceleration.** Lays out how DPS renewal dynamics and accelerating logs adoption can translate consumption momentum into improving reported growth over time.
- **Elastic evaluated through the lens of competition and execution.** Assesses Search & AI, Observability, and Security to highlight where Elastic's platform remains well positioned, how competitive pressure from broader platforms is intensifying, and how execution in these markets shapes the path toward more durable cloud-led momentum.

订阅研报添加微信 LCD123321Cad 04-21

订阅研报添加微信 LCD123321Cad 04-21

As we launch coverage on Dynatrace and Elastic, we think both companies are positioned against market backdrops that are becoming more consequential, albeit in different ways. In observability, the center of gravity is shifting from collecting and visualizing telemetry toward automating diagnosis, remediation and operational response; in search, the rise of AI is increasing the strategic value of retrieval, relevance and access to large pools of unstructured data. These are distinct categories, but both are becoming more important as enterprises are forced to convert rapidly rising volumes of machine data into faster decisions, more effective automation, and durable operating leverage, rather than simply absorbing higher monitoring and data infrastructure costs.

Dynatrace and Elastic overlap most directly in observability; we focus on six key themes that we believe best capture where the observability market is heading and how the basis of competition is changing: 1) Observability is shifting from visibility toward operations, with value moving from dashboards to prevention, remediation, and automation; **2)** Security and observability are converging around shared telemetry, increasing the strategic importance of platforms central to logs, metrics, and runtime context even as buying motions remain distinct; **3)** OpenTelemetry is pushing differentiation up the stack, reducing the defensibility of proprietary instrumentation and shifting competition toward analytics, workflows, and economics; **4)** AI is increasing operational complexity and raising the bar for observability, as AI-native workloads introduce new failure modes, cost dynamics, and reliability challenges; **5)** Data economics are becoming a core area of competition, with customers seeking broader visibility without unsustainably rising data costs; and **6)** Platform consolidation is accelerating, as enterprises rationalize tool sprawl in favor of fewer platforms built on shared data with simpler operations and clearer total cost of ownership.

订阅研报添加微信 LCD123321Cad 04-21

Theme #1: From Visibility to Operational Control

We believe observability is moving up the stack from visibility toward execution. The category was originally built to make increasingly complex software environments easier to see and diagnose. That remains foundational, but we believe the center of gravity is shifting. As architectures have become more distributed and release cycles more continuous, vendors are increasingly orienting product roadmaps around faster diagnosis, remediation, and operational response rather than pure telemetry visualization. In our view, the value proposition is moving closer to shortening the path from signal to action – and, over time, automating more of that loop altogether.

订阅研报添加微信 LCD123321Cad 04-21

The shift is increasingly visible in how the market is being built. Dynatrace's launch of Dynatrace Intelligence (Jan. 2026) framed the platform around agentic operations and trusted automation, while Datadog's Workflow Automation and Bits AI SRE point in a similar direction by unifying observability data to support investigation and response. Elastic's Observability AI assistant is a more incremental example, but the direction of travel is clear. We believe observability is gradually evolving from a system of insight into more of an operating layer for modern software environments, and the vendors best positioned will be those that can pair broad telemetry coverage with strong context, automation, and workflow depth – not simply collect more data but help customers operationalize it.

As observability transitions from a system of record to a system of action, the surface area for value capture expands materially. By integrating earlier into the software development lifecycle - shaping testing and application design - and becoming more central to live operations, platforms are evolving from visibility tools into a mission-critical operating layer. This broader footprint provides a more robust framework for underwriting ROI. In our discussions with private companies, value is increasingly framed through two primary levers: the quantifiable cost of downtime and the compression of SRE labor. In this model, human effort becomes a measurable cost pool that software can actively reduce. We believe the category is heading toward a participation model: as automation becomes more tangible, the best-positioned platforms will be those that reliably translate telemetry into trusted action, allowing them to capture a portion of the significant productivity gains they create.

订阅研报添加微信 LCD123321Cad 04-21

Exhibit 1: Observability is shifting from reactive to autonomous operations

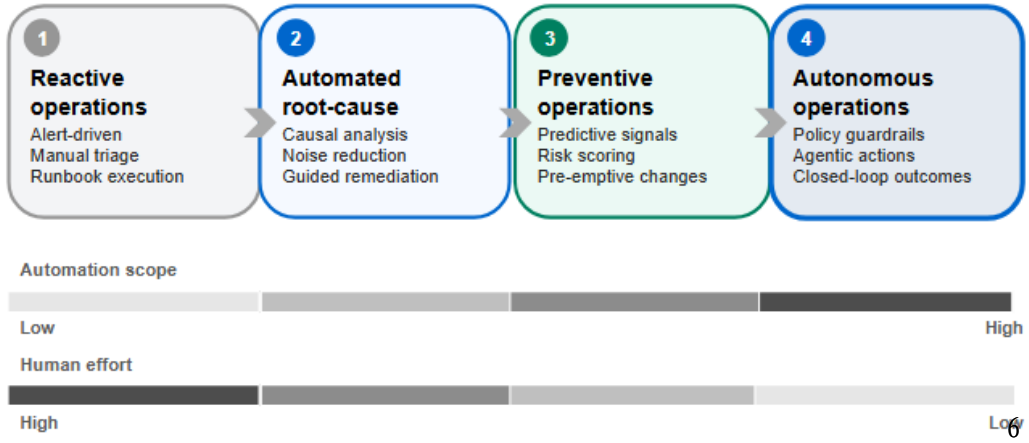

Theme #2: Security and Observability Convergence

Security and observability are increasingly drawing from the same underlying telemetry. Logs, traces, runtime events and infrastructure signals now sit near the center of both operational monitoring and threat detection, which is pulling the two markets closer together. That is visible from the security side in moves such as Palo Alto's acquisition of Chronosphere and CrowdStrike's expansion around high-scale log analytics through Humio (acquired 2021) and Onum (acquired 2025). Both play to the same M&A philosophy – own the telemetry data plumbing (ingest, filter, correlate) inside a broader platform, in line with buyer demand for fewer vendors and faster paths from detection to automated response. The key point, in our view, is that the underlying data plane is becoming more strategic as more teams rely on it for monitoring, troubleshooting, and security analysis.

订阅研报添加微信 LCD123321Cad 04-21

Exhibit 2: Observability and security workflows are converging on shared telemetry

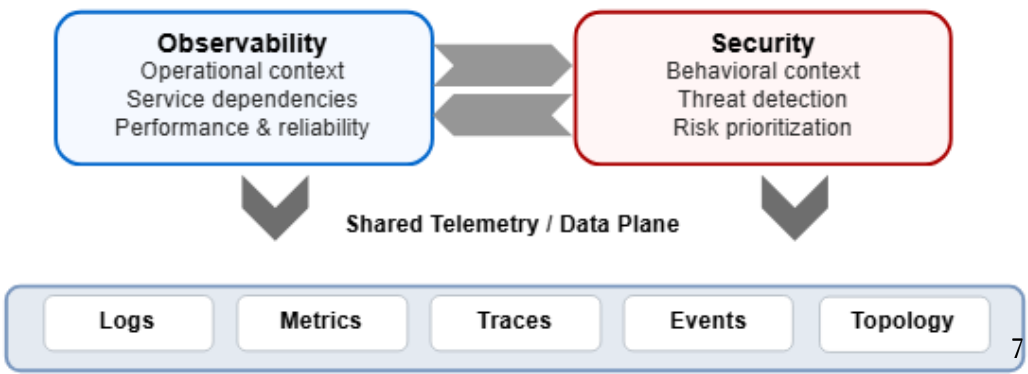

workflow depth, differentiated context or stronger economics.

Exhibit 3: Timeline of Recent Industry Consolidation and Product Launches in Observability
Security vendors, data-platforms and Hyperscalers all moving more aggressively into security and observability workflows, raising the competitive bar

Event	Timeline	Strategic Implications
Databricks launches Lakewatch (agentic SIEM) and acquires Antimatter & SiftD.ai	Mar. 24, 2026	Pushes SIEM into the lakehouse; frames security telemetry, detection and investigation as open data workflows rather than a standalone security silo. Anthropic-powered agentic workflows plus the Antimatter and SiftD.ai acquisitions signal a broader push to make data platforms credible control points for security operations.
Snowflake announces acquisition of Observe Inc. (~$1B)	Feb. 2026	Pull observability data into the data cloud; logs, metrics, and events increasingly treated as strategic enterprise data rather than a separate monitoring layer. Signals that telemetry is becoming a first-class workload inside broader data and AI architectures.
AWS re:Invent 2025: Expands into Gen-AI observability & unified security-ops data lake	Dec. 2025	Deepens native-cloud observability and SIEM-adjacent tooling through CloudWatch, OpenSearch, and AI observability enhancements; raises the floor for built-in functionality. Signals that Hyperscalers are becoming more aggressive competitors for monitoring, log analytics, and telemetry-driven workflows, especially inside their own installed base.
Palo Alto Networks acquires Chronosphere ($3.35B)	Jan. 2025	Pulls observability into the security control plane; pairs Chronosphere's telemetry, cloud-native metrics and AI-era observability with Palo Alto's much larger security distribution. Signals production telemetry is becoming core to detection, response and automation - and that security vendors want to own more of that data upstream.
CrowdStrike acquires Onum (~$290M)	Aug. 2025	Moves detection closer to the data pipeline to reduce SIEM noise, improve performance and lower cost before data reaches expensive downstream tools. Highlights the growing strategic value of first-mile telemetry optimization as event volumes and AI-driven workflows scale.
ClickHouse launches ClickStack (Open-source observability stack)	May 2025	Turns Clickhouse from an observability backend into a more complete observability product; packages logs, metrics, traces and session replay on top of its high-performance analytical engine. Signals that observability is increasingly being treated as a data and query problem, with lower-cost, high-scale architectures pushing further up the stack.
Cisco completes acquisition of Splunk ($28B)	Mar. 2024 (Announced Sep. 2023)	Bring observability, machine-data analytics, and security telemetry into Cisco's much larger network and security footprint; turns Splunk from a stand-alone analytics platform into a broader control-plane asset. Signals that scaled telemetry platforms are increasingly more valuable when tied to an installed base, distribution engine, and adjacent workflows.

Source: Company Reports, Goldman Sachs Global Investment Research

订阅研报添加微信 LCD123321Cad 04-21

Theme #3: OpenTelemetry Raises the Competitive Bar

OpenTelemetry is reshaping observability by decoupling telemetry collection from proprietary agents that historically defined the category. For much of observability's history, platforms differentiated through tightly integrated agents that controlled how telemetry was collected, enriched, and correlated, which created strong lock-in but limited flexibility. OpenTelemetry changes that dynamic by standardizing how logs, metrics, and traces are generated and routed, giving customers more freedom in how telemetry is instrumented and where it is sent, and reducing proprietary collection as a standalone point of control. That said, open collection does not eliminate the value of approaches that simplify deployment, enrich telemetry with system-level context, and deliver a more consistent and reliable view of production behavior, which helps explain why Dynatrace's OneAgent remains relevant. Both Dynatrace and Elastic have been aligned with this shift for some time: Dynatrace supports OpenTelemetry ingestion alongside OneAgent-based enrichment, while Elastic has long supported native OTLP workflows and introduced its Distributions of OpenTelemetry (EDOT) in 2024. In our view, the key takeaway is that OpenTelemetry is no longer an alternative architecture, but the market baseline, with differentiation increasingly accruing above the collection layer rather than at it.

Exhibit 4: Which of these incubating CNCF projects is your organization using in production or evaluating?
OTel adoption trends underscore growing standardization on open collection

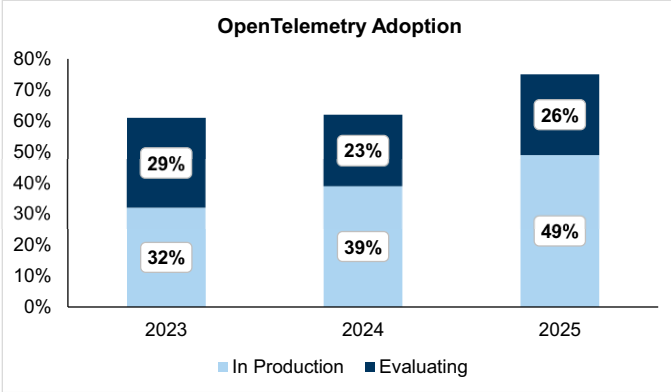

Data obtained by CNCF for their Annual Cloud Native Survey. As a part of the survey, 395-502 respondents answered this question in 2025, 689 respondents answered in 2024.

Source: Cloud Native Computing Foundation

The more critical implication is that competitive advantage is migrating to the highest levels of the stack. This shift is not a new phenomenon, but it has become the central industry tension: as telemetry becomes portable, the focal point of competition moves past the collection layer and even the processing plane toward the decision and action tier. While proprietary agents were never the primary source of value, the transition toward Open Standards as the market baseline has increased the onus on vendors to distinguish themselves through the intelligence that drives automated remediation and business outcomes. In our view, the processing plane, where data is shaped, routed, and governed, is becoming a strategic prerequisite, but the true competitive moat lies in the platform's ability to translate that data into decisive actions. Dynatrace's April 2026 acquisition of Bindplane is a clear signal of this transition, reflecting the need to control telemetry upstream so that the downstream decision engine can operate on high-fidelity, cost-optimized context. The vendors best positioned for this era are those that treat open collection as a commodity baseline and focus their innovation on the analytical intelligence and automated workflows that transform raw signals into proactive operational decisions.

Prometheus was an early signal that observability was already moving towards more open and modular architectures. It remains the foundational open-source standard for cloud-native metrics, and vendors such as Grafana, Chronosphere and Datadog all support Prometheus-compatible workflows in some form. OpenTelemetry pushes that same direction further by extending portability beyond metrics and across multiple signals. In our view, both technologies are best understood as part of the same broader architectural trend: more open collection, more modular telemetry stacks and a higher burden on platforms to differentiate above the data layer.

Exhibit 5: As Open Standards become the market basline, competitive advantage and value capture moves closer to the decision and action layer

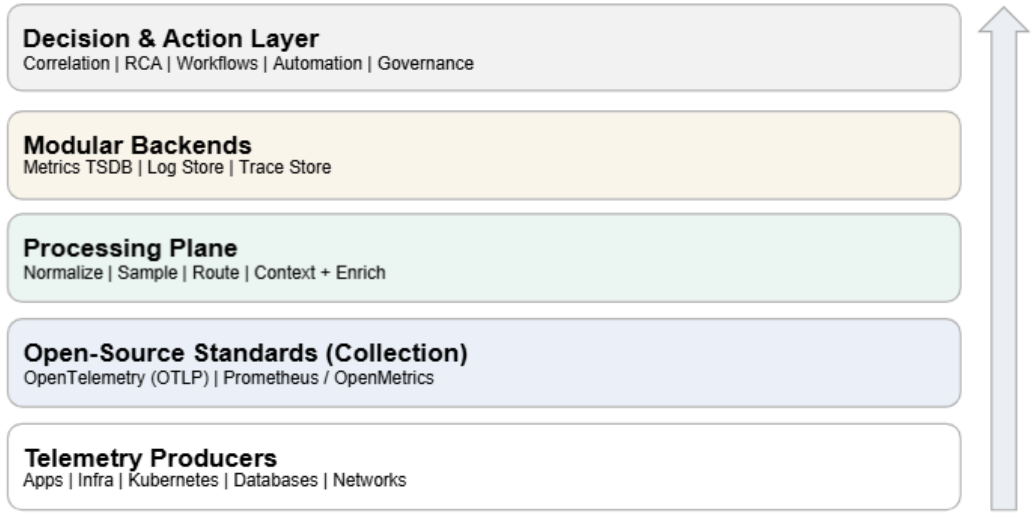

Source: Goldman Sachs Global Investment Research

Theme #4: AI Expands Both Markets

We believe AI is making modern software environments harder to monitor and more important to understand. On the observability side, AI workloads introduce new sources of complexity around latency, token consumption, reliability, model behavior, agent reliability, and cost control. That expands both the amount of telemetry enterprises need and the importance of interpreting that telemetry correctly across models, orchestration layers, tools and downstream infrastructure. Dynatrace has leaned directly into that trend with its AI and LLM observability offering, which is designed to provide real-time visibility into LLMs, AI agents, orchestration layers and their downstream application and infrastructure impact, and more recently broadened support for additional agentic AI frameworks and protocols. In our view, AI is not simply adding incremental workload volume; it is introducing a new class of production systems with non-deterministic behavior, opaque failure modes, and cost profiles that traditional application monitoring was not designed to manage.

订阅研报添加微信 LCD12321Cad 04-21

At the same time, AI is raising the strategic value of retrieval infrastructure. RAG and agentic workflows depend on fast, relevant access to enterprise data, which increases the importance of hybrid search, vector search, filtering and reranking quality. We believe this should be a tailwind for Elastic at the back-end retrieval layer, where its search heritage and hybrid capabilities are directly relevant. At the same time, this remains a highly competitive part of the stack, with alternatives spanning open-source (Postgres), MongoDB, Hyperscalers and specialist providers. Elastic is trying to differentiate with Agent Builder and tighter Jina AI-powered embeddings and reranking inside Elastic Inference Service – but we think the clearer benefit today is at the retrieval layer, while the eventual winners higher up the agent stack remain less settled.

订阅研报添加微信 LCD12321Cad 04-21

Theme #5: Data Economics Shape the Next Leg of Growth

We believe data economics have become one of the category's defining fault lines. The historical friction started in the Splunk era, where ingest-and-index economics often forced customers to ration visibility because retaining and querying data volumes

became expensive quickly. The shift to cloud-native platforms improved usability and broadened workflow value, but it did not eliminate the underlying tension once log volumes scaled. In our view, data economics are no longer a secondary consideration but a gating factor in platform selection, as customers increasingly evaluate whether they can afford to retain, query, and operationalize more data over time.

Observability is entering a phase where architecture matters at least as much as product breadth. Some companies are attacking the problem at the storage and query layer through lower-cost log architectures, lighter indexing, or more modular back ends. Other are attacking it even earlier, at the point of collection, by helping customers filter, enrich, and route telemetry before it ever lands in a premium destination. In our view, that first-mile shift is important because it changes where observability economics are determined. Once noisy or low-value data reaches an expensive back end, much of the cost has already compounded. Companies like Cribl are pressing that point directly, and the broader implication is that destination platforms are facing pressure not just from cheaper storage, but from upstream control over where data gets sent, retained and queried in the first place.

Against that backdrop, we think the market is likely to support both modular architectures and broader premium platforms. The modular path should appeal to customers optimizing for lower-cost back ends, more composable tooling, and tighter control over telemetry spend. The broader platform path should appeal to customers willing to pay more for unified data, tighter workflow integration, reliability, and a more complete operating layer. In our view, Dynatrace fits firmly in the latter camp, while Elastic straddles a middle-ground between a high-performance data backend and an increasingly capable operating layer. We do not believe being the lowest-cost endpoint is a prerequisite for success. But it does require both companies to deliver enough value through unification, context and workflow depth to justify a premium as data volumes continue to grow.

Exhibit 6: Observability Spend as % of Compute Infrastructure; Average: 17%, Median: 10%
As IT environments grow more complex, organizations are increasingly focused on optimizing observability spending

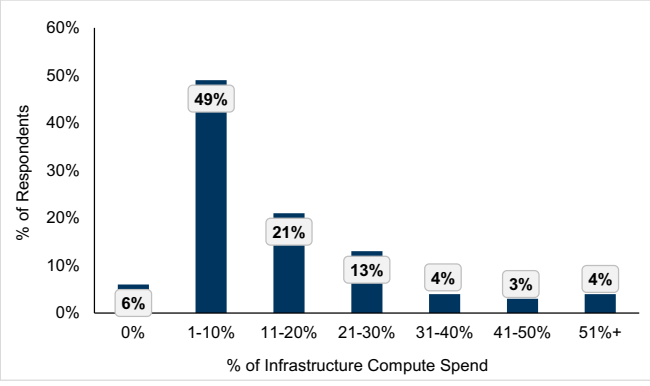

Source: Grafana 2025 Observability Survey

Exhibit 7: Cost has become the primary consideration when selecting new observability tools
Favors unified platforms and modular stacks providing attractive unit economics with feature depth and breadth

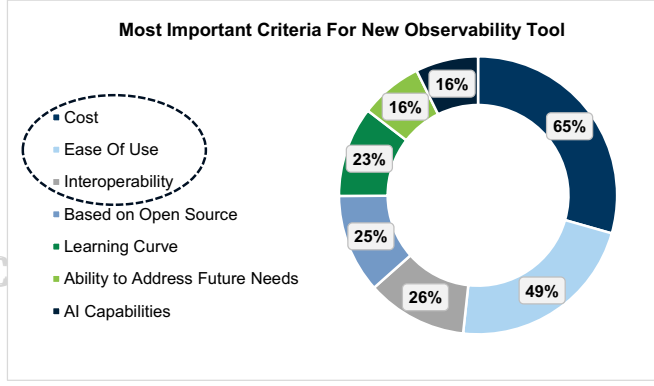

Source: Grafana Observability Survey 2026

Theme #6: Buyers Want Fewer Platforms

We believe observability is entering a more explicit consolidation cycle. Tool sprawl was tolerable when the category was younger and buyer were optimizing for

best-of-breed functionality at the point-solution level. That is less true today. As observability estates have become larger and more expensive, the cost of fragmentation becomes harder to ignore – not just in vendor spend, but in workflow friction, duplicated data and operational complexity. We believe this is central to why the market is beginning to place greater value on platforms that can unify more of the workflow on shared data and shared context, with clearer governance and cost controls over what gets collected, stored, and acted upon.

订阅研报添加微信 LCD123321Cad 04-21

Market signals underscore this is becoming a real source of demand. Grafana's Observability survey (2025) showed 64% of respondents now have less than five observability technologies (up from 44% last year) and ETR's 2025 survey indicates more than half of respondents are planning to consolidate further over the next two years. The market is increasingly responding to that same trend: Cisco's acquisition of Splunk pulled a scaled observability asset into a broader networking, security, and operations platform; Palo Alto's acquisition of Chronosphere reflects the growing value of telemetry inside security control planes; Snowflake's acquisition of Observe points to observability becoming more relevant inside data and AI stacks. In our view, consolidation is changing the basis of competition from feature leadership in isolated tools to ownership of shared telemetry, context, and response workflows across the estate. Vendors increasingly need to prove they can consolidate spend, displace adjacent tools, and own a broader share of the telemetry estate.

Exhibit 8: Number of Observability Technologies In Use
Buyers continue to reduce the number of tools in their observability stack, favoring full-stack observability vendors

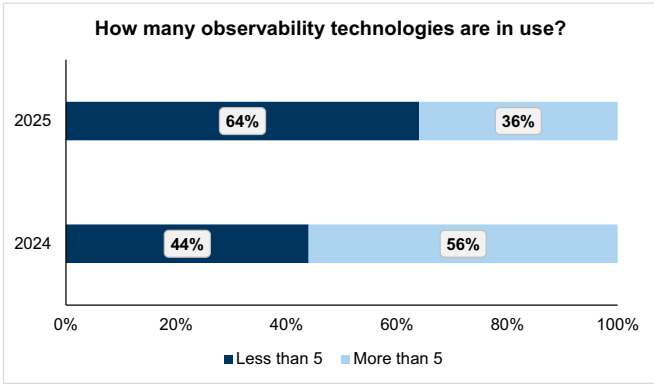

Exhibit 9: Health, Performance and Analysis Market
Observability market remains highly fragmented with 40+ vendors, half of which generate >$200M in Revenue

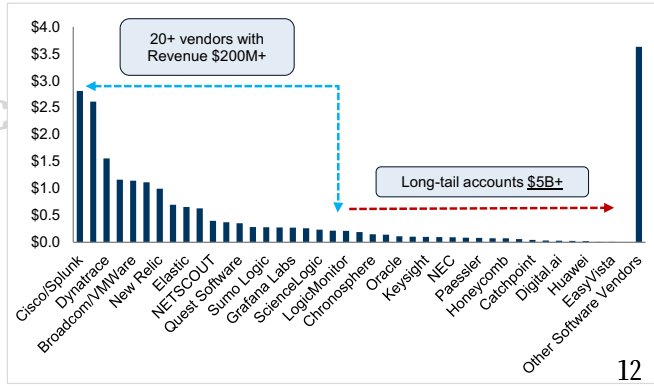

The setup across both markets remains attractive, shaped by a mix of rising infrastructure intensity, faster data creation and a broader enterprise need to manage complexity more effectively. As cloud environments scale, architectures become more distributed, and digital workflows became more data-intensive, customers are dealing with a larger operating and information footprint at the same time. That is the common thread underpinning both Dynatrace and Elastic. Complexity is rising across the stack and spend is increasingly following the platforms that can help customers monitor it, search it, secure it, and act on it.

订阅研报咨询微信 LCD123321Cad 04-21

Infrastructure growth remains the clearest tailwind to observability, with both direct and indirect benefits to the category. The Infrastructure (IaaS) market is already $190B+ and expected to grow at a near-20% rate through CY29 (Gartner), which provides a healthy baseline for observability given how closely those budgets have historically tracked underlying infrastructure spend. The direct benefit is straightforward: more workloads, services, and dependencies create more telemetry and environments to monitor. The indirect benefit is where we think the opportunity becomes more interesting. As those environments become harder to operate manually, the value shifts towards platforms that can compress diagnosis, improve prioritization and automate a greater share of remediation. **We therefore think the category can sustain mid-teens-plus growth even before considering a more meaningful AI-driven uplift, and why the wallet opportunity can expand over time as observability platforms like Dynatrace take on more of the operating burden than simply monitoring volume.**

订阅研报咨询微信 LCD123321Cad 04-21

Exhibit 10: Traditional IaaS Market Forecast
Observability budget outlook looks favorable mid-term with tight linkage to growing compute infrastructure spending (+18% CAGR through CY29)

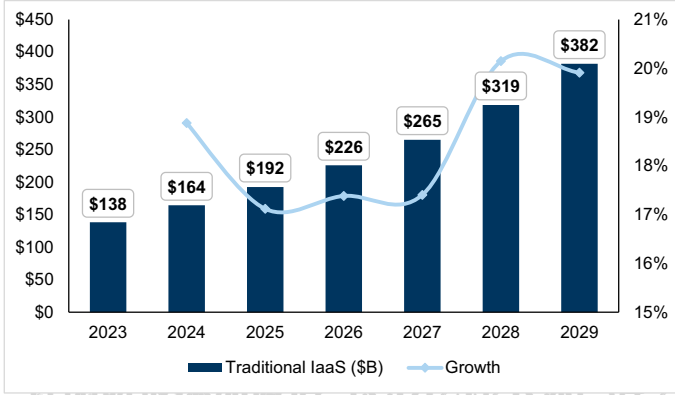

订阅研报咨询微信 LCD123321Cad 04-21

Data growth is the parallel tailwind on the search and analytics side, particularly as AI raises the value of retrieval. IDC projects total data generation to reach 527 ZB's by 2029, implying a +25% 5-year CAGR, with unstructured data remaining the dominant share of the information base. That aligns with the growing share of enterprise context now sitting outside traditional systems of record – in documents, tickets, chats, logs, alerts, transcripts and other free-form data sources. As enterprises move from AI experimentation toward production, retrieval becomes central because Copilots and Agents are only as useful as the context they can access. Elastic fits naturally here. The

company sits at a layer that is becoming increasingly strategic: retrieve relevant context, and make that information usable inside operational, security and AI-driven workflows.

Exhibit 11: Structured/Unstructred Data Generation Forecast
Data growth compounding mid-20's; supportive for data platforms operationalizing unstructured data for AI

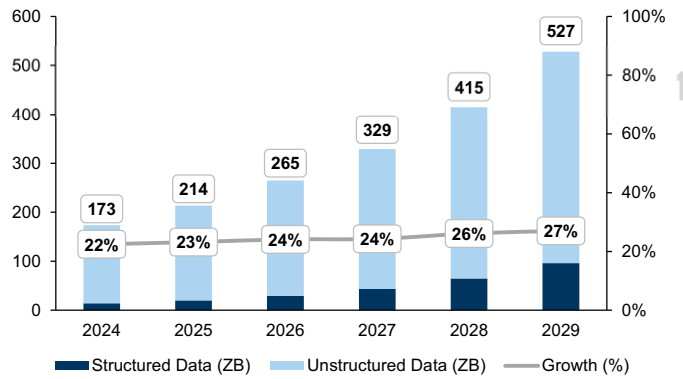

Source: IDC, Goldman Sachs Global Investment Research

Observability and SIEM remain distinct markets, but both are becoming more strategic as enterprises rely more heavily on telemetry and event data to run and secure modern environments. Observability budgets are supported by a very practical need: application environments are becoming harder to operate, outages are more expensive and customers increasingly want to reduce time between detection and remediation. SIEM budgets are supported by a similarly practical need on the security side: event volumes are rising, threats are moving faster, and security teams need better ways to search, correlate and prioritize what matters across larger data sets. The common thread is that logs, traces, runtime context and machine data are becoming more central to both operational and security workflows. That overlap is more directly supportive of Elastic today, but it also creates a longer-term tailwind for Dynatrace as runtime context, automation, and app-security workflows become more important.

Exhibit 12: Health, Performance and Analysis (Observability) Market Forecast
Observability spend remains supported by workload growth, rising telemetry and greater automation demand

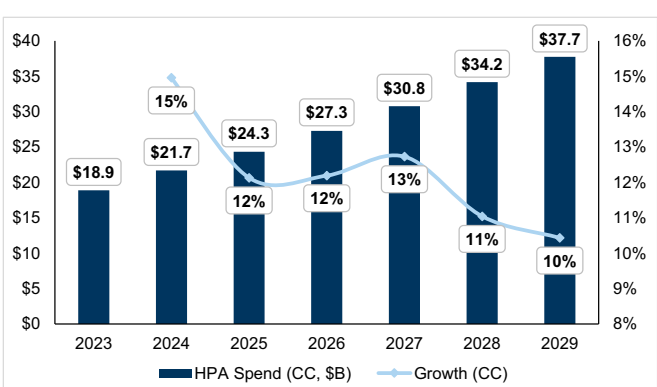

Exhibit 13: SIEM Market Forecast
The category contines to benefit from growing machine-data intensity and security complexity

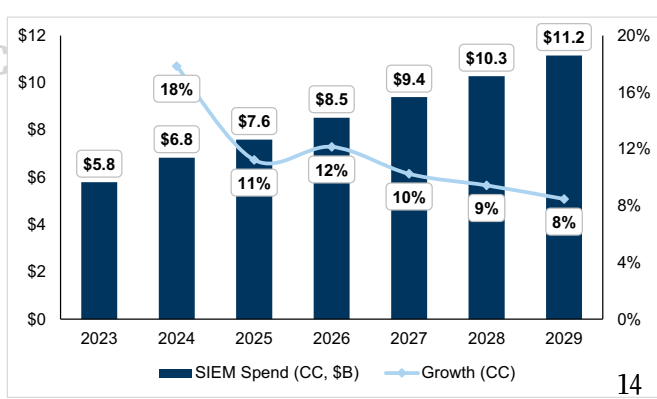

Our review of the private ecosystem suggests the more meaningful startup signal today is coming from observability, where customer pain points remain acute and the market is still actively searching for better answers around telemetry, economics, pipeline control, workflow simplicity, and AI-assisted operations. We believe that helps explain why the private landscape remains crowded and well-funded across several distinct wedges. Vendors such as Cribl, Onum (acquired by CrowdStrike) and Bindplane (acquired by Dynatrace) are attacking telemetry routing and data control, Clickhouse and Coralogix are pressing on lower-cost storage and analytics, while Honeycomb, Traversal, and Resolve AI are pushing the category toward faster diagnosis, intelligent investigation, and more automated operations. The broader takeaway is that these companies are not all trying to build a direct full-platform replacement. In most cases, they are targeting specific parts of the workflow where customers remain most sensitive to cost, complexity, or time-to-resolution, even as incumbent platforms remain deeply embedded at the center of enterprise environments.

请加微信 CD123210 获取

Exhibit 14: Observability Private Startup Ecosystem

Company	Founding Year	Valuation	Valuation Source	Strategic Angle	Why It Matters
Grafana Labs	2014	$6B (As of Aug-24)	TechCrunch	Open / Composable Observability	Set the pace on open, interoperable observability; makes portability, OSS support, and pricing discipline table stakes.
Cribl	2018	$3.5B (As of Aug-24)	Company Data	Data Pipeline / Telemetry Routing Layer	Controls the data flow between sources and destinations; gives buyers leverage to decouple ingestion from any single vendor.
ClickHouse	2021	$15B (As of Jan-26)	Company Data	Real-Time Analytical Database Engine	Provides the performance backbone for next-gen observability stacks; underpins cost-efficient log and trace storage at scale.
Coralogix	2014	$1B (As of Jun-25)	Company Data	Streaming-First / Cost-Optimized Observability	Proves that in-stream analytics can slash storage costs; pressures incumbents on price without sacrificing data fidelity.
Sentry	2011	$3B	Company Data	Developer-First Performance Monitoring	Owns the developer on-ramp to observability via open-source error tracking; expands TAM by bridging into full APM territory.
Honeycomb	2016	$0.5B (As of Apr-23)	TechCrunch	High-Cardinality / Observability 2.0 Pioneer	Redefined how engineering teams explore production data; championed the shift from siloed pillars to unified, queryable telemetry.
Arize AI	2020	--	--	AI / ML Model Observability & LLM Evaluation	Purpose-built for monitoring AI models in production; captures the emerging budget line for LLM and agent observability.
Resolve AI	2024	$1B (As of Dec-2025)	TechCrunch	Autonomous SRE / AI-Native Incident Response	Applies agentic AI to automate on-call and root-cause analysis; tests whether AI can replace manual SRE workflows at scale.
Traversal	2023	--	--	AI SRE Agent / Causal Root-Cause Analysis	Uses causal ML to pinpoint root causes across massive telemetry; addresses the growing complexity from AI-generated codebases.
Dash0	2023	$1B (As of Mar-2026)	Company Data	OpenTelemetry-Native / Agentic Observability	Built OTel-native from day one with autonomous AI agents; challenges incumbents on vendor lock-in, pricing, and data portability.

around vector, hybrid retrieval and AI-native architectures versus legacy keyword-only approaches. At the same time, competitive pressure is not limited to that cohort. AI-native application vendors like Glean can compete for portions of enterprise search and discovery budgets by owning the end-user workflow, while broader developer and data platforms are increasingly embedding retrieval capabilities natively, which could limit how much value ultimately accrues to standalone features. In our view, this is the key distinction versus observability: search is still an active market, but the startup landscape is more fragmented, with no clear path yet to platform consolidation or category leadership.

Exhibit 15: Vector Database / Search Platforms

Company	Founding Year	Valuation	Valuation Source	Strategic Angle	Why It Matters
Glean	2019	$7.2B (As of Jun-2025)	Company Data	Enterprise Knowledge / Assistant Layer	Shows AI assistants can compete for portions of search and discovery budgets by owning the end-user workflow.
Pinecone	2019	$750M (As of Apr-2023)	Company Data	Managed Vector Database for AI Applications	Provides the retrieval backbone for RAG and LLM apps; captures infrastructure spend as every AI product needs vector search.
Algolia	2012	$2.25B (As of Jul-2021)	Company Data	API-First Search & Discovery Platform	Proved search-as-a-service at scale for e-commerce and SaaS; now pivoting to AI-native search to defend a large install base.
Qdrant	2021	--		Open-Source Vector Search Engine	Offers high-performance, self-hosted vector search with advanced filtering; appeals to cost-sensitive and privacy-focused teams.
Weaviate	2019	--		Open-Source AI-Native Vector Database	Combines vector and structured search in a single schema layer; positions itself as the knowledge graph for AI-native applications.
Zilliz/Milvus	2017	--		Open-Source Vector Database at GPU Scale	Backs the most-starred open-source vector DB; targets massive-scale enterprise workloads where raw throughput is the bottleneck.

订阅研报添加微信 LCD12332[illegible]

observability, AI, and operations. Two conversations in particular - with Grafana Labs at the Software Bus Tour and Traversal at our Private Software Conference - were especially informative in illustrating how innovation at the private layer is responding to the same structural shifts reshaping the public observability landscape. While these companies operate at very different stages of maturity, both offer useful signals on how architectures, economics, and workflows are evolving as AI increases system complexity, raises query intensity, and places greater emphasis on cost discipline and deterministic outcomes. The takeaways below highlight where we see these dynamics emerging most clearly.

订阅研报添加微信 LCD12321Cad 04-21 Cad 04-21

- **Our discussion with Grafana Labs highlighted how observability architectures are adapting as AI-driven query patterns and cost discipline become first-order concerns.** Grafana's evolution from Prometheus-native tooling to a fully OpenTelemetry-aligned platform mirrors the broader standardization underway across enterprise telemetry. Where differentiation is emerging is economics and control: dynamically filtering telemetry before it reaches storage directly addresses the data-volume problem that has driven observability costs higher as environments scale, particularly as AI agents push query intensity well beyond human baselines. Coupled with a graph-based data model designed for agentic workflows and an Assistant that has demonstrated credible root-cause accuracy in live incidents, Grafana's openness-plus-economics positioning is increasingly resonating with buyers reassessing total cost of ownership and lock-in risk alongside feature completeness.
- **Traversal represents an earlier-stage but conceptually clear signal in the shift toward faster, more deterministic operations.** While much of today's AI tooling in observability remains correlation-driven, Traversal is built around a causal search engine operating on a live production world model, enabling it to trace root-cause chains rather than surface probabilistic suspects. That distinction is especially relevant in SRE workflows where explainability, repeatability, and MTTR reduction are paramount - particularly for uptime-sensitive customers like financial services, where downtime costs are immediately quantifiable. Near term, Traversal's decision to sit on top of existing observability stacks lowers adoption friction; longer term, the architecture points toward deeper infrastructure ownership as the demands of maintaining and querying the underlying data layer grow. Together with Grafana, Traversal underscores how innovation at the private layer is increasingly concentrated on the economics, control planes, and operational determinism that incumbent platforms will need to address.

«¢– x b¥m0R _®Oá LCD123321Cad 04-21

AI Impact Framework

Applying the GS AI Impact Framework to Dynatrace and Elastic. In conjunction with our analysis of the private startup ecosystem and the emerging competitive threats reshaping observability, search and security, we apply the GS AI Impact Framework to Dynatrace and Elastic to assess how each platform is positioned against the most important vectors of AI-driven disruption. The framework evaluates six dimensions: orchestration risk, monetization model, system of record ownership, data and integration moat, AI execution and budget alignment. Both companies benefit from consumption-based models that scale with data intensity and workload complexity, a

structural advantage as AI expands the volume and urgency of machine-generated data. Dynatrace's deep workflow entrenchment, proven causal AI engine and expanding agentic capabilities positions it well across most dimensions of the framework. Elastic's foundational role in search and retrieval becomes more strategically important as AI raises the value of context and unstructured data, and the company has made considerable product progress building toward higher-value workflows, we think the degree to which it can capture that opportunity relative to broader platforms remains one of the more important debates in the story. Our detailed assessment is outlined below.

请阅读最后评级声明 LCD12221Ced 04.21

Exhibit 16: GS AI Impact Framework
Both platforms are well-positioned against the most important vectors of AI-driven disruption, with consumption-based models that scale with data intensity and a growing suite of AI-native capabilities across observability, search and security

AI Impact Framework	Dynatrace	Elastic
Orchestration Risk Could an external AI layer capture value?	**Low-Moderate.** Davis AI embedded in operational workflows, limiting abstraction risk. Emerging AIOps vendors targeting the remediation layer directly, making Dynatrace Intelligence's execution increasingly consequential.	**Low-Moderate.** Search and retrieval are foundational infrastructure layers, but agent frameworks and Hyperscaler-native stacks could abstract value if retrieval becomes commoditized inside broader platforms.
Monetization Model Is revenue tied to seats or assets/data/consumption?	**Insulated.** DPS consumption model benefits directly from AI-driven infrastructure complexity and rising telemetry volumes, expanding the addressable workload opportunity.	**Insulated**. Consumption tied to data volume and query intensity; AI workloads are inherently data-intensive, driving incremental demand for search and retrieval infrastructure.
System of Record Ownership Who owns execution and governance?	**Strong.** Owns the operational telemetry layer for large enterprises; Davis AI sits at the center of diagnosis, remediation and autonomous operations.	**Moderate.** Elasticsearch is the data and retrieval foundation across search, observability, and security. Ownership of the action and governance layer above the data tier remains less established.
Data and Integration Moat How embedded is the platform?	**Deep.** OneAgent and Smartscape create deep technical entrenchment; switching costs increase as customers rely on Davis AI for automated remediation; OTel's rise adds instrumentation flexibility, gradually tempering the proprietary moat.	**Moderate.** Deeply embedded in log analytics and search workflows; open-source heritage drives broad adoption and developer mindshare but limits exclusivity relative to more proprietary architectures.
AI Execution Is AI embedded into core workflows?	**Meaningful Progress.** Davis AI has powered automated RCA and remediation at scale for years. Dynatrace Intelligence adds natural language querying, AI-assisted workflow automation and agentic resolution across observability and security.	**Progressing.** Elastic supports vector search, hybrid retrieval, reranking, and native inference. Agent Builder and Workflows extend the platform into AI agent construction and automated action across enterprise systems.
Budget Alignment Does AI increase category urgency?	**Positive.** AI increases infrastructure complexity, telemetry volume, and remediation demand - all expanding Dynatrace's relevance and addressable opportunity.	**Positive.** AI expands the strategic importance of retrieval, context engineering and unstructured data handling, though rising competitive intensity across all three end-markets may temper the net benefit.

Dynatrace combines deep roots in APM with an increasingly broader full-stack and AI-powered platform. The company's heritage is in application performance monitoring (APM), where it built a strong position with large, complex enterprises that needed visibility into highly distributed, mission-critical environments. Over time, that foothold has expanded well beyond classic APM into infrastructure observability, digital experience monitoring, log management, application security and automation, with Grail and Davis AI now serving as the architectural backbone for a more unified platform. Importantly, Dynatrace is no longer just selling monitoring; it is increasingly selling a system designed to help enterprises understand, prioritize, and act across complex environments. In our view, that combination of enterprise entrenchment, platform breadth, and a growing emphasis on AI-driven workflows gives Dynatrace a more durable strategic position than many narrower observability vendors.

订阅研报添加微信 LCD12321Cad 04-21

Exhibit 17: Gartner's Product or Service Score by Use Case (Observability)

Dynatrace ranks #1 or #2 across all measured observability use cases, underscoring its market leadership

Capability	Rank (n=20)	Score (Out of 5)
SRE (Site Reliability Engineering)	#1	4.30
Platform Operations	#2	4.28
Software Engineering	#2	4.29
Business Insights	#1	4.30
Cost Optimization	#1	4.32
AI Engineering	#1	4.29
Average	**#1**	**4.30**

Source: Gartner

Exhibit 18: Dynatrace is the #3 scaled player in Observability

Observability market remains highly fragmented; Dynatrace positioned as a consolidator with platform depth and breadth

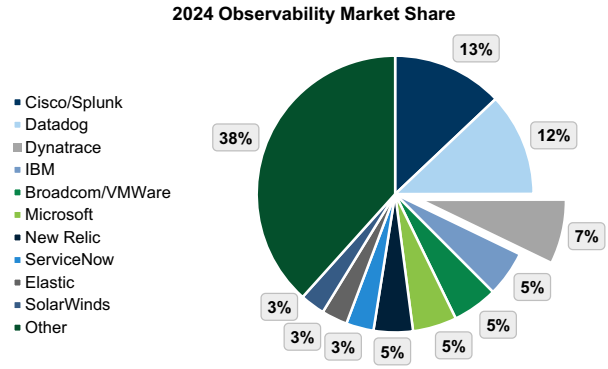

Source: Gartner

Our Buy rating reflects a platform moving in the right direction and a clearer line of sight to potential growth acceleration. Dynatrace has spent roughly the last two years putting several of the key pieces in place. The GTM changes increasingly look to be in the rearview, and we think that is showing up in the numbers, with CC NNARR growing double-digits for three straight quarters and the model carrying better stability than it did a year ago. Dynatrace Platform Subscription (DPS) has also made the platform easier to adopt and should begin to create a clearer renewal tailwind in FY27 as the earliest cohorts cycle, which gives the company a more visible path to accelerating ARR growth. Logs is also emerging as a meaningful upside node, crossing $100M in annualized consumption (+100%), helped by the same broader platform adoption dynamics. Grail and Davis AI reinforce the strategic case by giving Dynatrace a stronger data and AI foundation as customers look for more unified, automated operations. With a healthy non-GAAP OpM profile (~30%) and expanding GAAP profitability, we think the stock offers a compelling balance of improving growth potential and healthy FCF generation.

订阅研报添加微信 LCD12321Cad 04-21

Exhibit 19: Constant-Currency NNARR Growth
GTM changes implemented in FY25 drove volatility in NNARR performance; maturation of new model supporting better consistency

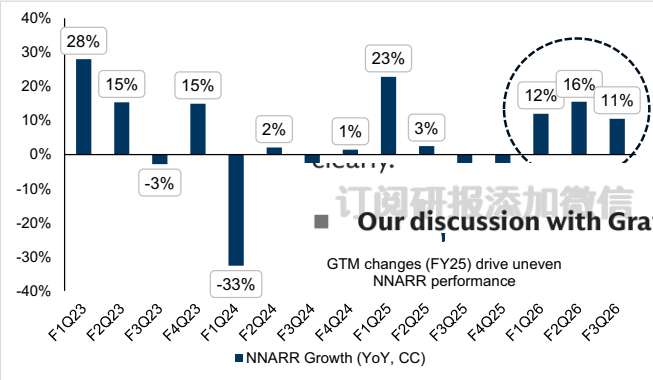

Source: Company Data

Exhibit 20: Constant Currency ARR Growth
Recent consistency in NNARR driving three consecutive quarters of stable ARR growth

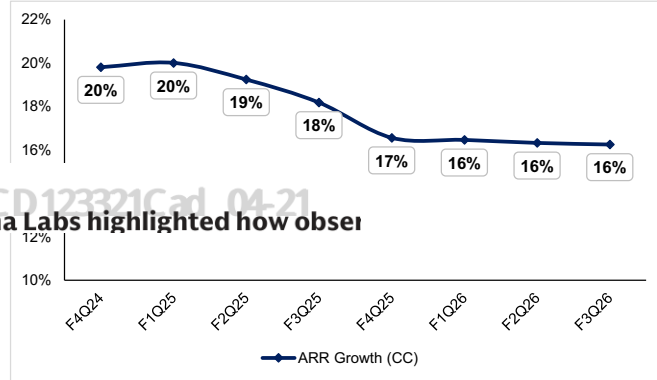

Source: Company Data

订阅研报添加微信 LCD123321Cad 04-21

Dynatrace Through the Lens of the GS Framework for Investing in Software

Dynatrace screens well across the GS Framework for Investing in Software. The company addresses a large and growing observability market underpinned by structural demand – cloud complexity, AI workload proliferation, and the convergence of observability and security are all expanding the opportunity set in ways that favor platforms with deep technical differentiation and broad workflow coverage. Dynatrace's combination of full-stack instrumentation, proprietary causal AI and a contract structure that deepens consumption over time positions it as a consolidating platform rather than a point solution, with switching costs that increase as customers rely more heavily on Davis AI for automated remediation and operational decision-making. Unit economics are already strong, with Non-GAAP OpM's near-30% and a Rule of 40 profile consistently above 40, which compares favorably to most software peers at a similar stage of growth. Looking ahead, the DPS renewal cycle and an underpenetrated logs opportunity provide a more visible path to growth acceleration, while the extension into agentic operations through Dynatrace Intelligence broadens the platform's relevance as the category shifts toward autonomous IT. Taken together, these attributes support a defensible competitive position and a compelling risk-reward at mid-teens EV/FCF.

订阅研报添加微信 LCD123321Cad 04-21

Exhibit 21: Mapping Dynatrace to the GS Framework for Investing in Software

	OUR VIEW ON DYNATRACE	
8 KEY FACTORS	**TANGIBLE**	**INTANGIBLE**
TAM	$51B TAM for Observability & $14B TAM for Security; AI workload monitoring expands addressable opportunity	TAM expansion via Logs, AppSec, AIOps, DevOps Focus on complex hybrid estates & AI automation
SECULAR THEMES	Traditional compute infrastructure spend compounding at near-20% rate (Gartner); multi-year tailwind to observability budgets	Cloud complexity, microservices proliferation, AI workload growth, vendor consolidation, security-observability convergence
ENTRY/EXIT POINTS	Stock trades at 4x EV/Sales FY28 and 14x EV/FCF FY28	We see a favorable risk-reward at current levels given DPS renewal tailwind, underpenetrated logs opportunity and durable FCF profile
LT FRAMEWORK RETURN	Compounded Returns: Bear: -5%, Base: 11%, Bull: 21%	Path to accelerating ARR growth as DPS base matures; potential for sustainable mid-to-high-teens growth with best-in-class FCFMs
UNIT ECONOMICS	Operating at Rule of 40+ with Non-GAAP CAC at ~2x vs. broader software at 2-3x	DPS renewal cycle drives natural expansion within existing accounts; E2E observability deployments driving higher ACV deals
PLATFORM/ BEST OF BREED	Full-stack observability spanning APM, ITIM, Logs, DEM, and AppSec; Davis AI purpose-built for causation / remediation	Dynatrace Intelligence extending into agentic operations with domain-specific agents for SRE, development, security
COMPETITIVE MOAT	OneAgent instrumentation, Smartscape topology, Grail and Davis AI create deep technical entrenchment; consistent mid-90's GRR	Deterministic causal AI difficult to replicate; switching costs rise as Davis AI becomes central to operational decision-making
ESG	Helps enterprises reduce IT inefficiency, energy consumption and infrastructure waste	Platform directly supports digital reliability and sustainability mandates large enterprises are increasingly measured against

Source: Gartner, Company Data, Goldman Sachs Global Investment Research

Platform Strategy: Leveraging Enterprise Entrenchment and Hybrid Complexity

End-to-end observability remains Dynatrace's core strength, and the market is moving in its direction. Dynatrace built its core franchise in APM, but the platform has expanded well beyond that starting point into infrastructure, user experience, logs, security and automation. We believe that breadth is strategic because organizations are increasingly looking for fewer tools, tighter operational workflows, and broader visibility across complex environments. Dynatrace is well-positioned for that shift, in our view, because it has long been strongest in large enterprises running a mix of on-prem, legacy and cloud workloads rather than only clean cloud-only environments. The combination of broad product coverage and credibility in hybrid estates helps support the platform's stickiness (95% GRR) and reinforces Dynatrace's appeal as a consolidation choice for customers looking for a more unified way to run operations. Management has framed end-to-end observability as an emerging shift two years ago but a prevalent one now, and we think that is consistent with the recent momentum in strategic pipeline growth (up ~50% YoY in F1Q26 and +45% YoY in F2Q26).

That position should support both durable expansion and high-quality new-logo activity over time. Dynatrace remains geared toward larger enterprises rather than a high-velocity, developer-led motion, which means new-logo volume trails peers such as Datadog. More important is the quality and yield of those accounts once they land. Average land size was $200K versus ~$500K in ARR/Customer across the installed base (F3Q26), which points to meaningful expansion headroom if newer cohorts mature in a way that is even directionally consistent with older ones. Importantly, Management has noted that today Dynatrace typically observes only ~20-30% of customer workloads, largely those deemed mission-critical and customer-facing, with broader environments often monitored via lighter-weight or home-grown tooling. As AI-native and agentic workloads proliferate, we think that definition of mission-critical naturally expands, creating a tangible path for workload-based expansion over time. The GTM changes over the last two years fit naturally with that backdrop: sharper account focus and greater customer intimacy should continue to improve Dynatrace's ability to expand

within these large environments; meanwhile, DPS lowers the friction to consume more of the platform over time. In our view, that combination supports a healthier and more durable expansion path than logo counts alone would suggest.

Exhibit 22: New Logos Added
New logo activity has trended lower vs. prior year as the company leans into consolidation deals with higher ACV potential

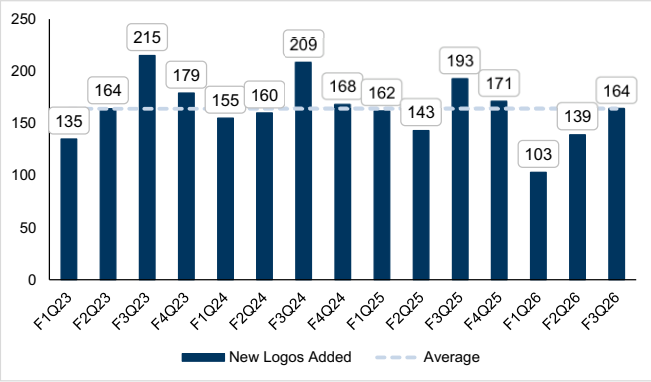

Source: Company Data

Exhibit 23: Avg. New Logo Land Size and Avg. ARR Per Customer
Land size reached a record $200K in F3Q26 (vs. $160K TTM), an early sign of momentum in larger platform consolidation deals

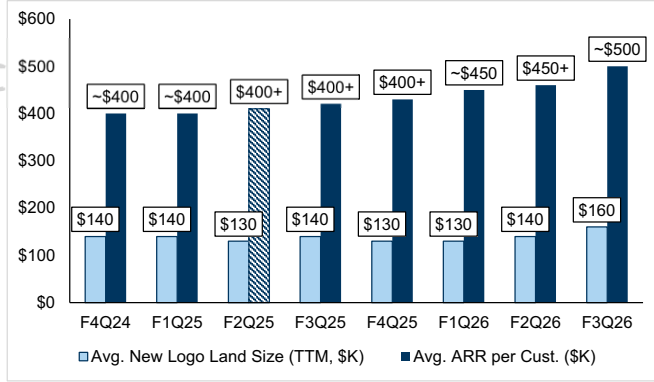

Source: Company Data

Unpacking the Logs Product Cycle

Grail and logs give Dynatrace a stronger path to workload consolidation. Grail is central to the story because it gives Dynatrace a common data foundation across logs, metrics, traces, and security data rather than requiring customers to work across separate systems. The logs business is the clearest example of how that architecture can create a new leg of growth. Dynatrace did not have a viable logs product until Oct. 2024 but the ramp since then has been notable, with logs now above $100M in annualized consumption and growing 100%+ (F3Q26). The value-proposition is straightforward: when logs sit on the same architecture as metrics and traces, customers can analyze more of their environment in one place, reduce tool sprawl, and improve both performance and economics.

Logs are becoming more important because they improve both platform value and quality of automated layered on top of it. When logs sit alongside metrics and traces, teams can move from signal to context to root-cause faster and with less friction than they can across separate tools. That becomes more valuable as customers aim to standardize on fewer platforms and automate more of their operations. In our view, these dynamics are what makes logs strategically important for Dynatrace: they are not only another workload the company can win, but also a deeper source of context that should improve the value of the broader platform as customers push further toward AI-driven analysis and automation. We therefore see logs as both a direct growth driver and a practical way to deepen platform standardization over time.

Channel checks suggest Dynatrace is winning logs primarily as an APM/DEM attach - customers add logs to put more context around workflows they already run on the platform, which we think also sets up better downstream insights and automation via Dynatrace Intelligence. The same checks underscore that as logging scales, cost control becomes central, which is why more customers are looking to be deliberate about what data they send into premium back-ends. We think Dynatrace is well positioned in that environment because logs are still a smaller, growing piece of the platform rather than a

massive legacy estate customers are trying to cut back. Bindplane should be an incremental tailwind by enabling more filtering and governance earlier in the data flow, improving the economics of scaling logs as adoption broadens.

Exhibit 24: Timeline of Logs Disclosures
Logs crossed the $100M milestone in F3Q26, still growing 100%+; constructive on durable growth potential as DPS reduces adoption friction

Period	Disclosures
F3Q26	Annualized logs consumption officially surpassed the $100 million milestone The logs business continued to grow at over 100% year-over-year
F2Q26	Logs remained the fastest-growing product category, growing well over 100% year-over-year The number of customers spending over $1 million annually on logs doubled compared to the prior year The number of customers spending over $500,000 annually on logs increased nearly 4x year-over-year
F1Q26	Logs consumption increased by 36% sequentially and well over 100% year-over-year Three 7-figure ACV wins in the quarter each included planned annualized logs consumption of at least $1 million Major airline customer exceeded several million dollars in annualized logs consumption <18 months on the platform
F4Q25	Over one-third of the total customer base was leveraging the logs solution Customers leveraging logs increased by 18% compared to the prior quarter ~50% of all new logos added in Q4 deployed logs in their initial implementation, up from ~20% in year-ago period
F3Q25	The company reached well over 1,000 customers leveraging log solutions The logs customer base grew 17% sequentially compared to the prior quarter More than 50% of the new logos added in TTM leveraged logs in their initial deployments
F2Q25	Nearly 25% of the total customer base was using the logs product at this time Customer Growth: The number of logs customers increased by ~20% sequentially
F1Q25	Logs consumption was noted as growing well above 100% year-over-year
F4Q24	Dynatrace reported having roughly 600 customers leveraging the logs solution This 600-customer mark represented a 30% quarter-over-quarter increase in the number of log customers Exiting fiscal 2024, consumption for the emerging logs offering was growing very rapidly at north of 100%
F3Q24	The company saw a 50% sequential quarter-over-quarter increase in paying log management customers There was a 30% increase in log management Proof of Concepts (POCs) during the quarter

Source: Company Data

Causal AI and Davis: Delivering Trusted Answers Before Guided Action

AI should reinforce Dynatrace's relevance in large, hybrid enterprise environments. Dynatrace has been embedded for years in enterprises running a mix of legacy systems, on-prem infrastructure, cloud workloads and increasingly AI-driven applications. We view this as an important advantage since AI is being layered onto existing estates, not built into a clean slate. In practice, this tends to increase operational complexity: release cycles compress, dependencies multiply, and the need for trusted visibility across the stack becomes more important. In our view, this plays directly to Dynatrace's strengths in end-to-end observability and should support continued demand from customers looking to manage AI-driven change inside already-complex environments.

订阅研报请加微信 TCJD12321Cao 14-21

Dynatrace's architecture unifies data through a single system model, eliminating traditional telemetry silos by leveraging Grail as a unified fabric and OneAgent for deep, automated discovery. While OneAgent remains a core differentiator, the

steps become automated.

Ultimately, we view this as a tangible opportunity for Dynatrace to capture incremental value. By enabling customers to streamline operational workflows and enhance engineering efficiency via Dynatrace Intelligence, the company can effectively participate in the long-term cost savings derived from reduced labor intensity in SRE and platform engineering functions.

Exhibit 26: Dynatrace Intelligence Platform Architecture

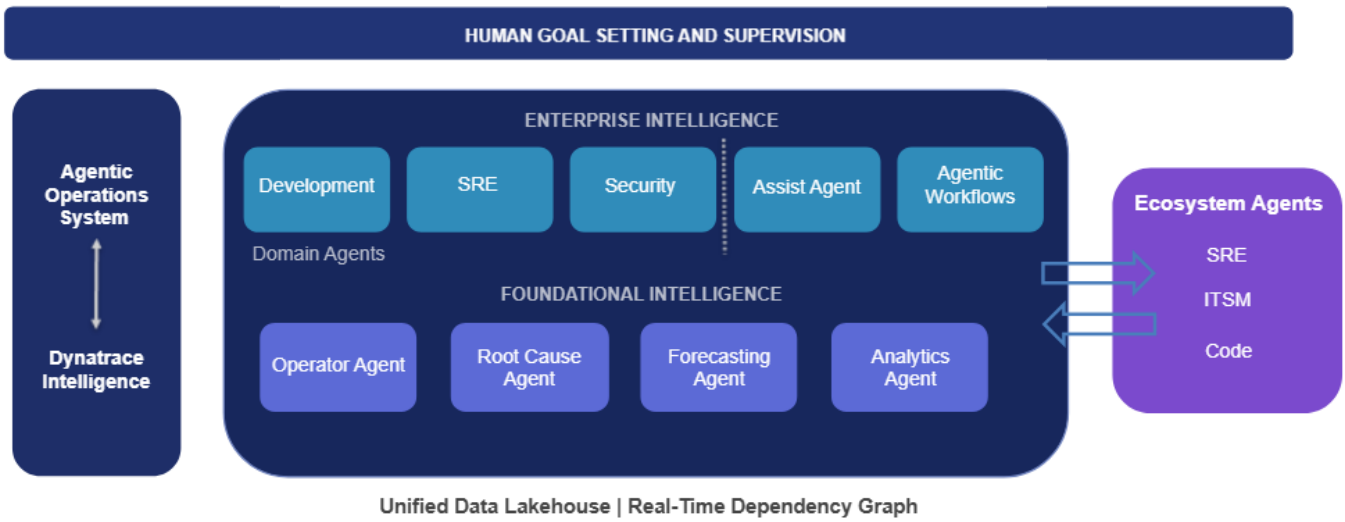

Source: Company Data, Goldman Sachs Global Investment Research

订阅研报添加微信 LCD123321Cad 04-21

The DPS Transition: Bridging the Gap Between Consumption and Reported ARR

Dynatrace's transition to DPS creates a timing lag between usage and reported ARR. Dynatrace Platform Subscription (DPS) customers commit to a minimum annual spend and consume against a common platform rate card rather than buying individual products under a traditional SKU-based subscription. If usage rises above that annual commitment, customers can either continue consuming on-demand at the same unit price or step up the contract earlier, though Management has indicated that customers earlier in the three-year contract term more often stay on on-demand consumption (ODC) and defer the contractual reset. Excess usage still flows through revenue, but it is excluded from ARR until the commitment is rebased higher. With DPS now representing 75% of ARR, more of Dynatrace's growth is flowing through a structure where usage can run ahead of committed ARR. Management's disclosures punctuate this: annualized dollar consumption is growing 20%+ and DPS customers are consuming roughly 2x faster than legacy SKU-based, while reported ARR has been stable at +16% CC growth through FY26.

订阅研报添加微信 LCD123321Cad 04-21

Dynatrace is sustaining an industry leading margin profile while deliberately investing for growth acceleration. Management has signaled that 29% non-GAAP OpM and 32% pre-tax FCFM should be viewed as a floor as Dynatrace reinvest back into the business – a commitment supported by emerging AI efficiencies and productivity gains across business functions that are effectively self-funding the investment cycle. Those dollars are being directed with clear intent. R&D remains the primary area of resource allocation and headcount growth, with the recent launch of Dynatrace Intelligence a tangible proof point of what that investment is producing. On the GTM side, Management continues to invest in S&M while generating leverage – driven by a

Feb. 2026 Pull observability data into the data cloud; logs, metrics, and event enterprise data rather than a separate monitoring layer. Signals the class workload inside broader data and AI architectures.

订阅研报添加微信 LCD123321Gao 04-23

strike teams across Logs, AppSec and Digital Experience Monitoring that are driving high-margin platform attach. The combination of partner leverage, focused sales capacity, and the low-friction DPS consumption model is allowing Dynatrace to invest for growth while holding margins at levels few software peers can match. We view the current posture as the right balance given where Dynatrace is in its growth cycle. Long-term, we see room for incremental margin expansion as AI efficiencies compound, the business scales, and the DPS expansion motion becomes increasingly dominant.

Exhibit 30: Non-GAAP Operating Margin and Free Cash Flow Margin
Dynatrace has maintained high-20's OpM's even through a growth investment cycle, underscoring strong operational discipline

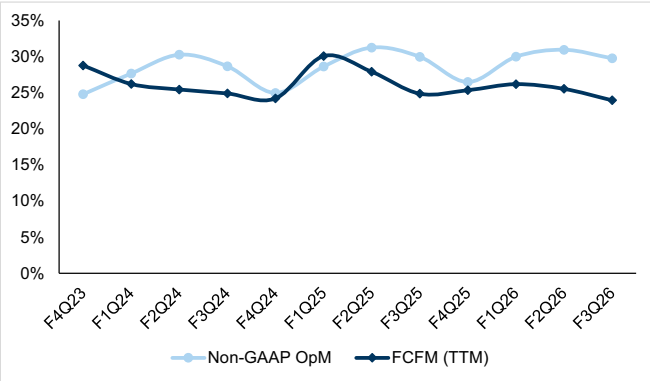

Source: Company Data

Exhibit 31: Non-GAAP R&D Spending
R&D is a key investment priority as Dynatrace extends the platform toward agentic operations, security and business observability

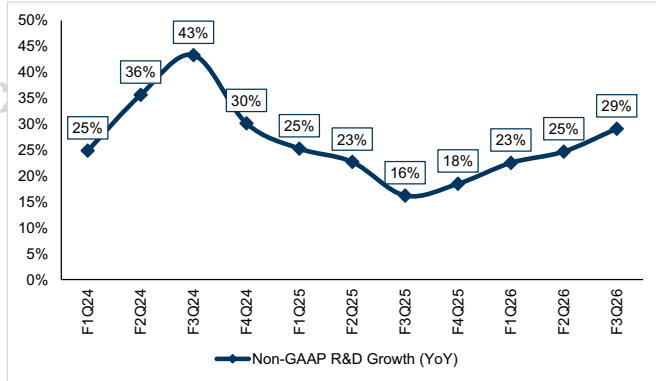

Source: Company Data

Dynatrace has de-rated to levels that we view as increasingly compelling given the quality of the business and improving growth outlook. Against a backdrop of severe multiple compression across software on AI disruption fears, Dynatrace's NTM EV/FCF multiple has contracted to 16x (from 21x Jan. 2026) - a 31% discount to the one-year average and 50% discount to three-year average. We are not dismissive of what is driving that compression – the broader software de-rating reflects ongoing investor debate around terminal value in an AI-disrupted world and GAAP profitability as a measure of business quality. On both dimensions, we think Dynatrace stacks up better than most. The platform operates at the operational and remediation layer of enterprise IT, where AI is more likely to expand demand than displace it, while GAAP profitability is already a reality with 8% GAAP OpM's in FY26. At current multiples, we think the market is underappreciating the combination of 30%+ pre-tax FCFM's, a balanced Rule of 40+ profile, and a growth setup we expect to improve as the DPS renewal cycle matures and log management scales.

订阅研报添加微信 LCD123321Gao 04-23

Exhibit 32: Dynatrace EV / FCF (NTM)
DT has de-rated to 16x EV/FCF (NTM), -31% below the 1Y average

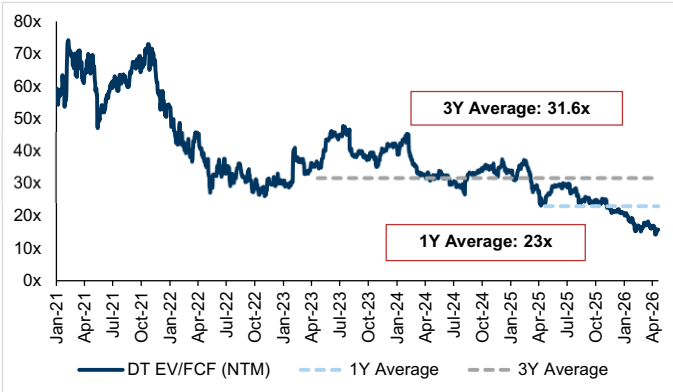

Source: FactSet

Estimates and Management Guidance

We outline our estimates and Management guidance for Dynatrace below in Exhibit 33

订阅研报添加微信 LCD12332ICad 04-21

Exhibit 33: Estimates vs. Consensus

All figures in $ mns	4Q26 (E) Guidance			4Q26 (E) GSe	4Q26 (E) Cons.	FY26 (E) Guidance			FY26 (E) GSe	FY26 (E) Cons.	FY27 (E) GSe	FY27 (E) Cons.	FY28 (E) GSe	FY28 (E) Cons.
Total ARR				$2,058	$2,054	$2,053	--	$2,061	$2,058	$2,054	$2,355	$2,356	$2,696	$2,704
YoY				*19%*	*18%*	*18%*	--	*19%*	*19%*	*18%*	*14%*	*15%*	*14%*	*15%*
QoQ				*4%*	*4%*									
Total ARR (CC)				$2,009	$2,010				$2,009	$2,016	$2,355	$2,350	$2,696	$2,683
YoY (CC)				*16%*	*16%*				*16%*	*16%*	*14%*	*14%*	*14%*	*14%*
NNARR (CC)				$77	$79				$273	$278	$297	$298	$341	$327
YoY (CC)				*4%*	*7%*				*10%*	*11%*	*9%*	*8%*	*15%*	*10%*
Total Revenue	$518	--	$523	$521	$521	$2,005	--	$2,010	$2,008	$2,008	$2,310	$2,305	$2,639	$2,637
YoY	*16%*	--	*17%*	*17%*	*17%*	*18%*	--	*18%*	*18%*	*18%*	*15%*	*15%*	*14%*	*14%*
QoQ	*0%*	--	*1%*	*1%*	*1%*									
Subscription Revenue	$493	--	$498	$496	$496	$1,917	--	$1,922	$1,920	$1,920	$2,214	$2,209	$2,536	$2,529
YoY	*16%*	--	*18%*	*17%*	*17%*	*18%*	--	*18%*	*18%*	*18%*	*15%*	*15%*	*15%*	*14%*
QoQ	*0%*	--	*1%*	*0%*	*1%*									
Professional Services Revenue				$25	$25				$88	$87	$96	$96	$104	$104
YoY				*17%*	*15%*				*15%*	*14%*	*10%*	*10%*	*8%*	*8%*
QoQ				*14%*	*12%*									
Non GAAP Gross Profit				$435	$436				$1,686	$1,689	$1,939	$1,940	$2,215	$2,217
Gross Margin				*84%*	*84%*				*84%*	*84%*	*84%*	*84%*	*84%*	*84%*
Non GAAP Operating Income	$133	--	$138	$136	$137	$582	--	$587	$585	$586	$684	$683	$800	$800
Operating Margin	*26%*	--	*26%*	*26%*	*26%*	*29%*	--	*29%*	*29%*	*29%*	*30%*	*30%*	*30%*	*30%*
Non GAAP EPS	$0.38	--	$0.39	$0.39	$0.39	$1.67	--	$1.69	$1.68	$1.68	$1.95	$1.91	$2.29	$2.20
GAAP EPS				$0.09	$0.15				$0.56	$0.63	$0.66	$0.88	$0.87	$1.07
Total Billings (Balance Sheet)				$816	$817				$2,142	$2,140	$2,457	$2,439	$2,792	$2,860
YoY				*14%*	*14%*				*20%*	*20%*	*15%*	*14%*	*14%*	*17%*
QoQ				*46%*	*46%*									
Cash Flow from Operations				$216	$215				$551	$554	$645	$646	$753	$741
CFO Margin				*41%*	*41%*				*27%*	*28%*	*28%*	*28%*	*29%*	*28%*
Free Cash Flow				$205	$204	$520	--	$525	$522	$522	$612	$609	$715	$705
FCF Margin				*39%*	*39%*	*26%*	--	*26%*	*26%*	*26%*	*26%*	*26%*	*27%*	*27%*

Total ARR, Total ARR (CC), NNARR (CC) Consensus - Visible Alpha Consensus Data, Rest- FactSet

Elastic combines a durable search franchise with a broader platform opportunity. Built around Elasticsearch, the company now spans Search and AI, Observability, and Security, and that footprint lines up with several of the structural shifts reshaping infrastructure software: unstructured data is proliferating, relevance is becoming more important in AI applications, open telemetry architectures are gaining ground, and observability and security are converging around shared telemetry and workflow layers. Elastic's heritage matters because it starts from a real architectural strength in search and unstructured data, with a large installed base, meaningful developer mindshare and platform that can be applied across multiple data-heavy workflows.

Exhibit 35: Stack Overflow Developer Survey: Which database environments have you done extensive development work in over the past year, and which do you want to work in over the next year?
Elastic has maintained steady developer mindshare over the past few years

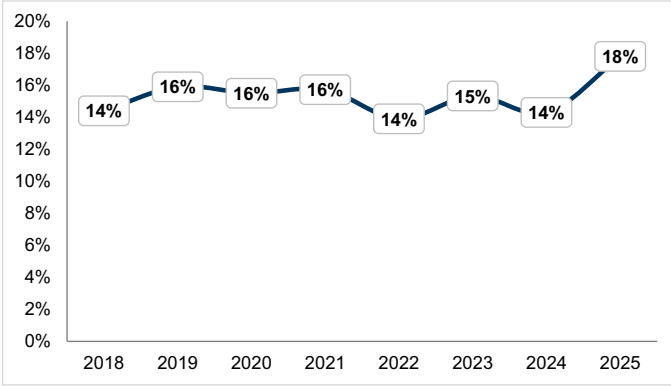

Source: Stack Overflow

Our Neutral rating reflects a strong platform direction and steadier execution, but still limited proof of a durable growth inflection. Elastic has moved well beyond its roots in search and log analytics, adding products such as Elastic Inference Service, Agent Builder, Workflows, AI Assistant and Attack Discovery that push the platform closer to higher-value workflows across Search and AI, Observability and Security. The company has also made tangible progress on the GTM front in the past two years through sharper segmentation, tighter sales focus and a more structured approach to expansion across the installed base. That's showing up in steadier new logo activity, RPO growth and net expansion rates. However, we believe that the company still needs to prove it can turn broader product relevance and better GTM discipline into sustained growth acceleration, particularly as Search/AI becomes more crowded, observability gets pressured from both up-stack and lower-cost alternatives, and security grows more competitive as platforms converge. We therefore remain on the sidelines for now, looking for more consistent evidence that Elastic can translate its product momentum into sustained growth acceleration while navigating the competitive pressures of an increasingly crowded and converging market landscape.

Exhibit 36: Large Customer Net Adds
$100K+ new logo activity grew +5% on TTM basis, pointing to steady momentum in cross-sell and new customer wins

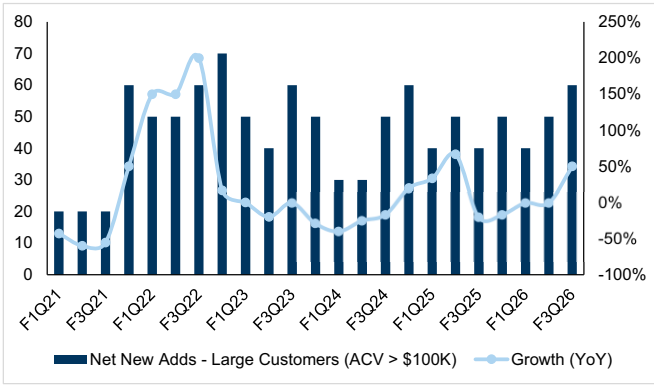

Source: Company Data

Exhibit 37: Constant Currency RPO Growth
Elastic's RPO growth has stabilized in the mid-to-high teens as GTM execution has improved

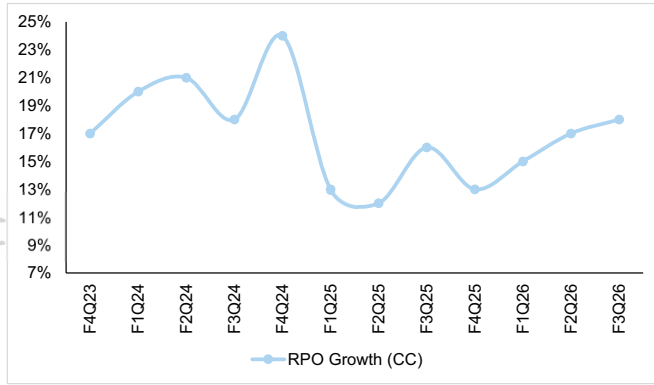

Source: Company Data

Elastic Through the Lens of the GS Framework for Investing in Software

Elastic screens as a strategically relevant but still-evolving platform story across the GS Framework for Investing in Software. The company spans search, observability and security through a common architecture built around Elasticsearch, one of the most widely adopted search and retrieval engines in enterprise software – a foundation that aligns well with several of the most important infrastructure themes today: unstructured data growth, AI-driven retrieval demand, open telemetry adoption and security-observability convergence. GTM execution has also improved over the last two years, with better segmentation, improving rep productivity, and more consistent deal activity reflecting. The medium-term framework targeting 20%+ Sales-Led Subscription growth, 20%+ non-GAAP OpM's (and 20%+ Adj. FCFM's) screens achievable, supported by GTM efficiency gains and G&A leverage as the business scales. At the same time, the competitive backdrop is intensifying across all three end markets, with broader platforms absorbing more of the workflow and data economics. Against that backdrop, Elastic Cloud growth continues to decelerate even as the AI narrative has broadened, and peers across the data and observability stack have shown cleaner demand reacceleration. We think Elastic has a credible platform and the right strategic direction, but we require clearer evidence that improving execution is translating into a more durable growth inflection before we can take a more constructive view.

订阅研报添加微信 LCD123321Cad 04-21

Exhibit 38: Mapping Elastic to the GS Framework for Investing in Software

	OUR VIEW ON ELASTIC	
8 KEY FACTORS	**TANGIBLE**	**INTANGIBLE**
TAM	TAM (Search, Observability, Security) (2026): $143B; TAM (Search, Observability, Security) (2029): $296B	Unstructured data, AI retrieval demand and converging observability-security workflows collectively expand TAM
SECULAR THEMES	Unstructured data growth, AI-driven retrieval demand, OpenTelemetry adoption, and security-observability convergence structurally	All three end-markets benefit from key secular tailwinds; competitive intensity rising in parallel across each
ENTRY/EXIT POINTS	Stock trades at 2x FY28 EV/Sales & 12x FY28 EV/FCF	Limited AI optionality priced in; cloud consumption inflection or accelerated AI monetization could be a meaningful catalyst
LT FRAMEWORK RETURN	Compounded Returns: Bear: -10%, Base: 4%, Bull: 15%	Rule-of-40+ target achievable but requires cloud-driven growth acceleration and margin expansion to underwrite LT compounding
UNIT ECONOMICS	Currently operating at Rule of 35% with Non-GAAP CAC at 2.6x with Software Median at 2-3x	Expansion within existing accounts, upmarket motion toward larger, multi-solution relationships driving higher ACV lands
PLATFORM/ BEST OF BREED	Unified architecture spanning enterprise search, log analytics, SIEM, and AI retrieval; extending into higher-value workflows	Architecture spans three large markets credibly; best-of-breed competitors remain more deeply embedded in respective workflows
COMPETITIVE MOAT	Deeply embedded in mission-critical workflows; OSS distribution advantage; OTel-first posture; early-mover in Vector Search	Open-source drives broad adoption but limits exclusivity; moat strongest in search, less established elsewhere
ESG	Hybrid and sovereign cloud deployment enables strict data residency; CISA award validates government positioning	Rising regulatory scrutiny around data governance, AI transparency, and cyber resilience aligns with platform strengths

Source: Company Data, Goldman Sachs Global Investment Research

In many respects, the debate on Elastic has not evolved considerably in recent years; the company is exposed to attractive markets, but competition is still the central risk. Elastic continues to span Search and AI, Observability, and Security with credible product breadth, but the market has long questioned whether that translates into best-of-breed status or durable control of any single category. What has changed more recently, in our view, is the intensity of competition, as larger platforms, adjacent vendors and emerging architectures push further into Elastic's markets and raise the bar around workflow ownership, data economics and strategic relevance. Against that backdrop, AI is increasingly the opportunity. If AI workloads continue to proliferate and Elastic can convert its search heritage, retrieval quality and broader platform investments into greater budget capture across the installed base, the business could move closer to its medium-term framework over time. **Our more measured view is that this upside remains plausible, but still unproven, while the rising competitive intensity across each of Elastic's end-markets remains the more immediate risk to sustained acceleration.**

We frame Elastic through the lens of its three core businesses: Search and AI, Observability and Security. In the sections that follow, we assess each business on its own merits, with a focus on where Elastic's product direction and strategic positioning are strongest, where the competitive backdrop is becoming more demanding, and how those dynamics shape our medium-term growth outlook. Taken together, these businesses explain both why Elastic remains strategically relevant and why we continue to see a more balanced risk-reward at current levels.

Search and AI

Search and AI remains Elastic's clearest right to win. This is where the company has the deepest heritage, the strongest technical credibility and the most natural role in the evolving AI stack. Search is Elastic's core DNA, and the company had vector capabilities in place as far back as 2017, well before ChatGPT reset market attention around AI.

Elastic has been investing in the stack for years and has continued to deepen it through hybrid retrieval, reranking, inference and a series of back-end optimizations designed to improve relevance, latency and cost at scale. In our view, that long development arc likely gives Elastic a stronger retrieval product that many alternatives, particularly for customers that care about production-grade search quality versus basic vector support alone. Elastic's recent launch of Agent Builder and Workflows extend that effort beyond retrieval alone by tying search and context directly into AI agents and automation, reinforcing the company's push to become a context engineering platform rather than another vector database.

Exhibit 39: Elastic- Key Products around Elasticsearch and Emerging AI Tools

Product / Feature	What It Does	Why It Matters
Elasticsearch (Core Search Engine)	Distributed, RESTful search and analytics engine built on Apache Lucene. Foundational technology underlying all three of ESTC's solution pillars.	ESTC's durable competitive moat. Enterprises that build on it become deeply embedded, it underpins everything from app search to RAG pipelines for generative AI.
Vector Search (kNN / Semantic Search)	Enables similarity-based retrieval using embedding vectors stored natively in Elasticsearch, supports k-Nearest Neighbor (kNN) search, hybrid vector search.	Vector search is the infrastructure layer for enterprise RAG, every company building an AI assistant on proprietary data needs this. Positions ESTC as default AI-native search infrastructure.
Elastic Inference Service (EIS)	Managed inference endpoints for embeddings/reranking and other ML models used in semantic search and RAG.	Reduces complexity and cost of running model infra; speeds time-to-value for AI search and agent workflows.
ELSER (Elastic Learned Sparse Encoder)	ESTC's proprietary sparse embedding model trained specifically for retrieval tasks, runs inside Elasticsearch with no external model hosting required.	Removes the need for a separate embedding service, reducing latency and cost for enterprises building AI search.
BBQ (Better Binary Quantization)	A compression technique for AI search data (vectors). Shrinks each data point by 32x while preserving the accuracy of search results.	Cuts memory needs by ~95%, making AI-powered search dramatically cheaper to run.
DiskBBQ (Disk-friendly ANN index)	An extension of BBQ that moves AI search data from expensive memory (RAM) to cheaper disk storage. Groups related data into clusters so the system only reads what it needs from disk.	Removes memory as a cost bottleneck for AI search at scale. Enables large-scale AI apps on lower-cost hardware.
Elastic Agent Builder	Elastic-native framework to build AI agents over Elasticsearch data, combining LLM reasoning with built-in/custom tools and integrations via MCP/A2A and APIs.	Extends Elastic from retrieval to agentic execution: answers stay grounded in indexed data with security controls, while agents can orchestrate multi-step tasks and trigger workflows.
Jina AI (Embeddings + Reranking via Inference API)	Uses Jina AI embedding and reranking models through Elasticsearch's Inference API / Elastic Inference Service, generating dense vectors for semantic/hybrid search.	Gives teams high-quality multilingual embeddings and reranking without self-hosting model infrastructure, improving relevance for vector search and production RAG pipelines.

sustainably closer to its mid-term framework.

Exhibit 40: Elastic has seen solid uptake of its Vector Search capabilities
While constructive on underlying adoption, it remains to be seen whether AI workloads can drive a sustained inflection in the growth profile

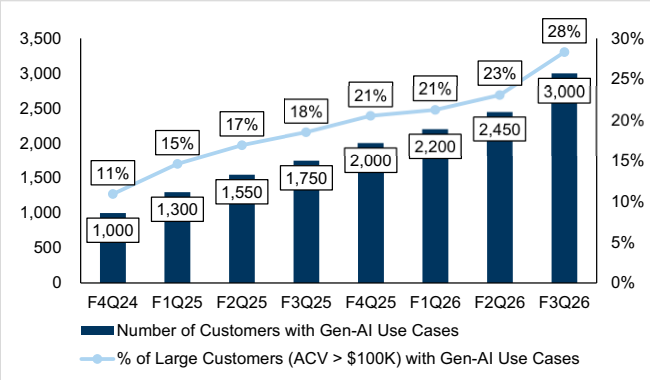

Source: Company Data

Exhibit 41: Gen-AI Customers are expanding 600 bps faster than non-AI customers
Early signals are positive around AI adoption and monetization but it is yet to surface materially in the overall growth profile

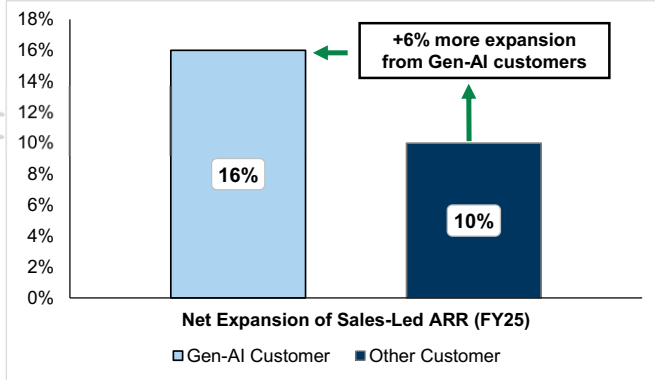

Source: Company Data

Exhibit 42: Elastic Vector Search leads in DB-Engines rankings among pure-play vector databases
We believe this speaks to Elastic's strong develop mindshare and feature depth

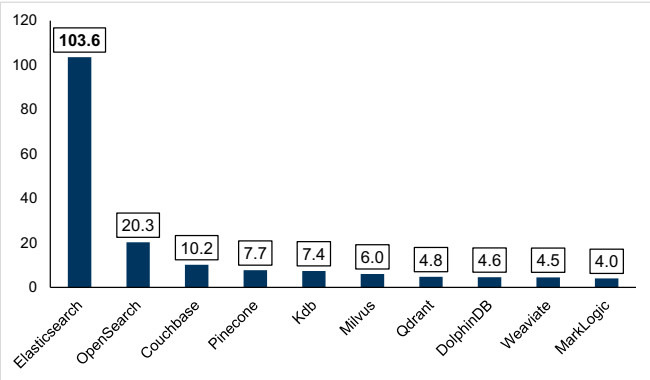

Source: DB-Engines

Exhibit 43: ...but competition is much more fierce when including broader data platforms
We continue to monitor whether Elastic's incumbent position in search can support growth acceleration as AI workloads scale

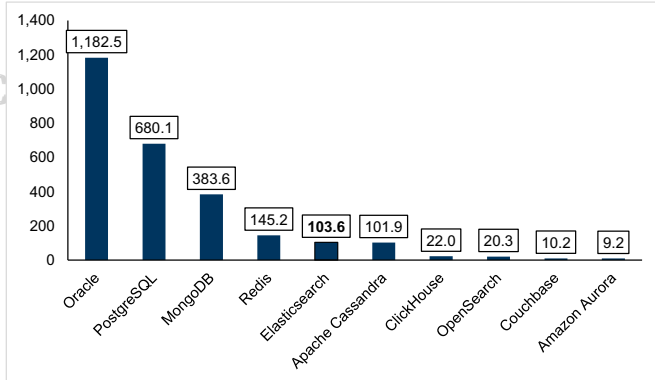

Source: DB-Engines

Observability

Observability fits Elastic's architecture, but the market is becoming more competitive from above and below. Elastic has built a scaled position in log analytics, and that footprint gives it a legitimate foundation in observability. Logs remain one of the richest and most widely used sources of operational data, particularly for investigation-heavy workflows, and Elastic's search expertise lends naturally to that use case. The company has leaned into that advantage from an OTel-first posture, ES|QL, AI Assistant and Streams, all aimed at making raw telemetry easier to investigate and faster to turn into root-cause analysis. The installed base reflects a logs-first entry point, with more than 90% of Elastic Cloud observability customers using log analytics and more than 35% already expanding beyond logs, which suggests Elastic's role in observability is established, though the path to broader platform ownership less so.

pushing the platform closer to a more automated SOC. Attack Discovery is worth highlighting - it uses AI to correlate disparate alerts into coherent attack narratives, an extension of Elastic's core strength in search and analytics across complex, multi-source data environments. Our industry conversations suggest that is where Elastic tends to find its strongest security traction - in specialized, data-intensive investigations that fall outside what more packaged solutions can readily handle. Security is one of the clearest places where ingest, search, correlation, and action fit naturally on a common architecture, and Elastic's SIEM footprint gives it a real foothold from which to expand. Elastic's two $20M+ security deals, including the CISA-related SIEM-as-a-Service award, are recent examples of that positioning showing up in practice, supporting the view that Elastic's security and data platform can resonate not just as a technical architecture, but also in large-scale environments.

Exhibit 47: Elastic - Key Products and Modules for Security Use Cases

Product / Feature	What It Does	Why It Matters	Competitors
SIEM	Aggregates security events from endpoints, network, cloud, and identity sources; correlates them against threat intelligence.	Elastic's search-first architecture enables faster queries across large securit datasets, a speed advantage over legacy SIEMs like IBM QRadar or Splunk ES.	Splunk, Microsoft Sentinel, CrowdStrike (Falcon Next-gen SIEM)
Endpoint Detection & Response (EDR / Elastic Defend)	Provides endpoint protection and detection using behavioral analytics. Telemetry feeds directly into the SIEM for correlated investigation.	Tight EDR-SIEM integration is a key differentiator, analysts pivot from an alert to raw endpoint telemetry in one click, eliminating tool-switching.	CrowdStrike (Falcon), SentinelOne, Palo Alto Networks (Cortex XDR)
Attack Discovery	Uses AI to automatically cluster and summarize related alerts into attack chains, presenting analysts with a narrative of an entire attack	Attack Discovery directly addresses alert fatigue, showcasing the 'AI for security' angle central to the security-observability convergence thesis.	CrowdStrike (Charlotte AI), SentinelOne (Purple AI)

Source: Company Data, Goldman Sachs Global Investment Research

The opportunity in security is real, even as the broader market is becoming more interconnected and competitive. The same convergence across telemetry, observability, security and automation that supports Elastic's platform thesis is also drawing larger vendors deeper into adjacent parts of the stack. CrowdStrike's Next-Gen SIEM is an illustrative example of the pressure Elastic is up against - the product exited FY26 at $585M in ARR growing more than 75%, underscoring both the scale of the modernization cycle and how quickly value is consolidating around broader security platforms. Palo Alto's Chronosphere acquisition, CrowdStrike's Onum deal and Databricks' recent Lakewatch launch all point in the same direction: the lines between telemetry, observability, security and workflow automation are blurring. We think that dynamic is strategically supportive of Elastic, but it also raises the bar. Security platforms would argue that their ownership of more first-party telemetry data across endpoint, identity, cloud and network gives them a natural advantage in consolidating adjacent workflows such as SIEM onto their platforms over time. Elastic remains relevant in that environment, but it will need to keep proving its strength in search, analytics and flexible data ingestion are differentiated enough to matter as security teams consolidate more of their workflow onto broader platforms. In our view, security is an important part of the thesis, but not yet a sufficient reason to underwrite a faster or more durable growth profile.

Exhibit 48: Gartner's Product or Service Score by Use Case (SIEM)
Similar to Observability, Elastic is relevant in SIEM but faces strong competition from a growing slate of security and data platform peers

Capability	Rank (n=17)	Score (Out of 5)
Out-of-the-Box SIEM	#9	3.36
Customizable SIEM	#7	3.61
Threat Detection (TDIR)	#7	3.63
Average	**#8**	**3.53**

Source: Gartner

Exhibit 49: Elastic's position as a top 10 vendor in SIEM underscores the credibility of its search technology for investigative workflows
Market share data, while dated, highlights a heavy-set competitive field with Cisco, Microsoft, Palo Alto and CrowdStrike

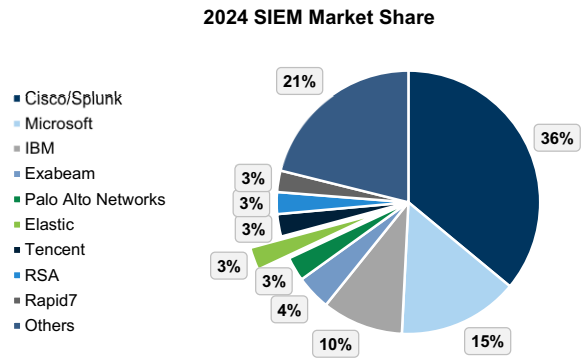

Source: Gartner

Exhibit 50: We use CrowdStrike's momentum in Next-Gen SIEM as a proxy for the heightened competition within the category
Larger platform peers like Palo Alto, Cisco (Splunk) and now Databricks are all vying for SIEM budgets

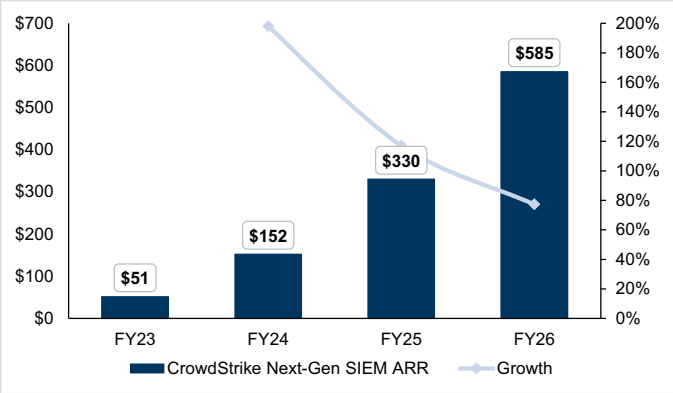

Data in MMs, except for percentages

Source: Company Data

The medium-term framework is credible, but the current model still looks more stable than inflecting. Elastic's mid-term targets call for 15%+ base Sales-Led Subscription (SLS) growth plus 5%+ from Gen-AI tailwinds, alongside 20%+ Non-GAAP OpM, 20%+ adjusted FCFM and a Rule of 40 above 40%. FY26 is not far off that path, with guidance implying roughly +18% CC Sales-Led growth. The nuance is that the stronger performance this year has come from a mix of cloud and self-managed, while Elastic Cloud – still the metric many investors gravitate toward as the cleanest read on AI momentum – has continued to decelerate from the mid-20's in FY25 to +21% (CC) in F3Q26. That leaves the model in a somewhat mixed place: execution is better, commitments are healthy, and the business looks more stable, but the cleanest signal of AI-driven acceleration has yet to improve. Margin expansion remains supportive, though likely not enough on its own to drive a more meaningful re-rating. For the stock to work from here, Elastic likely needs to pair continued operating leverage with clearer evidence

that growth in the cloud business can stabilize and improve.

Exhibit 51: Elastic's Mid-Term Framework
Elastic's framework calls for growth acceleration with AI serving as the incremental driver of growth; we see a business more stable than inflecting today

Metric	Mid-Term Target	Key Drivers
Sales-Led Subscription Revenue Growth	15% Base + 5% Gen-AI Tailwinds = 20% +	Existing customer expansion & Gen-AI driven growth acceleration across search, security and observability
Sales-Led Subscription Revenue % of Total	85%-90%	Majority of growth coming from sales-led subscription revenue
Non-GAAP Operating Margin	>20%	Increasing leverage as business scales
Adj. Free Cash Flow	>20%	Operating margin expansion Deferred revenue growth
Rule Of 40 (Total Revenue Growth + Adj. FCFM)	40% +	Balanced focus on both growth and profitability
Net Dilution	<2.5%	Disciplined headcount growth, balanced against investment in talent

Source: Company Data

Exhibit 52: Elastic Cloud vs. Self-Managed Revenue Growth
Self-Managed revenue growth is inflecting while Elastic Cloud continues to decelerate

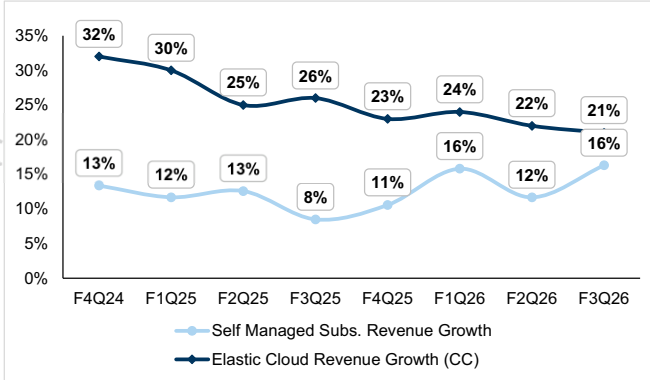

Source: Company Data

Digging into the underlying revenue trends adds further context. Elastic Cloud Net New Revenue (NNR) has been effectively range-bound for three consecutive quarters at roughly $37-$38M (in constant-currency) and broadly unchanged vs. prior year. Self-Managed, by contrast, has shown a more meaningful step-up with CC NNR reaching $23M in F3Q26, up +57%. That improvement reflects real strength landing in hybrid and regulated environments alongside some degree of AI-driven demand. Elastic runs the business on Sales-Led Subscription growth as its primary financial metric, and Self-Managed is a legitimate contributor to that motion. That said, the market has not meaningfully rewarded that outperformance, largely because cloud is still viewed as the more direct signal of AI-driven consumption demand – and on that metric the trend has yet to move.

Exhibit 53: Elastic Cloud Net New Revenue (YoY) in Constant Currency
Elastic Cloud NNR has been stable, driving growth deceleration

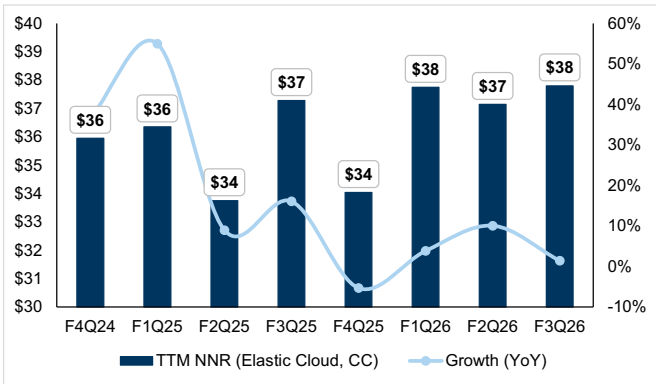

Data in MMs, except for percentages

Source: Company Data

Exhibit 54: Self-Managed Net New Revenue (YoY) in Constant Currency
Self-Managed accelerating; regulated and hybrid environments driving momentum

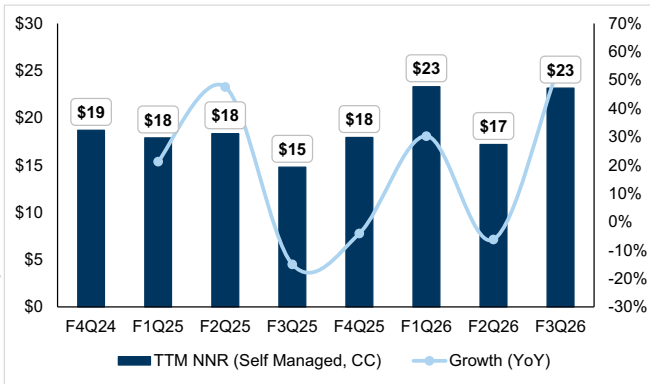

Data in MMs, except for percentages

Source: Company Data

Exhibit 62: Estimates vs. Consensus

All figures in $ mns	4Q26 (E) Guidance			4Q26 (E) GSe	4Q26 (E) Cons.	FY26 (E) Guidance			FY26 (E) GSe	FY26 (E) Cons.	FY27 (E) GSe	FY27 (E) Cons.	FY28 (E) GSe	FY28 (E) Cons.
Sales-led Subscription Revenue	$371	--	$373	$372	$373	$1,434	--	$1,436	$1,435	$1,436	$1,663	$1,664	$1,908	$1,928
YoY	*18%*	--	*19%*	*18%*	*18%*	*20%*	--	*20%*	*20%*	*20%*	*16%*	*16%*	*15%*	*16%*
QoQ	*-1%*	--	*-1%*	*-1%*	*-1%*									
Elastic Cloud (Annual)				$169	$169				$641	$640	$793	$791	$957	$965
YoY				*26%*	*26%*				*28%*	*28%*	*24%*	*24%*	*21%*	*22%*
QoQ				*0%*	*0%*									
Total Revenue	$445	--	$447	$446	$447	$1,734	--	$1,736	$1,734	$1,735	$1,969	$1,969	$2,221	$2,243
YoY	*15%*	--	*15%*	*15%*	*15%*	*17%*	--	*17%*	*17%*	*17%*	*14%*	*13%*	*13%*	*14%*
QoQ	*-1%*	--	*-1%*	*-1%*	*-1%*									
Non GAAP Gross Profit				$349	$347				$1,359	$1,356	$1,546	$1,536	$1,748	$1,748
Gross Margin				*78%*	*78%*				*78%*	*78%*	*79%*	*78%*	*79%*	*78%*
Non GAAP Operating Income	$65	--	$65	$65	$65	$283	--	$283	$283	$283	$342	$341	$407	$415
Operating Margin	*14.5%*	--	*14.5%*	*15%*	*15%*	*16.3%*	--	*16.3%*	*16%*	*16%*	*17%*	*17%*	*18%*	*19%*
Non GAAP EPS	$0.55	--	$0.57	$0.56	$0.56	$2.50	--	$2.54	$2.53	$2.52	$2.87	$2.85	$3.28	$3.32
GAAP EPS				($0.12)	--				($0.77)	--	$0.22	--	$0.67	--
Total Billings (Balance Sheet)				$616	$611				$1,861	$1,852	$2,111	$2,095	$2,378	$2,357
YoY				*15%*	*16%*				*13%*	*13%*	*13%*	*13%*	*13%*	*13%*
QoQ				*17%*	*17%*									
Total Deferred Revenue				$979	$977				$979	$977	$1,121	$1,112	$1,278	$1,263
YoY				*15%*	*15%*				*15%*	*15%*	*15%*	*14%*	*14%*	*14%*
QoQ				*21%*	*21%*									
Cash Flow from Operations				$142	$138				$316	$312	$365	$380	$428	$460
OCF Margin				32%	*31%*				18%	*18%*	19%	*19%*	19%	*21%*
Unlevered Free Cash Flow				$141	$141				$337	$336	$385	$397	$447	$474
FCF Margin				*32%*	*31%*				*19%*	*19%*	*20%*	*20%*	*20%*	*21%*

All Revenue Metrics- Visible Alpha Consensus Data, Rest- FactSet

Source: Company data, Visible Alpha Consensus Data, FactSet, Goldman Sachs Global Investment Research

M&A Framework

Across our global coverage, we examine stocks using an M&A framework, considering both qualitative factors and quantitative factors (which may vary across sectors and regions) to incorporate the potential that certain companies could be acquired. We then assign an M&A rank as a means of scoring companies under our rated coverage from 1 to 3, with 1 representing high (30%-50%) probability of the company becoming an acquisition target, 2 representing medium (15%-30%) probability and 3 representing low (0%-15%) probability. For companies ranked 1 or 2, in line with our standard departmental guidelines, we incorporate an M&A component into our target price. An M&A rank of 3 is considered immaterial and therefore does not factor into our price target and may or may not be discussed in research.

Within this context, we have ranked Elastic a 2 and weight our M&A value 15% (reflecting the probability implied by this ranking) within our price target methodology. We view Elastic as a strategically relevant asset within infrastructure software, where its core search and unstructured data technology could be complementary for a range of potential acquirers. In order to derive the M&A value, we apply a 3.5x NTM EV/S multiple to our Q5-Q8 revenue estimate, a one-and-a-half turn discount to the median of precedent transactions which is in-line with some lower multiple transactions in the space and indicates a ~35% premium to the current multiple at which Elastic trades. Our M&A multiple yields a $70 theoretical value.

Exhibit 63: Precedent M&A Transactions

Target	Acquirer	Closing Date	Deal Value ($M)	NTM Revenue ($M)	NTM FCF ($M)	EV/S (NTM)	EV/FCF (NTM)
Red Hat	IBM	07/09/19	$32,586	$3,703	$1,073	8.8x	30.4x
Sumo Logic	Francisco Partners	05/12/23	$1,195	$351	-$18	3.4x	--
New Relic	Francisco Partners / TPG	11/08/23	$5,767	$1,068	$152	5.4x	38.0x
VMWare	Broadcom	11/22/23	$68,978	$14,077	$3,892	4.9x	17.7x
Alteryx	Clearlake Capital Group	12/18/23	$4,400	$1,025	$50	4.3x	87.6x
Everbridge	Clearlake Capital Group	02/05/24	$1,500	$459	$75	3.3x	20.0x
Splunk	Cisco	03/18/24	$28,083	$4,255	$1,017	6.6x	27.6x
Hashicorp	IBM	02/27/25	$5,891	$669	$112	8.8x	52.7x
SolarWinds	Turn/River Capital	04/16/25	$4,200	$823	$259	5.1x	16.2x
Informatica	Salesforce	11/18/25	$8,314	$1,732	$496	4.8x	16.8x
Confluent	IBM	03/17/26	$10,245	$1,348	$110	7.6x	93.4x
Average						**5.7x**	**40.0x**
Median						**5.1x**	**29.0x**

Source: FactSet, Company Data

Valuation and Key Risks

We initiate on Elastic with a 12-month Price Target of $50. Our valuation methodology is based on an 85% weighting to EV/FCF and a 15% weighting to M&A. We apply a 12x EV/FCF multiple to our Q5-Q8 FCF estimate to arrive at our fundamental value. Our EV/FCF multiple is informed by a blend of both infrastructure software and comparable Rule-Of software peers. Our EV/FCF multiple represents a discount to the median of infrastructure software peers to account for Elastic's lower Rule-Of profile. We apply a 3.5x NTM EV/S multiple to our Q5-Q8 revenue estimates to derive an implied M&A valuation, which is approximately 1.5 turns below the median of precedent transactions while still implying a ~35% premium to the current EV/S multiple.

Key upside risks: 1) AI-Driven Cloud Acceleration: A more pronounced inflection in Elastic Cloud growth driven by AI and Vector Search workloads would be the most significant upside catalyst, potentially pulling forward the medium-term framework, **2) GTM Productivity:** If FY26 capacity investments compound faster than anticipated, S&M efficiency could improve more meaningfully, driving operating leverage above the current framework, **3) Security Momentum:** A more meaningful ramp in SIEM displacement and XDR adoption could expand Elastic's share of enterprise security budgets beyond current expectations.

Key Downside Risks: 1) Cloud Growth Deceleration: Continued cloud growth deceleration would undermine the AI narrative and raise questions about budget capture across all three end markets, **2) Competitive Displacement:** Intensifying competition across search, observability and security from broader platforms absorbing more of the workflow and data economics layer could compress win rates and limit installed base expansion, **3) GTM Execution:** If FY26 investments fail to translate into improved growth, S&M efficiency would remain elevated and the margin expansion story would be diluted, **4) Macroeconomic Sensitivity:** A deterioration in IT spending could disproportionately impact discretionary search and observability workloads, particularly as customers face pressure to rationalize telemetry spend.

订阅研报请加微信 LCD123321Cad 04-21