AI Series #35: Price Hike Continues; A Month of Multiple Models Come to Market

April 20, 2026

Highlights (1) ATH and Tencent released world models last week which is positive for film, video and gaming sectors; (2) Happy Horse is expected to start testing among enterprise users on 27 Apr. Official release is targeted in May-26; (3) BABA's Bailian platform raises price on 15 May; (4) Baidu Coding Plan resumes normal pricing after discounts ended on 18 Apr; (5) market pays attention to latest updates of DeepSeek V4 (e.g. launch date, potential financing).

订阅研报添加微信 LCD123321Cad

In this report, we provide investors with over 50 data points on AI models, OpenClaw, user trends by sub-sector, and other industry data for analysis.

World models catch attention; Happy Horse-1.0 to start testing on 27 Apr. BABA and Tencent released two models last week, which are important for the film, video, and gaming sectors. For BABA, ATH released Happy Oyster on 16 Apr which is multimodal and empowers users to create a 3D digital world. Happy Oyster allows users to navigate and explore AI-generated environments or adjust narratives and camera angles. Tencent launched HY-World 2.0 the same day, which enables users to create industrial-grade and editable 3D assets with a 94% consistency rate. Separately, Happy Horse-1.0 (video generation model) will be available for limited testing among enterprises in Bailian on 27 Apr. Official launch of Happy Horse is targeted in May-26.

BABA's Bailian platform to raise prices on 15 May. AliCloud's model aggregator platform Bailian, which has over 200 models, highlights it will raise prices by 2-7% for different MU (model units) based on fixed monthly packages and API calls. The rationale for the price hike of the Coding Plan is mainly due to a surge in demand amid supply constraints. In our view, there will be no impact on the agreed API price paid by Bailian to AI Labs, with the former acting as a distributor.

Standford: The gap between US-China models is close. According to 2026 Stanford AI Index Report (Apr-26), the top US model leads Chinese models by 2.7%. Separately, Anthropic (1503), xAI (1495), Google (1494), OpenAI (1481), Alibaba (1449) and DeepSeek (1424) occupy the top tier of Arena Elo ratings as of Mar-26.

OpenRouter: Token consumption was down 2% vs prior week. Based on the analysis of the week of 13 Apr (vs the week of 6 Apr) (Exhibit 1), the number of weekly tokens consumed decreased by 2% to 20.6tn. We believe this is due to (1) OpenRouter promotional campaigns for Mimo-V2-Pro model, which ended on 2 Apr; (2) OpenRouter promotional campaigns for Qwen3.6 Plus (free) between 2 and 16 Apr. There are limitations set on free API calls after 6 Apr. In terms of ranking, Anthropic's Claude Sonnet 4.6 ranked no. 1, followed by DeepSeek V3.2 and Claude Opus 4.6

订阅研报添加微信 LCD123321Cad

BAT, KC, and AI Labs are Beneficiaries. We expect a surge in token consumption on rising demand for AI models and agents to benefit CSPs like **Baidu, Alibaba, Tencent (BAT), Kingsoft Cloud (KC), and AI Labs.** According to analysis from Artificial Analysis, (1) there is a narrowing intelligence gap between US and China models; (2) API cost is just a fraction vs US benefiting from MoE, linear attention, and MFU (Exhibit 22-28); (3) Coding, Workspace scenario and mid- to long-form content have significant opportunities ahead (Exhibit 55-59). According to media reports (e.g. TechWeb), the industry expects DeepSeek V4 to be released in Apr.

Continued overleaf...

Exhibit 1 - Weekly token consumption over the past 12 months as of 20 Apr (tn)

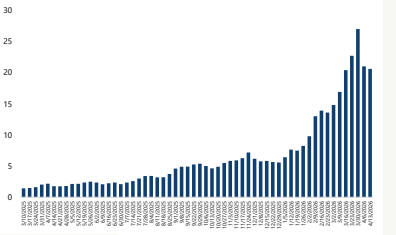

Source: OpenRouter, Jefferies

Exhibit 2 - Top 10 ranking in terms of token consumption by model in last week

No	Company	Model	Tokens
1	Anthropic	Claude Sonnet 4.6	1.38T
2	DeepSeek	DeepSeek V3.2	1.28T
3	Anthropic	Claude Opus 4.6	1.22T
4	Xiaomi	MiMo-V2-Pro	1.15T
5	Google	Gemini 3 Flash Preview	1.14T
6	MiniMax	MiniMax M2.5	1.05T
7	MiniMax	MiniMax M2.7	961B
8	Google	Gemini 2.5 Flash Lite	595B
9	xAI	Grok 4.1 Fast	573B
10	Google	Gemini 2.5 Flash	567B

Source: OpenRouter, Jefferies

China vs US comparison

Exhibit 3 - Alibaba Cloud revenue growth vs US hyperscalers between Sep-24 and Mar-26

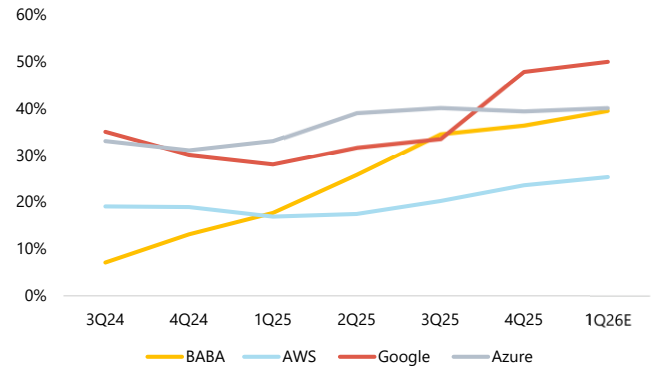

Source: Visible Alpha, Jefferies

Exhibit 4 - Alibaba Cloud revenue growth vs aggregate growth rate for US hyperscalers between Sep-24 and Mar-26

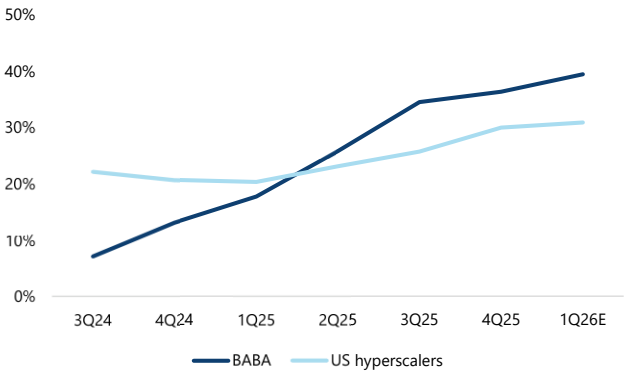

Source: Visible Alpha, Jefferies
US hyperscalers: Amazon, Microsoft, Google

Exhibit 5 - Capex as % of total revenue

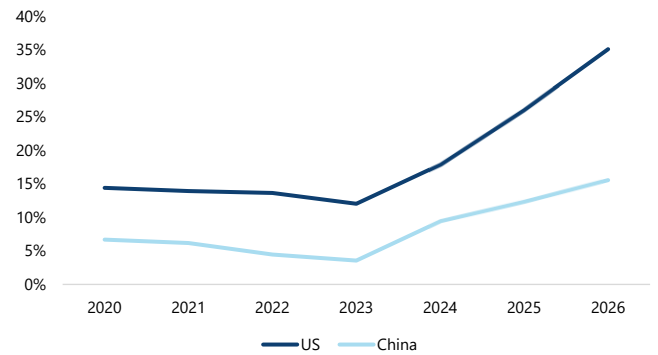

Source: Visible Alpha, Jefferies
China : Baidu, Alibaba, Tencent

US: Amazon, Microsoft, Google, Meta

Exhibit 6 - Capex as % of OCF

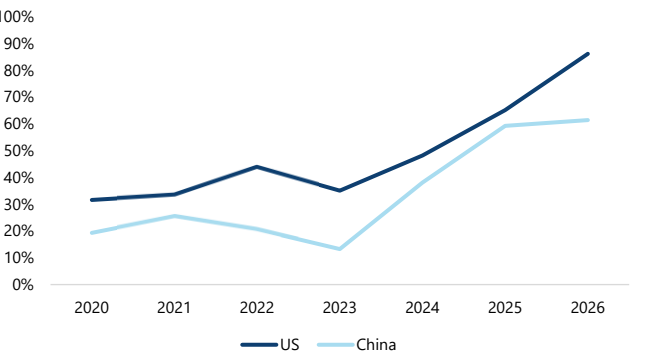

Source: Visible Alpha, Jefferies

订阅研报添加微信 LCD123321Cad 04-20

Input/Output Price for Different Models / Pricing for Coding Plan and Agents Offerings

Red color denotes price hike during the month

Exhibit 7 - Pricing plans for Alibaba, Tencent, Baidu and ByteDance

Models Provider	Model	Context	Apr-26	
			Input (RMB)	Output (RMB)
			(million tokens)	(million tokens)
Alibaba	Qwen3.6-Plus	(0, 256k]	2	12
		(256k, 1m]	8	48
	Qwen3.5-Plus	(0, 128k]	0.8	4.8
		(128k, 256k]	2	12
		(256k, 1m]	4	24
Tencent	HY 2.0 Think	(0, 32k]	3.975	15.9
		(32k,128k]	5.3	21.2
	HY 2.0 Instruct	(0, 32k]	3.18	7.95
		(32k,128k]	4.505	11.13
	Hunyuan-T1		1	4
	Hunyuan-TurboS		0.8	2
	Tencent HY Vision 1.5 Instruct		3	9
	Hunyuan-translation		1.2	3.6
Baidu	ERNIE-5.0	(0, 32k]	6	24
		(32k, 128k]	10	40
	ERNIE-4.5-Turbo		0.8	3.2
	ERNIE-4.5-Turbo-VL		3	9
ByteDance	Doubao-seed-2.0-pro	[0, 32k]	3.2	16
		(32k, 128k]	4.8	24
		(128k, 256k]	9.6	48
	Doubao-seed-2.0-lite	[0, 32k]	0.6	3.6
		(32k, 128k]	0.9	5.4
		(128k, 256k]	1.8	10.8
	Doubao-seed-2.0-mini	[0, 32k]	0.2	2
		(32k, 128k]	0.4	4
		(128k, 256k]	0.8	8
	Doubao-seed-2.0-code	[0, 32k]	3.2	16
		(32k, 128k]	4.8	24
		(128k, 256k]	9.6	48

Source: Company, Jefferies

LCD123321Cad
LCD123321Cad
订阅研报添加微信 LCD123321Cad 04-20

Exhibit 8 - Pricing plans for AI Labs MiniMax, Zhipu, Kimi, StepFun, Xiaomi

Models Provider	Model	Context	Apr-26	
			Input (RMB)	Output (RMB)
			(million tokens)	(million tokens)
MiniMax	MiniMax-M2.7		2.1	8.4
	MiniMax-M2.7-highspeed		4.2	16.8
	MiniMax-M2.5		2.1	8.4
	MiniMax-M2.5-highspeed		4.2	16.8
Zhipu AI	GLM-5.1	[0, 32k)	6	24
		[32k+)	8	28
	GLM-5-Turbo	[0, 32k)	5	22
		[32k+)	7	26
	GLM-5	[0, 32k)	4	18
		[32k+)	6	22
Kimi	Kimi K2.5		0.7	21
	kimi-k2-thinking		1	16
	kimi-k2-thinking-turbo kimi-k2-turbo-preview		1	58
	kimi-k2-0905-preview kimi-k2-0711-preview		1	16
	moonshot-v1-8k moonshot-v1-8k-vision-preview		2	10
	moonshot-v1-32k moonshot-v1-32k-vision-preview		5	20
	moonshot-v1-128k moonshot-v1-128k-vision-preview		10	30
StepFun	step-3.5-flash		0.14	2.1
	step-r1-v-mini		0.5	8
	step-3	input≤4k & output≤4k	0.3	4
		input≤4k & output>4k	0.3	8
		4k < input ≤64k	0.8	10
	step-2-mini		0.2	2
	step-2-16k step-2-16k-exp		7.6	120
	step-1-8k		1	20
	step-1-32k		3	70
	step-1-256k		19	300
Xiaomi	MiMo-V2-Pro	(0, 256k]	7	21
		(256k, 1m]	14	42
	MiMO-V2-Omni		2.8	14
	MiMo-V2-TTS		limited-time free	
	MiMo-V2-Flash		0.7	2.1

Source: Company, Jefferies

Exhibit 9 - Pricing plans for overseas players Anthropic, OpenAI and Google

LLM Provider	Model	Context	Apr-26	
			Input (US$)	Output (US$)
			(million tokens)	(million tokens)
Anthropic	Opus 4.7		5	25
	Sonnet 4.6		3	15
	Haiku 4.5		1	5
OpenAI	GPT-5.4		2.5	15
	GPT-5.4-mini		0.75	4.5
	GPT-5.4-nano		0.2	1.25
	GPT-5.4-pro		30	180
	GPT-5.2		1.75	14
	GPT-5.2-pro		21	168
	GPT-5.1		1.25	10
	GPT-5		1.25	10
Google	Gemini 3.1 Pro Preview	(0, 200k]	2	12
		(200k+)	4	18
	Gemini 3.1 Flash-Lite Preview	text/image/video	0.25	1.5
		audio	0.5	1.5
	Gemini 3 Flash Preview	text/image/video	0.5	3
		audio	1	3

Source: Company, Jefferies

LCD123321Cad
订阅研报添加微信 LCD123321Cad 04-20
订阅研报添加微信 LCD123321Cad 04-20

Coding Plans and Agents offerings

Red color denotes price hike during the month

Pricing Plans for different Coding and Agents offerings by leading players

Exhibit 10 - Tencent AI products pricing

Product	Version	Pricing
QClaw	Internal testing	TBD
WorkBuddy	Individual Trial	Limited-time free, 500 credits per month
	Individual Pro	RMB58 per month, 2,000 credits per month (incl. 500 credits for trial)
	Enterprise Flagship	RMB78 per month, 2,000 credits per month
	Enterprise Pro	RMB158 per month, 2,000 credits per month
Lighthouse	Starter	RMB45 per month, 2 vCPUs, 2GB MEM, 40GB SSD, 3Mbps Bandwidth, 200GB monthly data package
	Starter (Basic)	RMB50 per month, 2 vCPUs, 2GB MEM, 50GB SSD, 4Mbps Bandwidth, 300GB monthly data package
	General (Advanced)	RMB100 per month, 2 vCPUs, 4GB MEM, 100GB SSD, 7Mbps Bandwidth, 1,000GB monthly data package
	General (Pro)	RMB245 per month, 4 vCPUs, 8GB MEM, 180GB SSD, 12Mbps Bandwidth, 2,000GB monthly data package
	Razor Speed	RMB45 per month, 2vCPUs, 2GB MEM, 40GB SSD, 200Mbps Bandwidth, unlimited data package
	Razor Speed	RMB230 per month, 4 vCPUs, 8GB MEM, 120GB SSD, 200Mbps Bandwidth, unlimited data package

Source: Company, Jefferies

Exhibit 11 - BABA coding plan

Program	Price	Usage limitations
Coding Plan Pro	RMB 200/month	Every 5 hours: Maximum 6000 requests Every week: Maximum 45000 requests Every subscription month: Maximum 90000 requests

Source: Company, Jefferies

Exhibit 12 - Baidu coding plan

Program	Price	Usage limitations
Coding Plan Lite	RMB 40/month	Every 5 hours: Maximum 1200 requests Every week: Maximum 9000 requests Every subscription month: Maximum 18000 requests
Coding Plan Pro	RMB 200/month	Every 5 hours: Maximum 6000 requests Every week: Maximum 45000 requests Every subscription month: Maximum 90000 requests

Source: Company, Jefferies

Exhibit 13 - MaxClaw subscription prices

MaxClaw subscription prices					
Region	Program	Free	Basic	Pro	Ultra
China	Standard price	RMB0	RMB133 per month	RMB483 per month	
	Monthly program	RMB0	RMB39 per month	RMB119 per month	
	Privileges	200 credits per day 1,000 free credits (available for 3 days)	5,000 credits per month About 45 Pro mode tasks available 200 daily credits Peak hour priority Remove watermark Custom domain Deploy 24/7 Cloud MaxClaw	16,000 credits per month About 150 Pro mode tasks available 200 daily credits Peak hour priority Remove watermark Custom domain Deploy 24/7 Cloud MaxClaw	N/A
Overseas	Standard price	USD0	USD39 per month	USD119 per month	USD219 per month
	Monthly program	USD0	USD19 per month	USD69 per month	USD129 per month
	Yearly program	USD0	USD16 per month	USD58 per month	USD109 per month
	Privileges	200 credits per day 1,000 free credits (available for 3 days)	5,000 credits per month About 30 Pro mode tasks available 200 daily credits Peak hour priority Remove watermark Custom domain Deploy 24/7 Cloud MaxClaw	20,000 credits per month About 120 Pro mode tasks available 200 daily credits Peak hour priority Remove watermark Custom domain Deploy 24/7 Cloud MaxClaw	40,000 credits per month About 240 Pro mode tasks available 200 daily credits Peak hour priority Remove watermark Custom domain Deploy 24/7 Cloud MaxClaw

Source: Company, Jefferies

Exhibit 14 - MaxClaw Add-on prices

MaxClaw add-on prices			
Region	Starter pack	Booster pack	Mega pack
China	RMB59 for 5,000 credits (valid for 1 year)	RMB109 for 10,000 credits (valid for 1 year)	RMB149 for 15,000 credits (valid for 1 year)
Overseas	USD39 for 5,000 credits (valid for 1 year)	USD78 for 10,000 credits (valid for 1 year)	USD156 for 20,000 credits (valid for 1 year)

Source: Company, Jefferies

Exhibit 15 - Xiaomi MiMo token plan

Region	Program	Price	Content
China	Lite	RMB 39/month, RMB 34.32/month for the first purchase	60 million credits
	Standard	RMB 99/month, RMB 87.12/month for the first purchase	200 million credits
	Program	RMB 329/month, RMB 289.52/month for the first purchase	700 million credits
	Max	RMB 659/month, RMB 579.92/month, for the first purchase	1600 million credits
Overseas	Lite	USD 6/month, USD 5.28/month for the first purchase	60 million credits
	Standard	USD 16/month, USD 14.08/month for the first purchase	200 million credits
	Program	USD 50/month, USD 44/month for the first purchase	700 million credits
	Max	USD 100/month, USD 88/month for the first purchase	1600 million credits

订阅研报添加微信 LCD123321Cad 04-20

Exhibit 16 - GLM coding plan

Region	Program	Pricing	Content
China	Lite	RMB 49/month	5-hour limit: up to approx. 80 prompts, weekly limit: up to approx. 400 prompts
		RMB 132.3/quarter	
		RMB 470.4/year	
	Pro	RMB 149/month	5-hour limit: up to approx. 400 prompts, weekly limit: up to approx. 2000 prompts
		RMB 402.3/quarter	
		RMB 1430.4/year	
	Max	RMB 469/month	5-hour limit: up to approx. 1600 prompts, weekly limit: up to approx. 8000 prompts
		RMB 1266.3/quarter	
		RMB 4502.4/year	
Overseas	Lite	USD 18/month	5-hour limit: up to approx. 80 prompts, weekly limit: up to approx. 400 prompts
		USD 48.6/quarter	
		USD 172.8/year	
	Pro	USD 72/month	5-hour limit: up to approx. 400 prompts, weekly limit: up to approx. 2000 prompts
		USD 194.4/quarter	
		USD 691.2/year	
	Max	USD 160/month	5-hour limit: up to approx. 1600 prompts, weekly limit: up to approx. 8000 prompts
		USD 432/quarter	
		USD 1536/year	

Source: Company, Jefferies

04-20
订阅研报添加微信 LCD123321Cad 04-20
订阅研报添加微信 LCD123321Cad 04-20

AI Models

According to OpenRouter, token consumption went up significantly in Feb-26 (a) **surge in adoption of open-source OpenClaw**, which allows people using AI from "to chat" to "to work". It connects with large models which can be deployed locally or on servers which allows you to automate workflow through natural language chat; (b) **release a number of high performance and cost-efficient model** which includes Kimi K2.5 (Jan-26), M2.5 (Feb-26) and GLM-5 (Feb-26).

According to OpenRouter, rising adoption of OpenClaw and release of high-performance and cost-efficient China models drive token consumption

Exhibit 17 - Weekly token consumption over the past 12 months as of 20 Apr (tn)

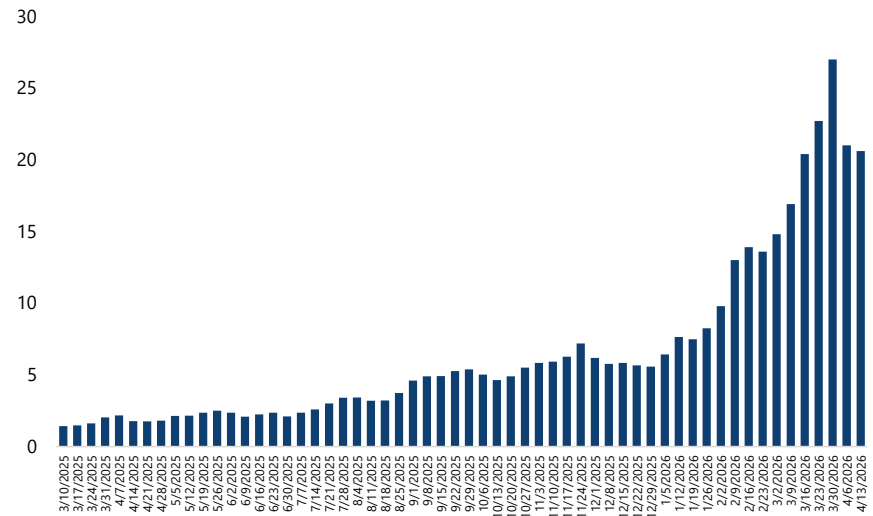

Source: OpenRouter, Jefferies

According to OpenRouter (between 13 Apr and 19 Apr), China models token consumption decreased by 23.8% compared to prior week reaching 4.4tn. We view this as due to (1) OpenRouter has promotional campaigns for Mimo-V2-Pro model, which ended on 2 Apr; (2) OpenRouter has promotional campaigns for Qwen3.6 Plus (free) between 2 and 16 Apr. There are limitations set on free API calls after 6 Apr. This compares to US models at 4.1tn token consumption.

China models surpass US models again in terms of token consumption the week of 13 April

Exhibit 18 - Market share: weekly token consumption by company as of 20 Apr

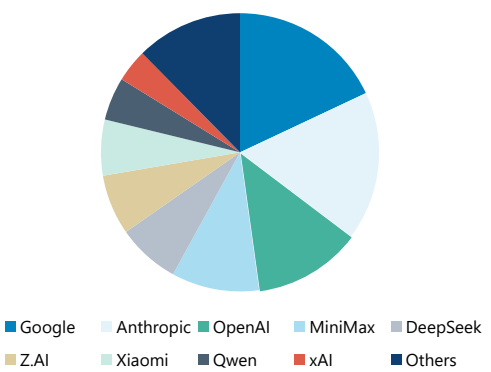

Exhibit 19 - Weekly token consumption since Jan-26 (bn)

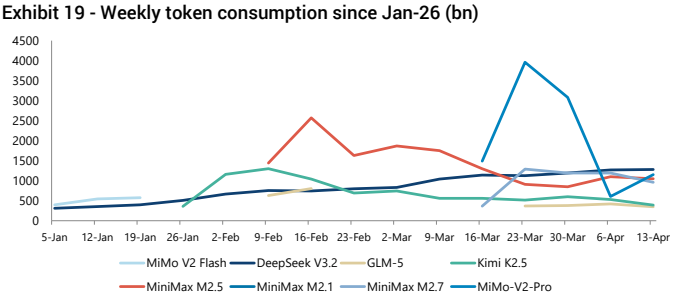

Source: OpenRouter, Jefferies

Source: OpenRouter, Jefferies

On token consumptions by company, Google (18%) comes first, followed by Anthropic (17.3%), OpenAI (12.5%), MiniMax (10.2%) and DeepSeek (7.3%). The top 5 weekly token consumption by model include Claude Sonnet 4.6 (1.38tn), followed by DeepSeek V3.2 (1.28tn) and Claude Opus 4.6 (1.22tn).

Claude Sonnet 4.6 ranked first in terms of token consumption the week of 13 April

Exhibit 20 - Top 10 ranking in terms of token consumption by model in last week

No	Company	Model	Tokens
1	Anthropic	Claude Sonnet 4.6	1.38T
2	DeepSeek	DeepSeek V3.2	1.28T
3	Anthropic	Claude Opus 4.6	1.22T
4	Xiaomi	MiMo-V2-Pro	1.15T
5	Google	Gemini 3 Flash Preview	1.14T
6	MiniMax	MiniMax M2.5	1.05T
7	MiniMax	MiniMax M2.7	961B
8	Google	Gemini 2.5 Flash Lite	595B
9	xAI	Grok 4.1 Fast	573B
10	Google	Gemini 2.5 Flash	567B

Source: OpenRouter, Jefferies

Exhibit 21 - No 11-20 ranking in terms of token consumption by model in latest week

No	Company	Model	Tokens
11	OpenAI	GPT-5.4	564B
12	OpenRouter	Elephant	564B
13	NVIDIA	Nemotron 3 Super (free)	535B
14	Z.AI	GLM 5.1	481B
15	OpenAI	gpt-oss-120b	425B
16	MoonshotAI	Kimi K2.5	391B
17	Google	Gemini 3.1 Pro Preview	387B
18	StepFun	Step3.5 Flash	364B
19	Z.AI	GLM 5	349B
20	Google	Gemini 3.1 Flash Lite Preview	305B

Source: OpenRouter, Jefferies

Based on Artificial Analysis, China models DeepSeek, Qwen, Kimi, MiniMax and Zhipu are closing the gap with US models in terms of intelligence. China models are highly cost-efficient on Mixture of Experts (MoE) architecture, linear attention, and MLU (Model Flog Utilization). According to 2026 Stanford AI Index Report (Apr-26), the top US model leads Chinese models by 2.7%. Separately, Anthropic (1503), xAI (1495), Google (1494), OpenAI (1481), Alibaba (1449) and DeepSeek (1424) occupy the top tier of Arena Elo ratings as of Mar-26.

China models are narrowing the gap with US models in terms of intelligence. On Input/ Output API, China large models are highly efficient vs US models.

Exhibit 22 - Frontier Language Model Intelligence over time as of 20 Apr

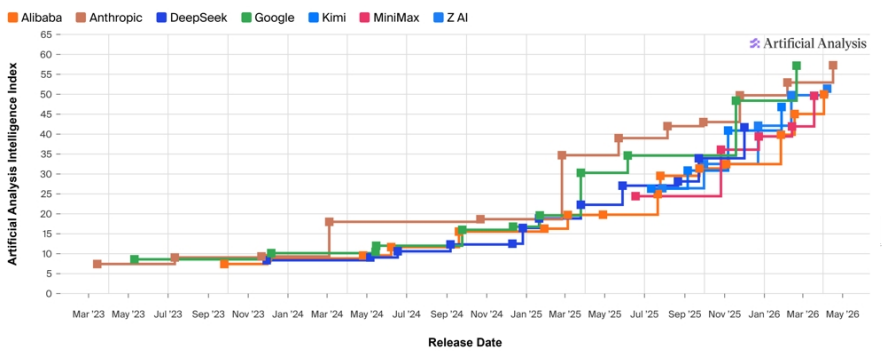

Source: Artificial Analysis, Jefferies

In terms of Input/Output API price, China models are just a fraction vs US models.

Exhibit 23 - Intelligence index for top 20 models as of 20 Apr

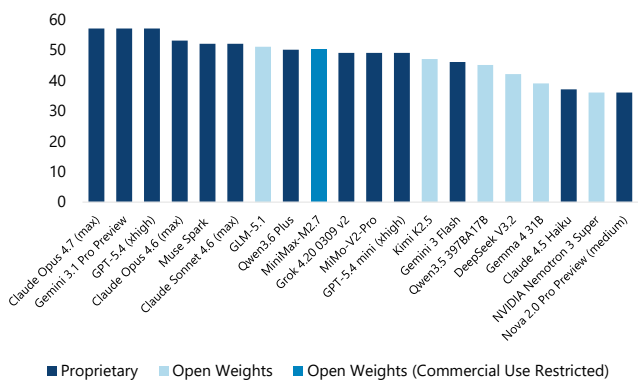

Exhibit 24 - Input/Output API price for leading models as of 20 Apr

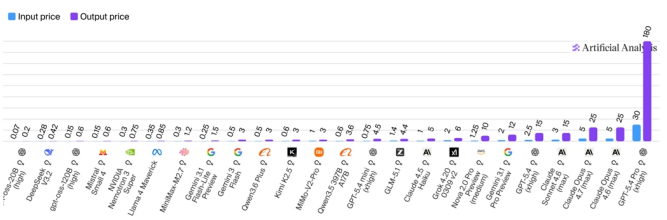

Source: Artificial Analysis, Jefferies

Source: Artificial Analysis, Jefferies

Among top 20 large models, there are 7 which are China models and open weighted as of 20 Apr.

Exhibit 25 - Agentic index for top 20 models as of 20 Apr

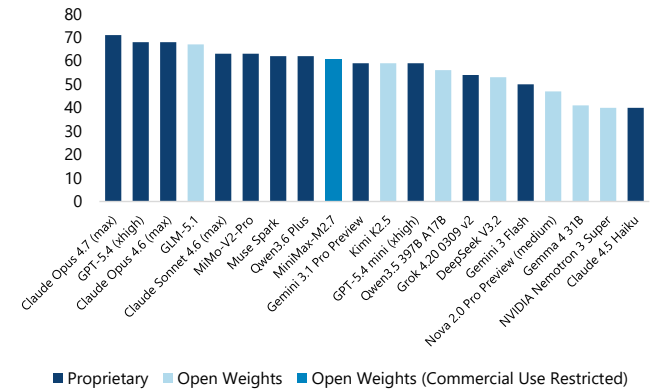

Source: Artificial Analysis, Jefferies

Exhibit 26 - Coding index for top 20 models as of 20 Apr

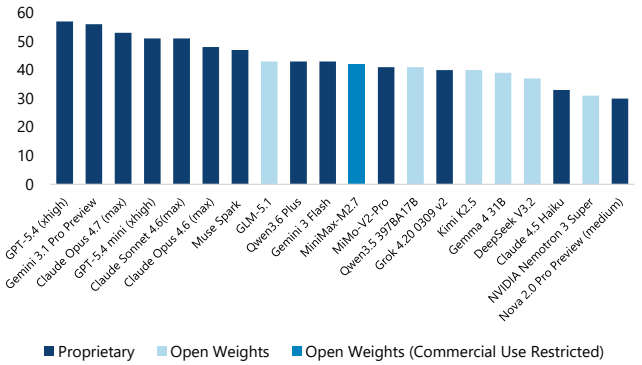

Source: Artificial Analysis, Jefferies

Exhibit 27 - Cost to run Artificial Intelligence Index as of 20 Apr

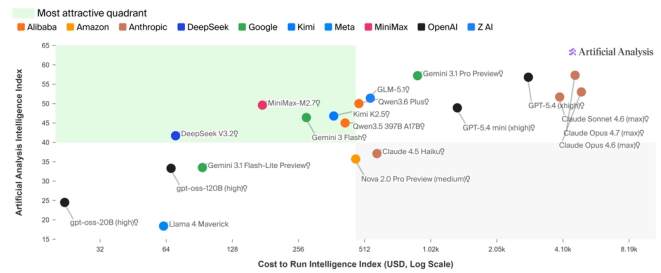

Source: Artificial Analysis, Jefferies

Exhibit 28 - Output speed to run Artificial Intelligence Index as of 20 Apr

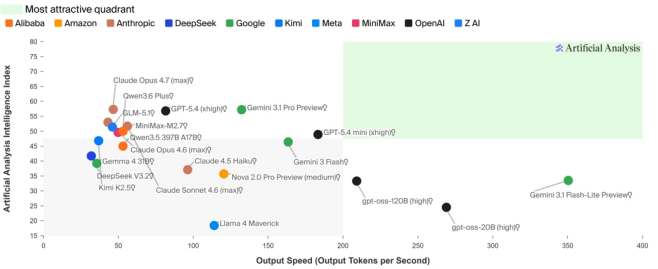

Source: Artificial Analysis, Jefferies

According to OpenRouter, top 3 models for OpenClaw based on token consumption come from China which include MiMo-V2-Pro, GLM 5 Turbo and Step 3.5 Flash. Other models include MiniMax M2.7, Qwen3.6 Plus, MiniMax M2.5, Kimi K2.5, DeepSeek V3.2, and MiMo-V2-Omni.

According to OpenRouter, MiMo-V2-Pro, GLM 5 Turbo and Step 3.5 Flash are ranked top 3 in terms of token consumption on OpenClaw

Exhibit 29 - Token consumption for OpenClaw (bn)

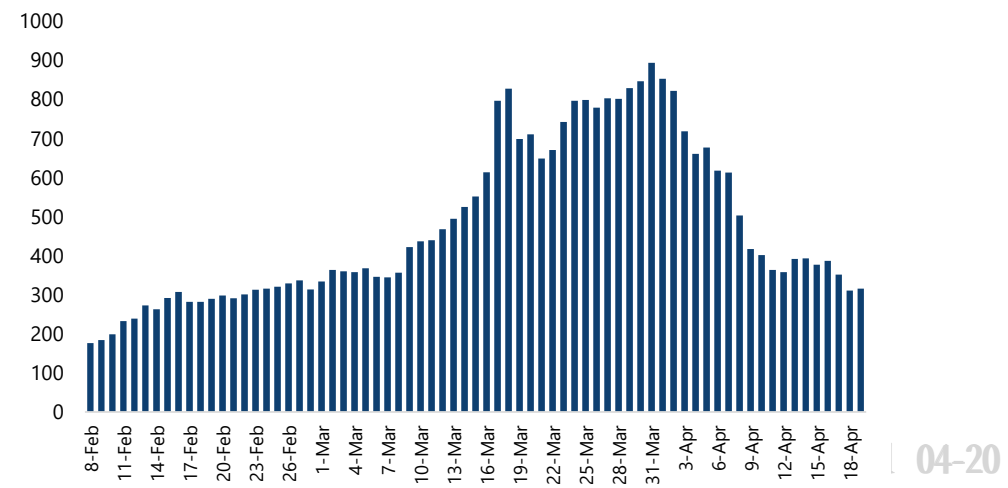

Source: OpenRouter, Jefferies

Exhibit 30 - Top 15 models used by OpenClaw over the last 30 days (as of 20 Apr)

No	Model	Tokens
1	MiMo-V2-Pro	2.68T
2	GLM 5 Turbo	2.29T
3	Step 3.5 Flash	2.28T
4	MiniMax M2.7	1.39T
5	Claude Sonnet 4.6	1.02T
6	Qwen3.6 Plus	965B
7	Nemotron 3 Super	743B
8	Claude Opus 4.6	709B
9	MiniMax M2.5	543B
10	Gemini 3 Flash Preview	409B
11	MiMo-V2-Omni	395B
12	Gemini 2.5 Flash Lite	364B
13	DeepSeek V3.2	326B
14	Kimi K2.5	312B
15	Trinity Large Preview (free)	294B

Source: OpenRouter, Jefferies

According to PinchBench, Anthropic's Claude Opus 4.6 ranked first in terms of success rate (93.3%), followed by Arcee AI Trinity Large Thinking (91.9%) and OpenAI GPT5.4 (90.5%).

Exhibit 31 - Global: Top 10 models by success rate as of 20 Apr

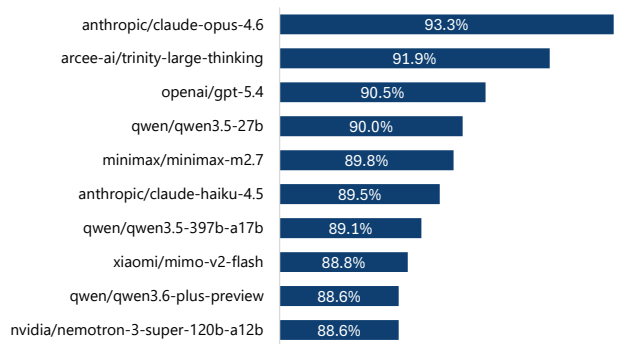

Source: PinchBench, Jefferies

In China market, Qwen 3.5-27B (90.0%) ranked first, followed by MiniMax M2.7 (89.8%) and Qwen 3.5-397B-A17B (89.1%). MiniMax is one of the top 10 models and OpenClaw's founder Peter Steinberger highlights MiniMax's models are recommended due to success rate, speed and cost efficiencies.

MiniMax models are recommended by OpenClaw on its success rate, speed, and cost efficiencies

Exhibit 32 - China: Top 10 models by success rate as of 20 Apr

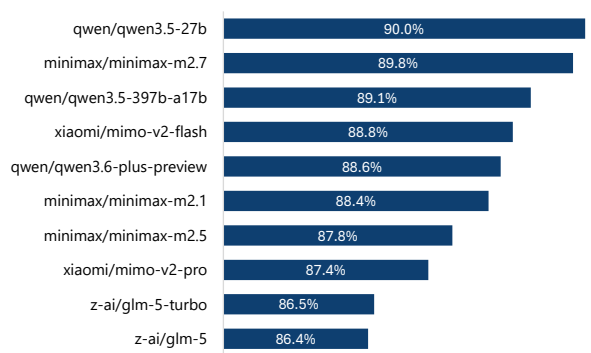

Source: PinchBench, Jefferies

AI Assistants

DAU for a number of leading players declined MoM in Mar after Feb CNY promotions

- MoM trend: DAU for Qwen reached 29m in Mar (vs Feb: 38m) and Yuanbao reached 8m (vs Feb:17m). For Doubao, it increased by 39% MoM to 133m (vs Feb: 96m). For DeepSeek, DAU increased by about 18% MoM to about 29m.

Exhibit 33 - Weekly DAU for leading AI Assistants in China (mn)

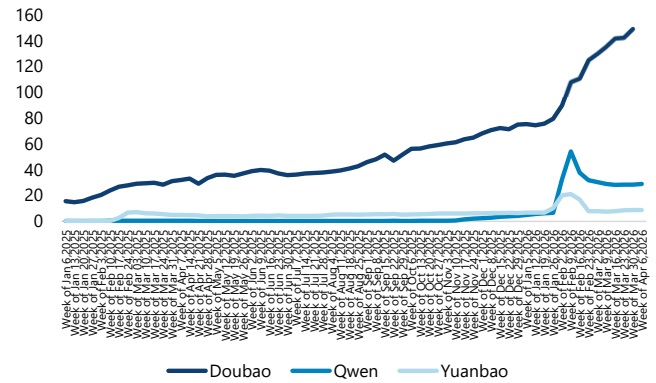

Source: QuestMobile, Jefferies

Exhibit 34 - Monthly DAU for Leading AI Assistants in China (mn)

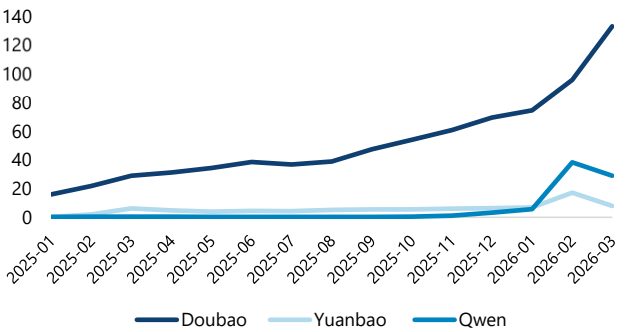

Source: QuestMobile, Jefferies

Exhibit 35 - Weekly user time spent per weekly DAU for leading AI Assistants in China (mins)

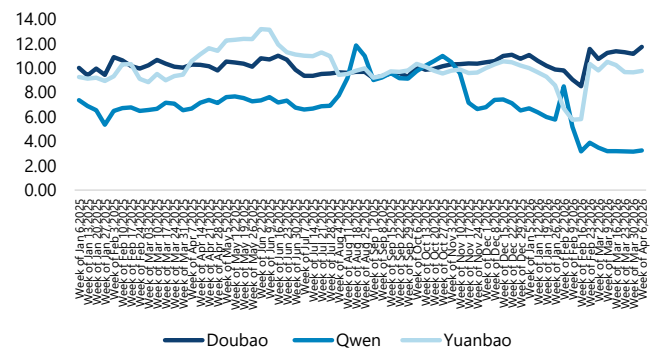

Source: QuestMobile, Jefferies

Exhibit 36 - Monthly user daily time spent per monthly DAU for leading AI Assistants in China (mins)

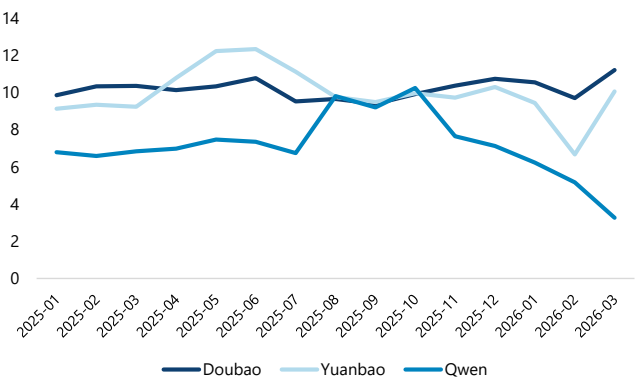

Source: QuestMobile, Jefferies

订阅研报添加微信 LCD123321Cad 04-20

Video generation models

According to Artificial Analysis (as of 20 Apr), HappyHorse-1.0 ranked top 5 in terms of text to video.
Others include Seedance 2.0, SkyReels V4, Kling 3.0 Omni 1080p, Kling 3.0 1080p, and Veo 3.1 Fast.

Exhibit 37 - Top 15 image to video leaderboard (with audio)

No	Model	ELO
1	Dreamina Seedance 2.0 720p	1,181
2	HappyHorse-1.0	1,167
3	SkyReels V4	1,090
4	grok-imagine-video	1,087
5	Veo 3.1 Fast	1,084
6	Veo 3.1	1,082
7	Kling 3.0 1080p(Pro)	1,075
8	Kling 3.0 720p(Standard)	1,067
9	PixVerse V6	1,067
10	Kling 3.0 Omni1080p (Pro)	1,057
11	Vidu Q3 Pro	1,048
12	Kling 3.0 Omni 720p (Standard)	1,047
13	Kling 2.6 Pro(January)	1,006
14	Seedance 1.5 pro	1,000
15	LTX-2.3 Fast	959

Source: Artificial Analysis, Jefferies

Exhibit 38 - Top 15 text to video leaderboard (with audio)

No	Model	ELO
1	HappyHorse-1.0	1,230
2	Dreamina Seedance 2.0 720p	1,223
3	SkyReels V4	1,138
4	Kling 3.0 Omni1080p (Pro)	1,102
5	Kling 3.0 1080p (Pro)	1,101
6	Veo 3.1	1,100
7	Veo 3.1 Fast	1,098
8	Kling 3.0 720p (Standard)	1,094
9	Kling 3.0 Omni 720p (Standard)	1,091
10	Vidu Q3 Pro	1,084
11	Sora 2 (December)	1,083
12	grok-imagine-video	1,083
13	Sora 2 Pro	1,080
14	PixVerse V6	1,067
15	Wan 2.6	1,038

Source: Artificial Analysis, Jefferies

订阅研报添加微信 LCD123321Cad 04-20

订阅研报添加微信 LCD123321Cad 04-20

Exhibit 39 - Video generation model pricing in USD

App	Model	Monthly standard plan (US$)	Monthly credit quota	Credit per video	Length per video	Cost (US$) per video
Kling AI	Video 3.0/3.0 Omni - 720p	8.8	660	27-135	3-15s	0.36-1.8
	Video 3.0/3.0 Omni - 1080p	8.8	660	36-180	3-15s	0.48-2.4
	Video O1 - 720p	8.8	660	30-60	5-10s	0.4-0.8
	Video O1 - 1080p	8.8	660	40-80	5-10s	0.53-1.07
	Video 2.6	8.8	660	50-100	5-10s	0.67-1.33
	Video 2.5 Turbo - 720p	8.8	660	15-30	5-10s	0.2-0.4
	Video 2.5 Turbo - 1080p	8.8	660	25-50	5-10s	0.33-0.67
Dreamina AI	Seedance 2.0 Fast	15	1575	56-210	4-15s	0.53-2
	Seedance 2.0	15	1575	68-225	4-15s	0.65-2.14
	Seedance 1.5 Pro	15	1575	100-240	5-12s	0.95-2.29
	Seedance 1.0 Mini	15	1575	25-50	5-10s	0.24-0.48
Hailuo AI	Hailuo 2.3 - 1080p	9.99	1000	160	6s	1.6
	Hailuo 2.3 Fast - 768p	9.99	1000	30-60	6-10s	0.3-0.6
	Hailuo 2.3 Fast - 1080p	9.99	1000	100	6s	1
	Hailuo 2.0 - 768p	9.99	1000	50-100	6-10s	0.5-1
	Hailuo 2.0 - 1080p	9.99	1000	160	6s	1.6
Google Gemini	Veo 3.1 Lite	With Google AI Plus (US$ 7.99 per month with 200 credits) and Pro (US$ 19.99 per month with 1000 credits) plan			4-8s	
	Veo 3.1	With Google AI Ultra plan at US$249.99 per month with 25,000 credits			4-8s	
PixVerse	PixVerse V6 - 720p	10	1200	10-150	1-15s	0.08-1.25
	PixVerse V6 - 1080p	10	1200	20-300	1-15s	0.17-2.5
	PixVerse C1 - 720p	10	1200	10-150	1-15s	0.08-1.25
	PixVerse C1 - 1080p	10	1200	20-300	1-15s	0.17-2.5
	PixVerse V5.6 - 720p	10	1200	10-100	1-10s	0.08-0.83
	PixVerse V5.6 - 1080p	10	1200	20-160	1-8s	0.17-1.33
	PixVerse V5 - 720p	10	1200	8-80	1-10s	0.07-0.67
	PixVerse V5 - 1080p	10	1200	16-128	1-8s	0.13-1.07
Runway	Gen 4.5	15	625	24-120	2-10s	0.58-2.88
	Gen 4 Video	15	625	60-120	5-10s	1.44-2.88
	Gen 4 Video Turbo	15	625	25-50	5-10s	0.6-1.2

Source: Companies, Jefferies

Exhibit 40 - Video generation model pricing in RMB

App	Model	Monthly standard plan (RMB)	Monthly credit quota	Credit per video	Length per video	Cost (RMB) per video
Kling AI	Video 3.0/3.0 Omni - 720p	58	660	27-135	3-15s	2.37-11.86
	Video 3.0/3.0 Omni - 1080p	58	660	36-180	3-15s	3.16-15.82
	Video O1 - 720p	58	660	30-60	5-10s	2.64-5.27
	Video O1 - 1080p	58	660	40-80	5-10s	3.52-7.03
	Video 2.6	58	660	50-100	5-10s	4.39-8.79
	Video 2.5 Turbo - 720p	58	660	15-30	5-10s	1.32-2.64
	Video 2.5 Turbo - 1080p	58	660	25-50	5-10s	2.2-4.39
	Seedance 2.0 Fast	69	725	20-75	4-15s	1.9-7.14
Jimeng AI	Seedance 2.0	69	725	32-120	4-15s	3.05-11.42
	Seedance 1.5 Pro	69	725	50-120	5-12s	4.76-11.42
	Seedance 1.0	69	725	50-100	5-10s	4.76-9.52
	Hailuo 2.3 - 768p	68	1000	50-100	6-10s	3.4-6.8
Hailuo AI	Hailuo 2.3 - 1080p	68	1000	160	6s	10.88
	Hailuo 2.3 Fast - 768p	68	1000	30-60	6-10s	2.04-4.08
	Hailuo 2.3 Fast - 1080p	68	1000	100	6s	6.8
	Hailuo 2.0 - 512p	68	1000	24-50	6-10s	1.63-3.4
	Hailuo 2.0 - 768p	68	1000	50-100	6-10s	3.4-6.8
	Hailuo 2.0 - 1080p	68	1000	160	6s	10.88
	Vidu Q3 Pro - 540p	59	800	10-160	1-16s	0.74-11.8
Vidu	Vidu Q3 Pro - 720p	59	800	25-400	1-16s	1.84-29.5
	Vidu Q3 Pro - 1080p	59	800	30-480	1-16s	2.21-35.4
	Vidu Q3 Turbo - 540p	59	800	8-128	1-16s	0.59-9.44
	Vidu Q3 Turbo- 720p	59	800	12-192	1-16s	0.89-14.16
	Vidu Q3 Turbo - 1080p	59	800	14-224	1-16s	1.03-16.52
	Vidu Q3 refernce - 540p	59	800	30-160	3-16s	2.21-11.8
	Vidu Q3 refernce - 720p	59	800	60-320	3-16s	4.43-23.6
	Vidu Q3 refernce - 1080p	59	800	75-400	3-16s	5.53-29.5

Source: Companies, Jefferies

订阅研报添加微信 LCD123321Cad 04-20

Global and China AI Mobile App MAU Ranking

Exhibit 41 - Global AI mobile app MAU ranking (Feb-26)

Rank	App	MAU (m)	MoM
1	ChatGPT	955.6	2.69%
2	Doubao	315.31	87.38%
3	Qwen	202.69	552.83%
4	Quark	170.42	18.24%
5	DeepSeek	132.57	2.26%
6	Gemini	114.95	8.23%
7	Tencent Yuanbao	109.45	24.03%
8	Baidu Drive	108.6	-23.40%
9	Grok	76.41	-2.63%
10	Nova	73.59	-1.35%
11	Jimeng AI	70.82	58.11%
12	Dola	59.98	12.61%
13	Meta View	39.77	9.49%
14	AI Mirror	36.92	-5.17%
15	Perplexity	32.39	-4.60%

Source: AICPB, Jefferies

According to aicpb.com, MAU for China AI apps, which include Doubao, Qwen, Quark, and Yuanbao, experienced strong MoM growth vs ChatGPT due to Chinese New Year in Feb

添加微信 LCD123321Cad 04-20
添加微信 LCD123321Cad 04-20
订阅研报添加微信 LCD123321Cad 04-20

Exhibit 42 - China AI Mobile App MAU Ranking (Feb-26)

Rank	App	MAU (m)	MoM
1	Doubao	315.31	87.38%
2	Qwen	202.69	552.83%
3	Quark	170.42	18.24%
4	DeepSeek	132.57	2.26%
5	Yuanbao	109.45	24.03%
6	Baidu Drive	108.6	-23.40%
7	Jimeng AI	70.82	58.11%
8	Kimi AI	24.73	5.77%
9	Nano AI Search	9.7	-4.19%
10	Baidu Wenku	9.45	-10.06%
11	ERNIE Bot	9.25	9.79%
12	Kling	6.09	108.21%
13	Zhipu Qingyan	5.76	-4.37%
14	Lovekey Keyboard	4.89	15.49%
15	Cat Box	4.4	1.14%

Source: AICPB, Jefferies

In domestic market, MAU growth for Kling and Jimeng AI experienced fast MoM growth in Feb

添加微信 LCD123321Cad 04-20
订阅研报添加微信 LCD123321Cad 04-20
订阅研报添加微信 LCD123321Cad 04-20

Mar 26: Monthly App Tracker from QuestMobile Analysis

Online shopping

JD maintains fast DAU growth YoY basis in Mar

- MoM trend: DAU for Taobao/PDD/VIPS increased by 11%/11%/7% respectively in Mar. DAU for JD increased by 1% MoM in Mar.
- YoY trend: DAU for JD increased 16% YoY and Taobao/PDD increased by 4%/9% YoY respectively.

订阅研报添加微信 LCD123321Cad 04-20

Exhibit 43 - MAU for Taobao / PDD / JD in Mar-26

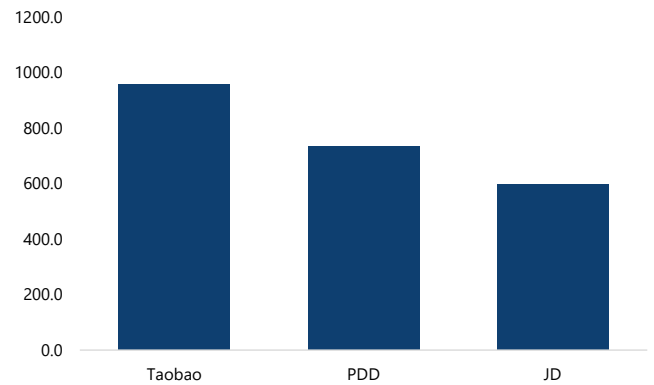

Source: QuestMobile, Jefferies

Exhibit 44 - DAU for Taobao / PDD / JD in Mar-26

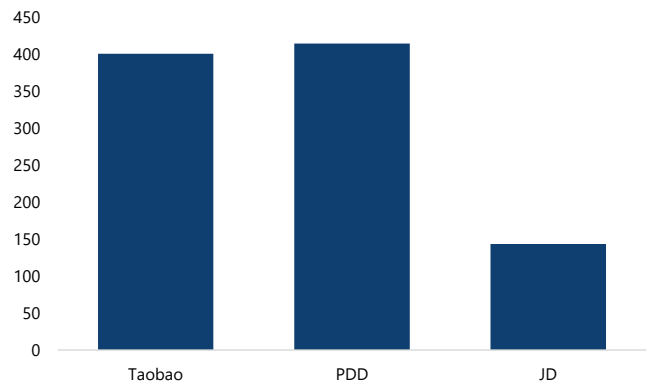

Source: QuestMobile, Jefferies

Food delivery

Meituan MAU is largely stable MoM in Mar

订阅研报添加微信 LCD123321Cad 04-20

- MoM trend: Meituan MAU is flattish MoM during the month.
- YoY trend: Meituan MAU increased by 3% (vs 5% YoY in Feb) in Mar.

Exhibit 45 - MAU for Meituan app in past 12 months

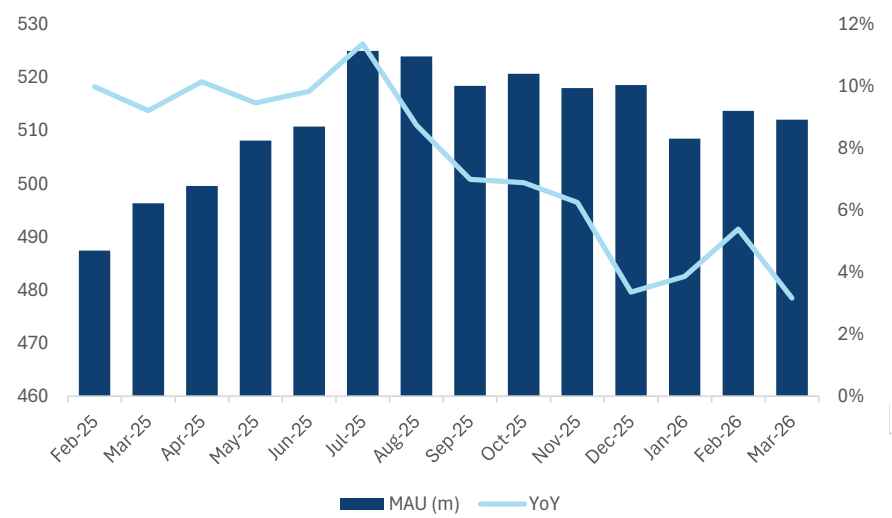

04-20

Source: QuestMobile, Jefferies

Online Music

Soda Music maintains fast MAU YoY growth while largely stable MoM in Mar

- MoM trend: MAU for Soda Music declined by 1% MoM to 156m in Mar. For TME, MAU for QQ Music/Kugou/Kuwo declined by 2%/2%/1% MoM respectively in Mar. MAU for NetEase Cloud Music declined 2% MoM during the month.

- YoY trend: MAU for Soda Music increased by 78% YoY while NetEase Cloud Music MAU increased by 1% YoY in Mar. For TME, MAU for QQ Music increased by 6% while that of Kugou/ Kuwo declined by 6%/15% YoY respectively.

Exhibit 46 - MAU for TME, Soda Music and NetEase Cloud Music in Mar-26

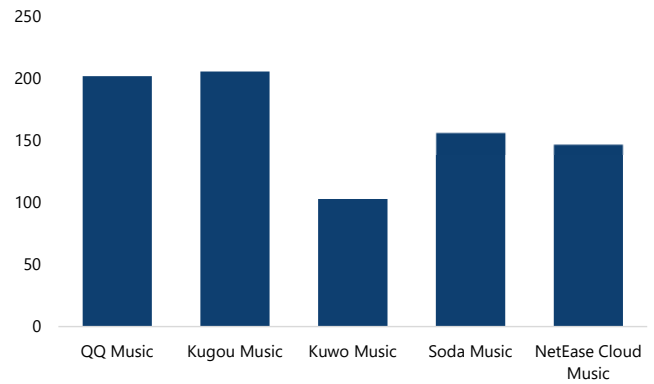

Source: QuestMobile, Jefferies

Exhibit 47 - MAU for Soda Music in past 12 months

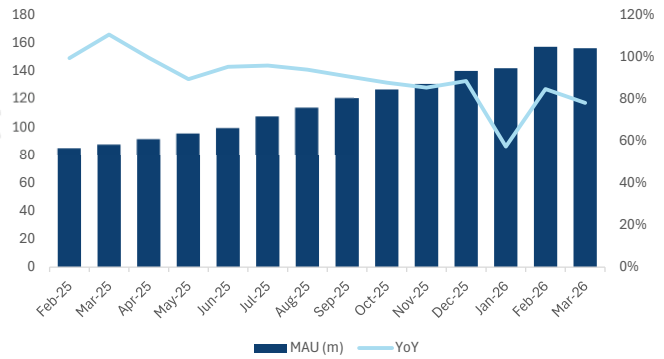

Source: QuestMobile, Jefferies

Online video

MAU for leading long form video platform declined YoY in Mar

- MoM trend: MAU for Tencent Video/IQ / Youku declined by 2%/4%/10% MoM. Mango TV experienced 4% MoM growth in MAU.
- YoY trend: MAU for leading long form video players declined YoY during the month.

Exhibit 48 - MAU for IQ, Tencent Video and Youku in Mar-26

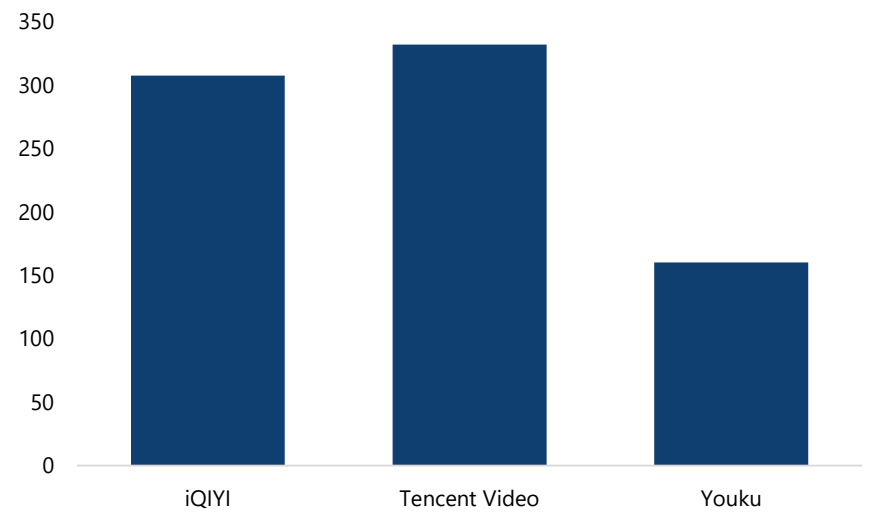

Source: QuestMobile, Jefferies

Live-streaming and short video

Douyin DAU maintained solid YoY growth in Mar

- MoM trend: DAU for Douyin main app increased 2% MoM in Mar. For game broadcasting, MAU for Huya/Douyu increased by 1%/4% MoM in Mar. For BILI, MAU declined 5% MoM during the month.
- YoY trend: DAU for Douyin main app maintains solid growth at 20% YoY. For game broadcasting, both Huya and Douyu experienced 6% and 10% YoY decline in MAU respectively. For BILI, MAU increased by about 3% YoY during the month.

Exhibit 52 - Weibo MAU over the past 12 months

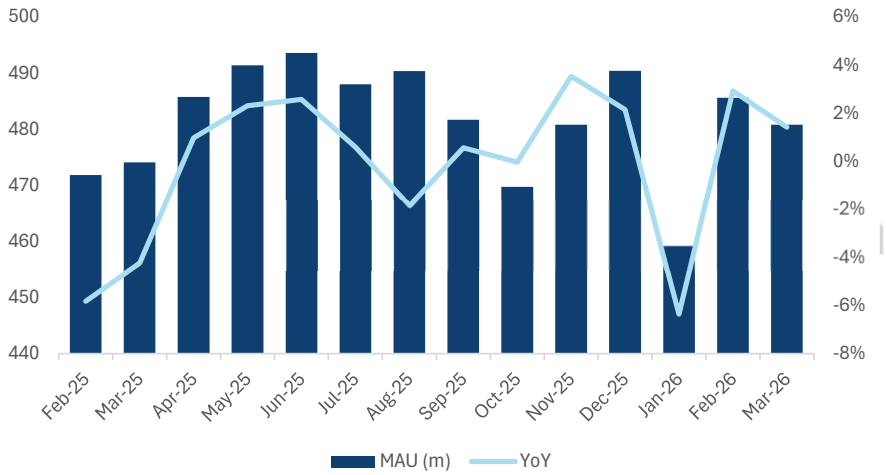

Source: QuestMobile, Jefferies

Baidu app

Baidu app MAU experienced MoM and YoY decline in Mar

- MoM: MAU declined by 1% and DAU increased by 2% in Mar.
- YoY: MAU declined by 6% YoY and DAU declined by 19% YoY in Mar.

Exhibit 53 - Baidu MAU over the past 12 months

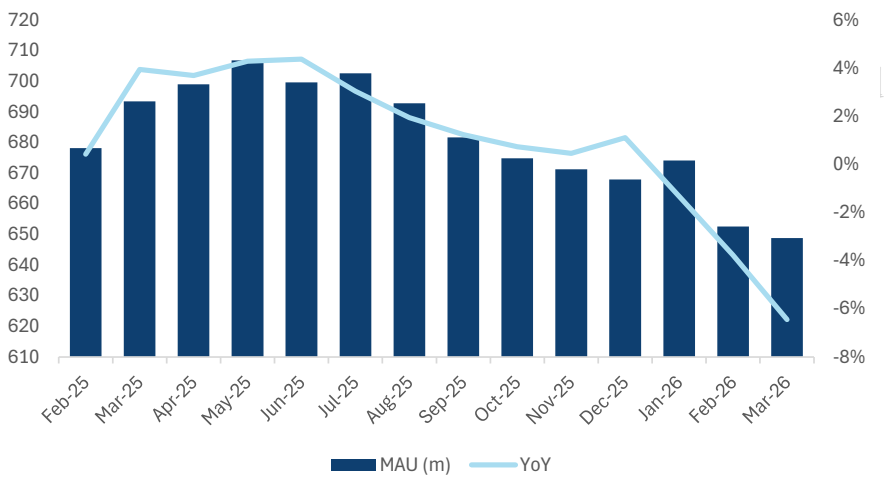

Source: QuestMobile, Jefferies

04-20
04-20
订阅研报添加微信 LCD123321Cad 04-20

Appendix

Omdia heatmap for agentic AI development platform in Asia and Oceania

According to Omdia, Alibaba Cloud is ranked as one of the leaders and top choice for building and deploying AI agents. It captures the highest ranking for 5 out of 7 metrics

According to Omdia, agentic AI software market is expected to increase at 105% CAGR from USD271m in 2025 to USD9.7bn by 2030. Alibaba Cloud captures the highest rankings in 5 areas: (1) context engineering, (2) model support, (3) multi-agent framework, (4) operation and lifecycle management, and (5) open source and community.

订阅研报请加微信 LCD123321Cad 04-20

Exhibit 54 - Omdia heatmap for agentic AI development platform in Asia and Oceania

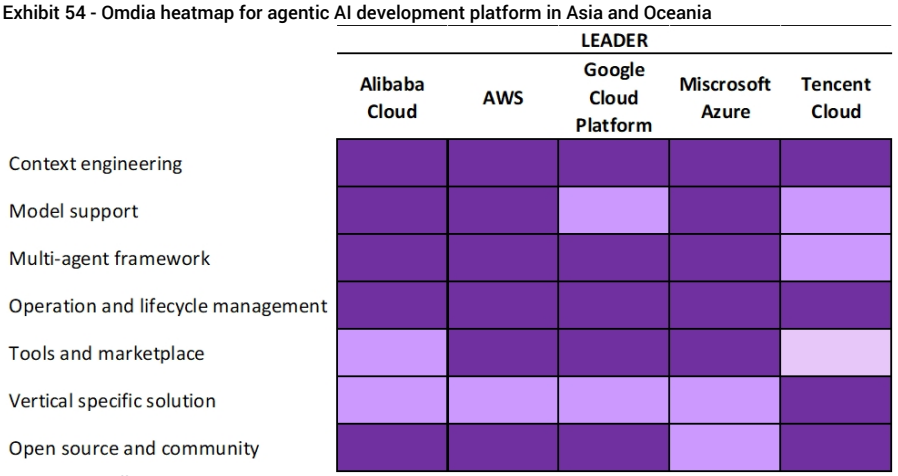

Source: Omdia, Jefferies

Market size: Coding, Workspace, and Content Creation

Global AI code tools market is expected to be 3x the scale in 2030 vs 2025

According to Grand View Research, the global AI code tools market is expected to increase by 3x to USD26bn from 2025 to 2030. On the other hand, F&S predicts that the China AI code generation market will reach about USD4.7bn in 2028. KA customized solutions are expected to be key revenue contributor.

Exhibit 55 - Global AI code tools market size (USDbn)

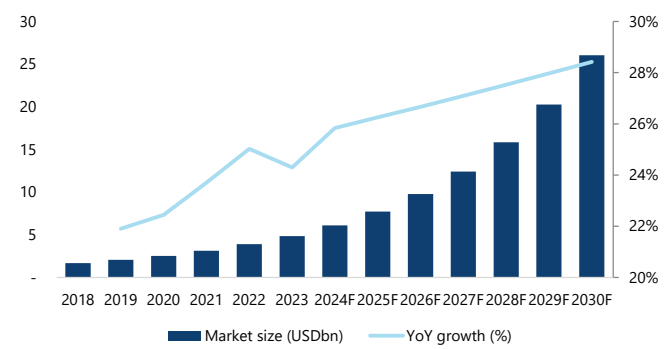

Source: Grand View Research, Jefferies

Exhibit 56 - China AI code generation market size by service type (RMBbn)

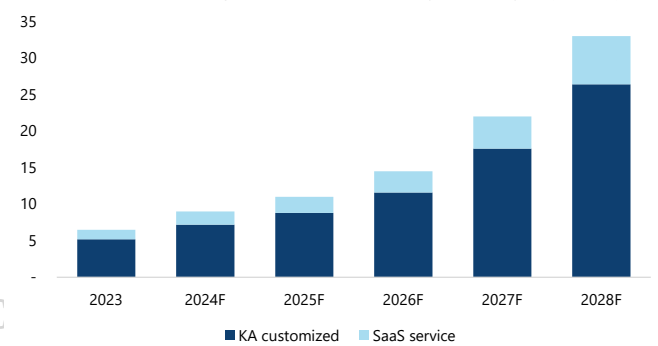

Source: Frost & Sullivan, Jefferies

Exhibit 57 - China AI code penetration in different industries

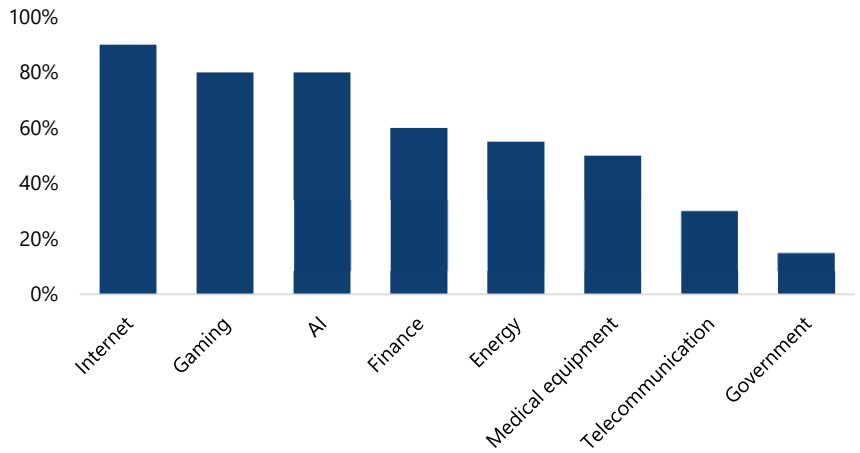

Source: Frost & Sullivan, Jefferies

In terms of AI code penetration, internet, gaming, AI, finance, and energy are top 5 key sectors

According to Gonyn Industry Research (2024), market size of AI+Workspace software is expected to reach RMB191bn in 2028, 6x the scale vs 2024.

Exhibit 58 - China AI + Workspace Market Size

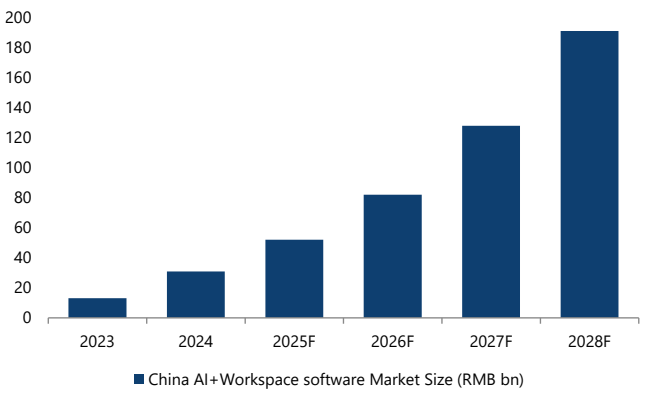

Source: Gonyn, Jefferies

Exhibit 59 - Generative AI in Content Creation Market Size 2025 to 2035 (USD bn)

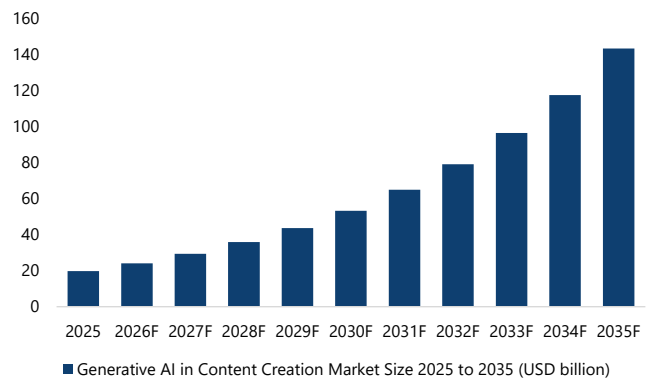

Source: Precedence Research, Jefferies

According to Precedence Research, global generative AI content creation market is expected to increase by 7x surpassing USD140bn in 2035 vs 2025.

Exhibit 60 - Daily token consumption in China by enterprises in 2H25 vs 1H25 (tn)

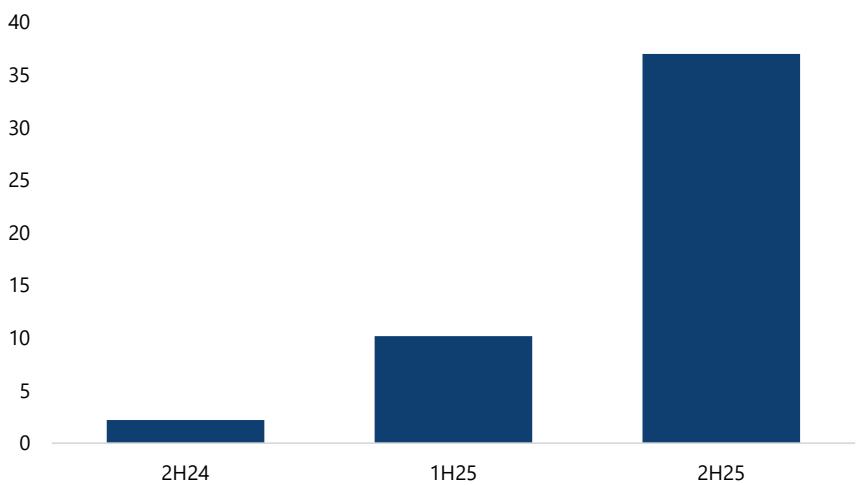

According to F&S, enterprise daily token consumptions increased by 263% HoH to 37tn, among which market share for Alibaba Cloud increased the most and captures no 1 position reaching 32% in 2H25 (1H25:17.7%), followed by 21.3% for Doubao (1H25: 14.1%) and 18.4% for DeepSeek (1H25: 10.3%).

Alibaba's Qwen ranked first in terms of token consumption among enterprises in 2H25

Exhibit 61 - Market share in terms of tokens used among enterprises in 1H25

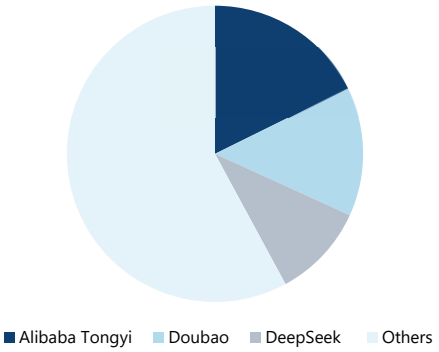

Source: Frost & Sullivan, Jefferies

Exhibit 62 - Market share in terms of tokens used among enterprises in 2H25

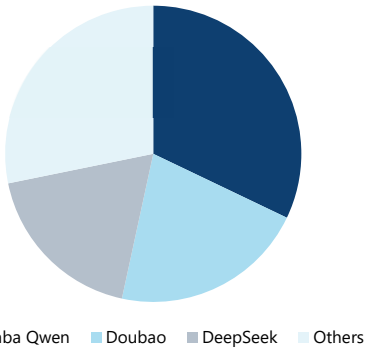

Source: Frost & Sullivan, Jefferies

«¢— x b¥mûR _®Oá LCD123321Cad 04-20

«¢— x b¥mûR _®Oá LCD123321Cad 04-20

China Large Language Models

According to F&S / CAICT, China LLM market is expected to increase by about 64% CAGR between 2024 and 2030 and surpass RMB100bn in 2030. By segment, consumer revenue is expected to grow at a slower rate vs enterprise as willingness to pay remains low.

China large language models are expected to increase by 64% CAGR between 2024 to 2030. By segment, on-premise is key revenue contributor

Exhibit 63 - Market size of China LLM (RMB billions)

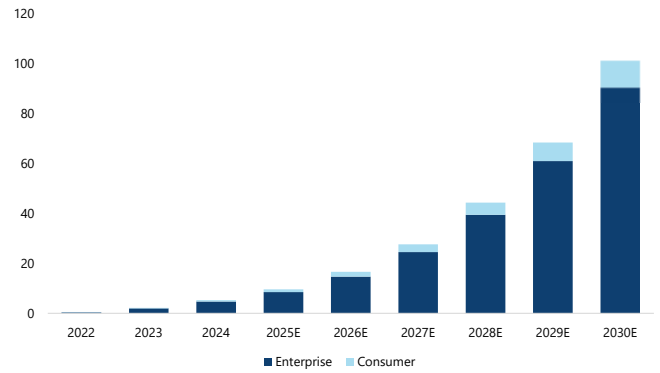

Source: Frost & Sullivan, CAICT, Jefferies

Exhibit 64 - Market size of China Enterprise LLM (RMB billions)

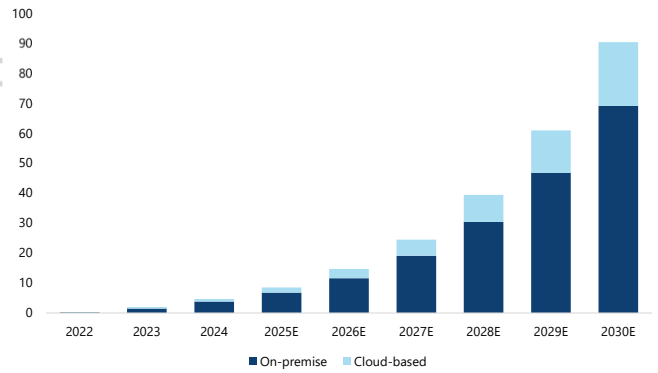

Source: Frost & Sullivan, CAICT, Jefferies

OpenAI

According to The Information, OpenAI revenue is expected to reach USD200bn in 2030, over 6x the scale in 2026. Rising Agent revenue contribution is expected over the next few years.

Exhibit 65 - OpenAI revenue between 2023 and 2030F (as of 3Q25)

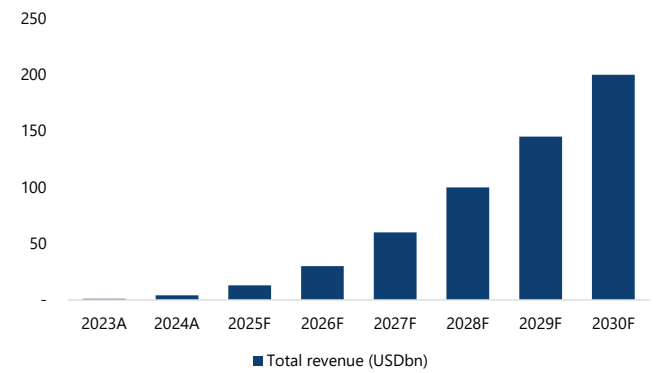

Source: The Information, Jefferies

Exhibit 66 - Revenue mix of OpenAI between 2024 and 2030F (as of 3Q25)

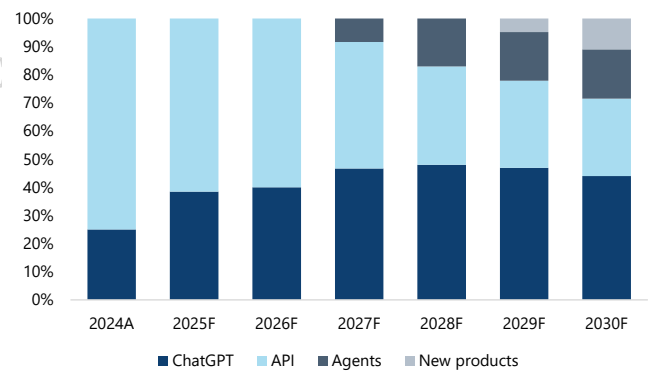

Source: The Information, Jefferies

Exhibit 67 - MiniMax: % of OpenAI revenue over the next few years

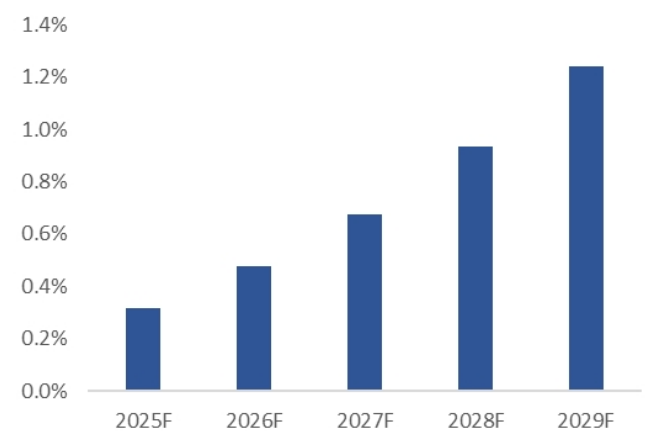

.CD123321Cad 04-20

Models of Leading Players in China

Exhibit 68 - Alibaba model list

Release Date	Model	Core Functions
Jan-24	Qwen-VL-Plus / Max	Enhanced visual reasoning, high-resolution OCR, and detailed image understanding via API.
Feb-24	Qwen1.5	First major open-source iteration with improved multilingual support and 0.5B to 72B sizes.
Jun-24	Qwen2	Optimized for coding, mathematics, and long-context processing (up to 128K tokens).
Aug-24	Qwen2-Audio	Native audio-to-text and audio-to-audio dialogue capabilities.
Aug-24	Qwen2-VL	Introduced dynamic resolution for visual inputs and long video understanding.
Sep-24	Qwen2.5	Significant upgrades in instruction following, structured data (tables) processing, and coding.
Nov-24	QwQ-32B-Preview	Reasoning-focused model using Reinforcement Learning for complex logic and chain-of-thought.
Jan-25	Qwen2.5-Max	Large-scale MoE (Mixture of Experts) flagship, matches or exceeds GPT-4o benchmarks.
Jan-25	Qwen2.5-VL	"Agentic" vision model capable of identifying and interacting with mobile/PC UI screens.
Mar-25	Qwen2.5-Omni	End-to-end full-modality model supporting seamless real-time text, audio, and visual interaction.
Apr-25	Qwen3 (Series)	New reasoning architecture with hybrid inference modes to reduce hallucinations in complex tasks.
Sep-25	Qwen3-Max-Preview	Trillion-parameter preview version optimized for massive enterprise-scale applications.
Jan-26	Qwen3-Max	Frontier reasoning, 1T+ params, agentic workflow mastery.
Feb-26	Qwen3.5	Latest flagship series featuring a 397B MoE model with superior cross-lingual reasoning.
Feb-26	Qwen3.5-Plus	Outstanding performance comparable to the very latest state-of-the-art models with signiicant leaps in both pure-text and multimodal capabilities compared to the 3 series
Feb-26	Qwen3-Coder-Next	Multi-turn tool interactions, future-ready development support.
Mar-26	Qwen3.5-Omni	Support for 256k long-text input and can process more than 10 hours of audio input and over 400 seconds of 720P audio-visual input at 1 FPS
Apr-26	Qwen3.6-Plus	Native multimodal, 1M context, agentic coding
Apr-26	Happy Oyster	AI model that creates interactive 3D scenes with realistic physics for game, film, and TV.

Source: Company, Jefferies

Exhibit 69 - Tencent model list

Release Date	Model	Core Functions
May-24	Hunyuan-DiT	The industry's first Chinese-native Diffusion Transformer (DiT) for high-quality text-to-image generation.
Sep-24	Hunyuan-Turbo	An upgraded LLM based on the Mixture-of-Experts (MoE) architecture, doubling training efficiency and reducing inference costs.
Nov-24	Hunyuan3D 1.0	First open-source 3D generation model supporting text-to-3D and image-to-3D workflows.
Dec-24	HunyuanVideo	A large-scale video generation framework capable of producing high-definition videos with complex motion.
Jan-25	Hunyuan3D 2.0	Significant upgrades in geometry accuracy and texture representation for 3D asset creation.
Feb-25	Hunyuan Turbo S	A "fast-thinking" model using Hybrid-Mamba-Transformer architecture, optimized for near-instant responses and competitive reasoning.
Mar-25	HunyuanVideo-I2V	An specialized image-to-video model built on the foundation of HunyuanVideo.
Jun-25	Hunyuan3D 2.1	Production-ready 3D asset generation with a new PBR (Physically Based Rendering) material model.
Sep-25	Hunyuan Image 3.0	Flagship image generation model (80B parameters) topping global rankings for semantic understanding.
Nov-25	HunyuanVideo 1.5	A highly efficient and powerful update to the foundation video generation model.
Dec-25	Hunyuan 2.0	Next-generation LLM featuring 406B parameters and a 256K context window for complex reasoning tasks.
Jan-26	Hunyuan Image 3.0-Instruct	An image-to-image editing model supporting precise single-sentence editing and multi-image fusion.
Mar-26	HY-WU	A scalable framework for on-the-fly conditional generation of low-rank(LoRA) updates.

Source: Company, Jefferies

Exhibit 70 - Baidu model list

Release Date	Model	Core Functions
Mar-24	ERNIE Speed / Lite / Tiny	Lightweight models optimized for low latency and high cost-efficiency.
Jun-24	ERNIE 4.0 Turbo	Efficiency upgrade for the 4.0 architecture with faster response speeds.
Mar-25	ERNIE 4.5	Baidu's first native multimodal foundation model (text, image, audio, video).
Mar-25	ERNIE X1	Deep-thinking reasoning model designed for complex logic and multi-step tasks.
Apr-25	ERNIE 4.5 Turbo / X1 Turbo	Performance-optimized version of X1 with significantly lower inference costs.
Sep-25	ERNIE X1.1	Major upgrade in factuality, instruction following, and agentic capabilities.
Nov-25	ERNIE 5.0 (Preview)	Preview of next-gen omni-modal model with advanced real-time reasoning.
Jan-26	ERNIE 5.0 (Stable)	Current flagship 2.4T parameter model; top-ranked Chinese model on LMArena.

Source: Company, Jefferies

Exhibit 71 - MiniMax model list

Release Date	Model	Core Function
Jan-25	MiniMax-Text-01/MiniMax-VL-01	New generation of text and vision language model
Feb-25	T2V-01-Director/I2V-01-Director	Director-level video generation model, better adherence to camera movement description instructions
Apr-25	Image-01	Support generating images of various sizes from text descriptions
Jun-25	MiniMax-M1	80K mind chain and 1M input
Jun-25	MiniMax-Hailuo 02	Native 1080p, SOTA instruction following, extreme physics mastery
Jun-25	Music-1.5 (Beta)	Support music generation by inputting musical inspiration and lyrics
Aug-25	Sppeech-2.5	Significantly enhanced multilingual performance, more realistic and accurate voice cloning, over 40 languages
Sep-25	Music-1.5	Extending song generation to a full 4minutes, one person is the band
Oct-25	MiniMax M2	Open-source MoE, strong code & agent
Oct-25	MiniMax-Hailuo-2.3	Complex video performance & media agent
Oct-25	MiniMax-Speech-2.6	Rea-time response, intelligent parsing, fluent LoRA voice
Oct-25	Music-2.0	A lead forward in model's understanding and expression of music
Dec-25	MiniMax M2.1	Exceptional multilingual capabilities to elevate code engineering
Jan-26	Music-2.5	Creation more precise and music more real
Jan-26	MiniMax M2-Her	Multi-character roleplay, immersive long-horizon interaction
Feb-26	MiniMax M2.5	Industry-leading coding and reasoning capabilities at a fraction of the cost
Mar-26	Music-2.5+	Enhanced creation experience & efficiency, smarter melodies, purer sound quality
Mar-26	MiniMax M2.7	M2.7 is capable of building complex agent harnesses and completing highly elaborate productivity tasks, leveraging capabilities such as Agent Teams, complex Skills, and dynamic tool search.

LCD123321Cad 04-20

Exhibit 72 - Moonshot AI model list

Release Date	Model	Core Function
Jan-25	Kimi k1.5	An o1-level multi-modal model
Apr-25	Kimi-VL	Offers advanced multimodal reasoning, long-context understanding, and strong agent capabilities
Apr-25	Kimina-Prover Preview	Can reason in a human-like way and prove mathematical theorems in the Lean 4 language
Apr-25	Kimi-Audio	Open-source audio foundation model excelling in audio understanding, generation, and conversation
Jun-25	Kimi-Dev	A strong and open-source coding LLM for issue resolution
Jun-25	Kimi-Researcher	End-to-end RL training for emerging agentic capabilities
Jul-25	Kimi K2	A state-of-the-art MoE language mode, featuring 32 billion activated parameters and a total of 1 trillion parameters
Sep-25	Kimi-K2-Instruct-0905	The post-trained model best for drop-in, general-purpose chat and agentic experience
Nov-25	Kimi K2 Thinking	Execute up to 200-300 sequential tool calls without human interference, reasoning coherently across hundreds of stpes to solve complex problems
Jan-26	Kimi K2.5	A multimodal model, state-of-the-art coding, vision capabilities, self-directed agent swarm paradigm

Source: Company, Jefferies

Exhibit 73 - Z.ai model list

Release Date	Model	Core Function
Jul-25	GLM-4.5-Air	Adopts a more streamlined design with 106B total parameters and 12B active parameters
Jul-25	GLM-4.5	Total parameter count of 355B with 32B active parameters per forward pass
Aug-25	GLM-4.5V	Achieved SOTA performance among open-source VLMs of the same level with 106B parameters and 12B activation parameters
Sep-25	GLM-4.6	Advanced agentic, reasoning and coding capabilities
Dec-25	AutoGLM	Autonomous foundation agents for GUIs
Dec-25	GLM-4.6V	Open source multimodal models with native tool use
Dec-25	GLM-ASR-Nano	A robust, open-source speech recognition model with 1.5B parameters
Dec-25	GLM-TTS	Zero-shot TTS with multi-reward reinforcement learning
Dec-25	GLM-4.7	Complex systems engineering and long-horizon agentic tasks
Jan-26	GLM-4.7-Flash	Open source and free
Feb-26	GLM-OCR	SOTA performance, mastering complex document recognition
Feb-26	GLM-5	MoE flagship, complex engineering & long-horizon agent
Mar-26	GLM-5-Turbo	GLM-5-Turbo is a foundation model deeply optimized for the OpenClaw scenario. It has been specifically optimized for the core requirements of OpenClaw tasks since the training phase, enhancing key capabilities such as tool invocation, command following, timed and persistent tasks, and long-chain execution.
Apr-26	GLM-5V-Turbo	It strengthens multimodal tool use and GUI task execution, allowing the model to better interpret interfaces, design mocks, documents, and charts, plan the next action, and complete end-to-end Agent workflows while maintaining strong pure-text coding and reasoning performance.
Apr-26	GLM-5.1	Built for long-horizon tasks, it runs autonomously for 8 hours, refining strategies through thousands of iterations, and it delivers significant leaps in coding and agentic performance with 6x performance boost. In high-performance database optimization, it reached 21.5k QPS over 600+ iterations and 6,000+ tool calls, which is 6x the performance of a standard 50-turn session

Source: Company, Jefferies

Exhibit 74 - Token share by region

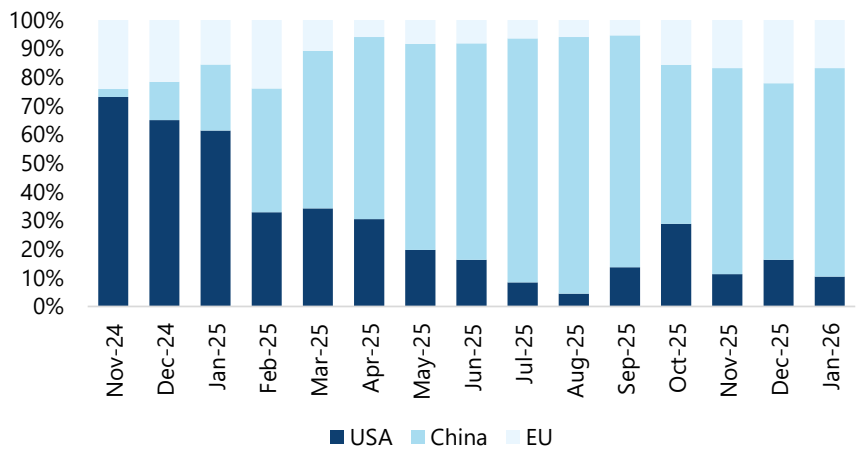

Source: OpenRouter, Jefferies

订阅研报添加微信 LCD123321Cad 04-20

Exhibit 75 - China Daily token usage

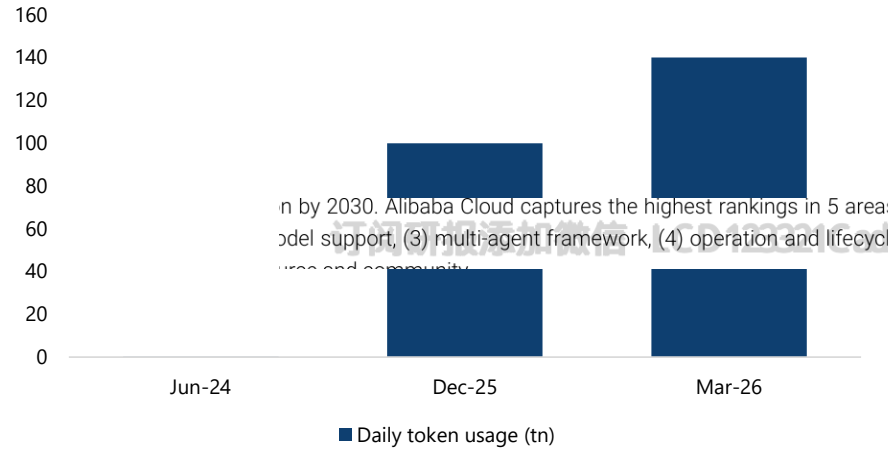

Source: National Data Administration (NDA), Jefferies

Exhibit 76 - Doubao daily token consumption

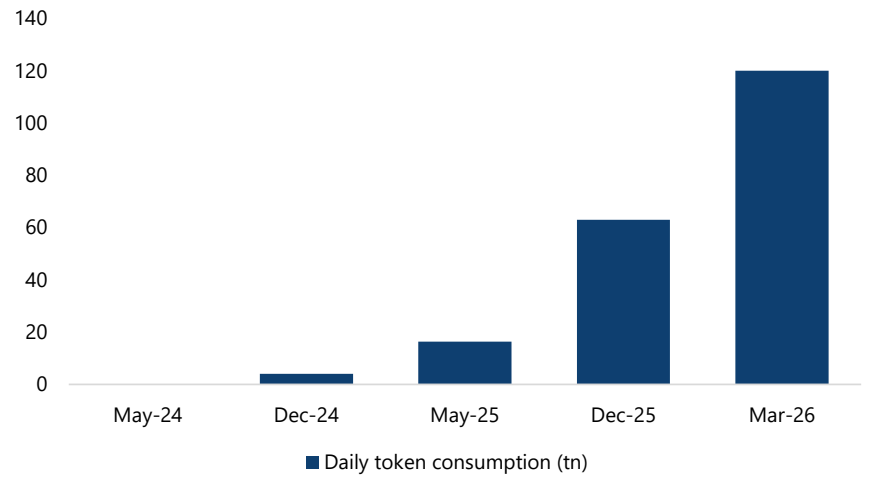

Source: VolcanoEngine, Jefferies

订阅研报添加微信 LCD123321Cad 04-20