Global Artificial Intelligence

17 Apr 2026 06:38:37 ET

CITI'S TAKE

Over 125 of the leading public and private AI companies, including Anthropic, OpenAI, Google, Amazon, IBM, Broadcom, Cursor, Micron, and others, joined the 4th Annual Citi AI Summit in Menlo Park over the last two days to discuss the future of artificial intelligence. While the dominant theme of the event was acceleration – in both the underlying AI technologies as well as the enterprise adoption of those technologies – that was tempered by cybersecurity concerns created by Anthropic's new Mythos model, constraints in compute capacity limiting AI utilization, and the rising price of frontier models. In the note below, we detail the major themes and investment conclusions from the last two days and share summaries of the 47 keynotes, fireside chats, and panels.

订阅研报添加微信 LCD123321Cad

订阅研报添加微信 LCD123321Cad

Contents

AI Summit Key Conference Takeaways 4
Key Internet Takeaways 5
AI Product Cycles Compress as Productivity Gains Accelerate 5
Key Application + Data Management Software Takeaways 5
Key Cybersecurity Takeaways 6
General Sentiment/Views on Mythos Across Sessions 6
Key Infra Software Takeaways 6
Stock selection "Rorschach Test" 6
Key Datacenter and Telco Operator Takeaways 7
Day One Takeaways 8
Thinking in Decades, Not Quarters: A Fireside Chat with Vinod Khosla 8
AWS Keynote: Applied AI at Scale 8
Five top insights from our discussion with IBM Research Chief Scientist Dr. Ruchir Puri 9
Revolutionizing Software Development 10
Lumentum: The Optical Backbone of AI Infrastructure 10
Defining the Frontier: How Foundational Models Are Evolving 11
DigitalOean: Building the Inference Factory 11
Agents in the Enterprise 12
The Future of Silicon Panel Perspectives 12
Building AI Factories 13
Replit: The Democratization of Software Creation 13
Keynote: Google Cloud 13
Anthropic Keynote: Enterprise AI at an Inflection Point 14
Allocating the AI Future: Where Growth Equity Is Going Next 15
Securing AI Panel Perspectives 16
Intelligent Finance Panel 16
Decagon: Redefining Customer Service 17
Vertical Intelligence Panel 17
Salesforce: Agents, Outcomes and the Enterprise 18
The Rise of Voice AI: Real-Time Intelligence in Action 18
Pegasystems: Determinism Over Autonomy 19
The Intelligent Enterprise 19
Perplexity: Search, Trust, and the Anti-Ad Model 20
Keynote: Cursor 20
Day Two Takeaways 22
OpenAI Keynote: From Hype to Habit 22
Fireside Chat with Adobe 23
Architecting Data Panel 23
Fireside Chat with Applied Intuition 24
Infrastructure Foundations of AI Panel 24
Fireside Chat with Mercor 25
Revolutionizing Workflows 25
Fireside Chat with Astera Labs 26
AI in Healthcare 26
Bringing AI Apps into Production 27
Glean Fireside: The Connective Tissue of AI 27
The VC Perspective on AI: What's Next? 28
Financing the AI Build-Out: Public and Private Credit Markets 29
Fireside Chat with Shield AI 29
Lightmatter Fireside: The Next 1,000X 30
Physical AI 30

订阅研报添加微信 LGD12321Cad 04-17

Intelligent Automation Panel 31
Conversation with Harvey: Transforming Legal Through AI 31
Embedding AI Into Workflows: Customer Perspectives 32
From Model to Motion: The Skild AI Approach 32
Broadcom Keynote: Custom at Scale 33
Appendix A-1 36

订阅研报添加微信 LCD123321Cad 04-17

订阅研报添加微信 LCD123321Cad 04-17

订阅研报添加微信 LCD123321Cad 04-17

AI Summit Key Conference Takeaways

AI acceleration is inverting the structure of work — The most consistent message across nearly every session was the acceleration of AI technology and enterprise adoption creating a fundamental inversion in how work gets done. Cursor's Michael Truell noted the acceleration as developers have gone from being 5-10% more productive with early models to running 2-3 agents at a time and accounting for 80% of enterprise customers code development. Ashley Kramer, Vice President of Enterprise at OpenAI, noted the pivotal moment we're in with "everybody moving from pilot to production." Vinod Khosla framed it most starkly: we are moving from giving AI to people to people working for AI as the organizing structure of the enterprise. AWS reinforced this with their AI teammate model, where everyone becomes a manager of AI agents, roles collapse across specializations, and the bottleneck shifts from execution to problem definition. Companies that are still treating AI simply as a productivity add-on are rearranging deck chairs. The leaders are rebuilding the ship. 2026 is the year enterprise AI scales and the gap between leaders and laggards grows larger.

Mythos signals a step-change in AI capability...and concerns — Anthropic's deliberate decision to withhold its most capable model, Mythos, from broad release, citing cybersecurity and other risks, was the conference's most-discussed event. We believe Anthropic's Mythos model represents a material leap in AI capability, but the more urgent implication is that Mythos-class models are expected to reach open source within 6–9 months, creating a risk of potentially putting frontier-level capability in the hands of threat actors. Enterprises should expect an increase of AI-driven breaches over the next 12 months, and the security paradigm must shift from detection to real-time runtime controls. Scott White, Anthropic's Head of Enterprise Product, noted the work they are doing with enterprises in Project Glasswing to help address those risks using these advanced technologies, a level of responsibility that other companies also recognized.

Capacity constraints — Compute demand is insatiable through 2028 at minimum, with the infrastructure panel placing supply/demand rebalancing at 2029 at the earliest or potentially never. The most-discussed binding constraint was power, not silicon. Lightning AI's William Falcon quantified the demand inflection: a single million-agent business deployment implies 100x growth in inference demand and a 1,000x need for compute. CoreWeave has secured 3GW of contracted power to expand from 850MW. The constraints are clearly showing up in higher prices for compute as well as just general unavailability for workload capacity at scale.

Physical AI — Physical AI emerged as the conference's most consensus-forward investment theme, with deployments already underway and capital formation accelerating. Applied Intuition's Peter Ludwig expects widespread autonomous trucking and ground defense vehicles by 2030, while Shield AI announced a next-generation autonomous aircraft with a first flight planned for fall 2026 and production targeted for 2029–2030. The Future of Silicon panel warned that physical AI will require entirely new model architectures that the market is materially underestimating.

Key Internet Takeaways

AI Product Cycles Compress as Productivity Gains Accelerate

Accelerating enterprise adoption (as agentic workflows scale to production), increasing productivity gains, and faster product releases given advances in compute were some of our core takeaways from the Summit. The net result, in our view, is that usage and revenue curves are converging at increasing speed leading to greater AI adoption cycles.

订阅研报添加微信 LCD123321Cad 04-17

The net-result, we emerge from the Summit incrementally positive on GOOGL and AMZN given Google Cloud and AWS adoption trends. We highlight investments in compute, infrastructure, and agentic capabilities as broad-based enterprise demand accelerates, and we highlight our conversations emphasized gains in inference optimization resulting in materially lowering unit economics, thereby broadening use cases – this, as we believe the launch of **Meta's** Spark Muse model and investments in agents have 'sparked' a re-emergence of Meta's AI product strategy.

On agentic adoption, multiple panelists highlighted productivity benefits, with managers increasingly directing agents to find improvements across multiple workstreams with one company calling out inbound agentic workflows leading to a 15% improvement in conversion rates of new sales. We note **Meta's Manus** (acquired in December) was highlighted as an early leader in agent-native applications, pushing beyond chat and research into real task completion resulting in a $0 to ~$100M ARR in ~8 months. Since Meta's acquisition of Manus, the platform has been integrated into Meta's Ads Manager, Instagram Creator Marketplaces, and WhatsApp Business, further accelerating Meta's agent ambitions.

订阅研报添加微信 LCD123321Cad 04-17

Additional takeaways from the AI Summit:

(1) Vertical depth — Larger models increasingly have greater vertical depth, though we note the emerg*ence of*, *and* demand for, vertical solutions.**(2) Human intelligence** remains a key part to model building with financial services, legal, life sciences, and healthcare the verticals we consistently heard were among the fastest growing use-cases for AI. **(3) Open and Closed** model use cases are increasingly based on task, cost, and risk, among other variables. **(4) Voice AI** benefits spans multiple verticals, including B2B (like customer service), consumer, eCommerce, coding, and scheduling creating a significant opportunity as latency, reliability, and context recognition scale. **(5) Physical AI** is answering labor shortage challenges across specific verticals (like mining, farming, and eCommerce, among others), as every machine / robot becomes autonomous.

订阅研报添加微信 LCD123321Cad 04-17

Key Application + Data Management Software Takeaways

Coming away from two days of meetings and presentations across unique privates, publics (CRM, ADBE) and IT decision makers, we retain our cautious view on most application software. First-party native AI solutions felt less of an emphasis at large incumbents, with some privates noting continued pressure on recently consumption /conversation-based pricing. We also heard more signs of budgetary pressure from AI spending and consolidation negatively influencing IT and traditional application software budgets. Within vertical software, our

conversations suggest the frontier model leaders are increasingly verticalizing, though the remains opportunity around offering more specialized/fine-tuned and cost effective models especially around proprietary data sets. Within data management, we continued to pick up positive signs on data volume/query growth led by increased app/code creation and AI democratization tailwinds. Finally across AI infrastructure, presenters across neoclouds and AI factory endorsed intensifying demand, persistent pricing strength but also multi-faceted bottlenecks (power shells, labor etc.).

订阅研报添加微信 LCD123321Cad 04-17

Key Cybersecurity Takeaways

General Sentiment/Views on Mythos Across Sessions

A Mythos-like offering was going to come sooner or later – More AI-driven capabilities in cybersecurity were generally expected across the vendor spectrum (from cyber start-ups to Google). Here, some believe that Mythos capabilities will quickly become generally available in open source (within NTM) as emerging AI apps/use cases will broaden the attack surface (e.g. OpenClaw) and require AI-level defense to properly address AI-powered threats. To better secure this new AI-coded attack environment, many see the need to shift security postures from a focus on defense/detection to runtime controls. We think this dynamic will disproportionately benefit cyber vendors at critical infrastructure control points (network, cloud, endpoint) who have incumbency with inline cyber defense competence and protocols.

Vulnerability identification/management is one (key) area of opportunity for frontier models – Along with incident investigation, and remediation. This largely validates our initial assessment of which cyber segments are likely to eventually be in the path of disruption.

Mythos's measured release was because of safety concerns – Mythos's limited release through Project Glasswing (vs. broad release) was because of the model's enhanced capabilities that naturally carried greater safety risks (cyber and CBRN concern). Anthropic does not intend to release it broadly until risks are fully understood.

Key Infra Software Takeaways

Stock selection "Rorschach Test"

Across two full days of proprietary presentations, we distill three high-level meta-themes/conclusions and tag positively levered companies: **(1)** Agent proliferation, scale, persistence, and autonomy all stand to magnify architectural complexity and consequently we think points of system-failure (DT, DDOG, tangentially RBRK). **(2)** Code assistants can equally and proficiently write exploits rapidly, pushing the marginal cost of creating/launching a cyberattack, which in tandem with prompt injection, data poisoning, and uncontrolled agent behavior risks (as enterprise AI/agent deployments scale) raises defensive stakes (PANW, CRWD). **(3)** Agents, latency-sensitive AI apps, and intelligence/compute-hungry AI native ISV startups necessitating purpose-built edge compute/highest-ROI-model-routing that command better/diverse economics than neoclouds underlying infra, where LLM competitive product intros (like Claude Managed Agents) are not competitive but category-validating (DOCN, NET, lesser extent AKAM).

Key Datacenter and Telco Operator Takeaways

We found multiple reasons to be enthusiastic about AI demand expansion: introduction of more capable agents, consumerization of agents for the enterprise, securitization of the data, and decentralization of workloads. A large focus across panels and fireside discussions was the opportunity for enterprise adoption and expansion of AI use cases to accelerate, especially as agents/agentic tools become more capable, available, and usable. While AI consumption has expanded rapidly, participants described that a super-majority of AI consumption to-date has focused on coding applications.

订阅研报添加微信 LCD123321Cad 04-17

Quantum computing (QC) seems likely to hit an inflection in 2029 in terms of capability. Participants at the conference shared the opportunity to combine QC with AI models to create a synergistic benefit to combine AI analytics with QC simulations. QC could present another demand opportunity for HPC deployments and demand for IT Load that could accelerate in the future.

We remain positive on the data center category. The expansion of enterprise and associated inference workloads should be a positive for data center colocation demand, including retail and mid-size enterprise deployments, that bode well for Buy-rated Equinix and Buy-rated Digital Realty. As demand remains strong for available IT load, we also learned that it's possible for smaller scale data center facilities (<20 MW) to be connected together within a market to create a synthetic campus that could be helpful for longer-dated facility vintages.

Commentary was positive on the investments in fiber and networking that need to follow the expansion of data center infrastructure. Commentary was encouraging for Neutral-rated Lumen on the prospects to continue to leverage its fiber assets and new NaaS platform to improve its share and capture a portion of the new demand to connect AI infrastructure with customers and other data centers. Stronger demand for connectivity could also be helpful for Uniti, and the business wireline segments at AT&T and Verizon. Each of these telcos is facing a varying degree of legacy wireline revenue headwinds that are diluting the flow-through from strengthening demand from fiber services for AI workloads.

订阅研报添加微信 LCD123321Cad 04-17

Day One Takeaways

WEDNESDAY, APRIL 15

Thinking in Decades, Not Quarters: A Fireside Chat with Vinod Khosla

We hosted Vinod Khosla, Founder of Khosla Ventures, for a fireside chat on Day One of the conference. Khosla opened by arguing that most enterprises still fail to grasp the scale of change ahead, citing $2.5T in capex as unimaginable five years ago, and set a benchmark of doubling revenue per employee by 2030 for large enterprises, and 10x for small ones. He projected 50% unemployment or underemployment in the US by 2035, framing AI not as a tool given to people but as the organization itself, with people increasingly working for AI. He pointed to an existing two-employee company generating $1.8B in revenue as an early proof point.

On policy, Khosla proposed eliminating the capital gains tax, instituting a token tax by 2030 to fund displaced workers, and establishing a national sovereign wealth fund owning 20% of every US enterprise by 2035. He expects the cost of medical, legal, and educational services to collapse to below $1/hour with AI, and sees a reasonable chance that by 2030, humans will no longer need keyboards or mice, with Khosla Ventures running 4 to 5 projects in silent speech interfaces. He identified AI security — cyber wars and AI in the wrong hands — as the single biggest danger, warning that all RSA encryption algorithms are likely to be broken well before 2030, and assigned greater than 50% probability that political friction, not technical failure, prevents the transformations he described.

订阅研报添加微信 LCD123321Cad 04-17

AWS Keynote: Applied AI at Scale

We hosted Colleen Aubrey, Senior Vice President of Applied Solutions at AWS, for a keynote conversation moderated by Ric Spencer, Vice-Chair of Tech Investment Banking and Head of AI Infrastructure Banking at Citi. The session covered AWS's philosophy for deploying AI inside the enterprise, the evolution from productivity tools to agentic “AI teammates,” real-world implementation results across retail and healthcare, and the strategic imperative of identifying durable, complex problems as the true long-term opportunity in enterprise AI.

An operator’s perspective. Aubrey brought an operator's lens to the AI conversation, drawing on her 15+ years scaling Amazon's advertising business into a $30B–40B franchise. The throughline from ads to AI, she argued, is the same: value creation comes not from giving customers exactly what they ask for, but from building on their behalf and iterating through failure. The inflection point in both cases arrived when the product became compelling enough that customers tolerated its rough edges — a signal, she suggested, that the industry is now approaching with enterprise AI.

订阅研报添加微信 LCD123321Cad 04-17

Philosophy and approach to applied AI solutions. The centerpiece of Aubrey's remarks was AWS's "AI teammates" framework — a vision in which hundreds of thousands of task-level agents are orchestrated into functional teams, effectively making every employee a manager of AI. This represents a deliberate departure from productivity-tool thinking; the goal is not to automate individual tasks but to fundamentally restructure how work gets done, collapsing traditional role boundaries across marketing, engineering, and finance functions. Notably, AWS has chosen to avoid desktop metaphors and anthropomorphized interfaces, betting instead on new collaboration patterns native to agentic AI.

Underlying all of it is a strategic question that Aubrey framed as **the most important challenge in enterprise AI right now: identifying which problems are truly durable.** The shift AWS is driving — getting AI to work *in* the business while people work *on* the business — only creates lasting value if it is applied to problems that won't be rendered moot by the next wave of tooling. Intermediate workflows that can be vibecoded away in the near term are not the target; the real opportunity lies in complexity that compounds over time, where regulatory friction, institutional knowledge, and high-stakes outcomes make human-AI collaboration irreplaceable. For AWS, healthcare and drug discovery are early proof points of that thesis — sectors where the cost of being wrong is high enough that durable, trust-based AI systems will always have a role.

订阅研报请加微信 LCD12321Cod 04.17

Implementation case studies. Aubrey used two concrete deployments to ground the vision. The Tumi Group implementation via Amazon Connect — spanning 10,000 agents and 30 million customers — delivered a 30% decrease in contacts per booking and a 25% lift in upsell, illustrating the commercial impact of consolidating on a single platform rather than stitching together point solutions. The second example, Amazon Health Connect, takes a horizontal approach across the full patient journey — from appointment scheduling through clinical documentation and billing — with a single AI teammate maintaining context end-to-end rather than handing off between siloed tools.

Healthcare emerged as a near-term beachhead for AWS, and Aubrey was candid about why: the sector's complexity, regulatory requirements, and administrative burden create durable, hard-to-commoditize problems that reward deep technical investment. The demo highlighted a key design choice — the AI raises its hand to flag gaps rather than auto-populating fields, preserving clinician control while surfacing meaningful clinical insights. AWS sees regulatory moats, HIPAA-compliant infrastructure, and FDA navigation as long-term competitive advantages, with Amazon Biodiscovery signaling that the same agentic architecture is being extended into drug and antibody discovery.

订阅研报请加微信 LCD12321Cod 04.17

Five top insights from our discussion with IBM Research Chief Scientist Dr. Ruchir Puri

We hosted Ruchir Puri, Chief Scientist of IBM Research, for a fireside chat moderated by Fatima Boolani, Managing Director, Co-Head of US Software Equity Research.

1. Quantum slated for a "Y2K Moment" in 2029 meaning a necessary spend cycle is on the come. Every organization running traditional encryption (i.e. everyone) will need to migrate to post-quantum cryptography (PQC) before fault-tolerant quantum computers arrive (~2029) that are capable of breaking prevailing encryption standards/frameworks. Conceding some real engineering challenges still persist and the hard-and-fast timeline has potential to slightly shift, the migration/adoption impetus remains. IBM already sells scanning/remediation tools today, and the vendor ecosystem is supportively building out PQC suites/productized expertise now, which stands to result in the type of compliance-driven spending upswing seen during Y2K. We think Anthropic's Mythos' model preview showcasing incredibly higher levels of performance and aptitude

订阅研报请加微信 LCD12321Cod 04.17

2. Quantum unlikely to displace classical computing/AI infra, poised to be more TAM-additive but revs opportunity still nebulous. Because quantum handles a specific class of problems (sampling, optimization, simulation) and is capable of working in concert with AI, it is unlikely to supplant classical IT architectures which

these topologies can't/won't efficiently solve. As such Quantum is likely to remain incremental rather than cannibalistic infrastructure spend, but perhaps narrower to start as mainstream commercialization momentum is still in early stages.

3. Pursuing a full-stack vs. "boutique" or piecemeal solution strategy. A vertically-integrated platform strategy (hardware → Qiskit SDK → compilers → functional libraries → natural-language coding interfaces) is an important differentiation and a conduit for mass adoption, analogous to what IBM achieved with its System/360 mainframe franchise launch 50+ years ago. ~200 quantum systems deployed globally and ~200 commercial partners (+ industry credibility from Nobel and Turing research accolades) are iteratively informing product/GTM strategy, alongside IBM's continuous efforts in building ecosystem breadth.

4. Life sciences/materials science use cases (drug discovery) appear to have the most commercial legs so far. While customer adoption commentary still skewed more experimental rather than widescale production-focused, these were cited as the most salient/promising near-term areas of commercial success thus far.

5. "Harvest Now, Decrypt Later" risk reinforces the catalyst to invest in Quantum infra as cyber risk ramps, but critical industries likely move first. Whilst broader enterprise spend inflection around PQC tools likely requires more visible catalysts, the opportunity is poised to be more tangible for security-conscious "critical" verticals (financial services, defense, healthcare).

Revolutionizing Software Development

We hosted Matan Grinberg, CEO of Factory; Jeanne Grosser, COO of Vercel; and Jyoti Bansal, CEO of Harness, for a panel on Revolutionizing Software Development. Bansal noted that only 30 to 40% of developer time is spent in the inner loop of engineering, with AI increasing code volume 3x to 10x and making the outer loop — DevOps, app security, FinOps — the critical bottleneck. Grinberg agreed that demand for software is absurdly high and AI is not deflationary in the short or medium term, while Grosser added that Vercel's fluid compute model charges only for specific tasks to address the cost concerns of running agents at scale. On multi-model strategy, Grinberg stressed that model agnosticism is critical for resilience and cost optimization, with Grosser noting enterprises should build stacks that allow switching on a weekly basis.

On build vs. buy, Grosser described a GTM engineering org of five engineers at Vercel that replaced a 10-person sales development team, at a cost of $1,000 per year with every KPI improving, and a three-person team that built a support tool now handling 97% of requests across 215,000 companies. Bansal warned of an explosion of vulnerabilities as code volume becomes infinite, calling for automated guardrails with clearly defined rules governing agent credentials. Grinberg acknowledged vibe-coding is a fantastic development but warned of a security hangover, noting the industry will have to do the unglamorous work of resolving it.

Lumentum: The Optical Backbone of AI Infrastructure

We hosted Michael Hurlston, CEO and President of Lumentum, for a fireside chat. The defining event of Hurlston's early tenure was a $2B strategic investment from NVIDIA, matched by a parallel $2B into Coherent, securing a multi-year R&D collaboration on co-packaged optics (CPO) and locking up a meaningful portion of Lumentum's supply for NVIDIA's roadmap. This triggered a "land rush" among hyperscalers scrambling to secure their own indium phosphide supply, with supply projected to trail demand through 2028 to 2029. The real inflection arrives at the

3.2T node in late 2028 or early 2029, when copper's effective reach drops to 1.5 to 2 meters and optics become the prevalent backplane solution.

Hurlston articulated a second derivative growth dynamic: the company's growth is not simply a function of how many data centers get built, but of increasing optical intensity within each one. Three independent drivers compound on top of the broader buildout — CPO solutions replacing traditional transceivers, optical backplanes creating an entirely new TAM inside the rack, and Optical Circuit Switching (OCS) for GPU failover, with significant growth expected in 2027. Financially, Lumentum is targeting $1.5 to $2B in quarterly revenue, and for the first time in the optics industry's history, suppliers are able to raise prices and pass through cost increases, a structural shift broadly bullish for Lumentum, Coherent, and Broadcom alike.

订阅研报添加微信 LCD12321Cad 04-17

Defining the Frontier: How Foundational Models Are Evolving

We hosted Aakanksha Chowdhery, Research Lead at Reflection AI; Ori Goshen, Co-CEO and Co-Founder of AI21; and Ren Ito, Co-Founder and COO of Sakana AI, for a panel on how foundational models are evolving. All three panelists agreed that overall employment will be higher in five years, though job profiles will change materially, with bankers becoming product managers, engineers sitting next to trading desks, and designers becoming more engineer-like. Goshen described AI21's Jamba architecture, joining Transformer and Mamba to achieve 3x to 5x lower token costs and higher throughput, and expects the GPU-to-CPU ratio to reach 1:1 as deterministic computation demand grows. Ito noted that Sakana produced a model through model merge methodology in March 2024 that performed nearly on par with GPT-3.5 for under $1M.

订阅研报添加微信 LCD12321Cad 04-17

On enterprise ROI, Goshen identified three key bottlenecks: evaluations (very hard to get right, requiring continuous automated testing), data (not just cleaning but making it ready for agentic systems), and people (willingness to give AI systems autonomy remains a challenging barrier). Chowdhery noted that open models are increasingly closing the gap with closed models for enterprises seeking ownership, auditability, and efficiency, while Goshen stated that open source will be the dominant method to activate and deploy production AI systems. Ito argued that model layer companies need to do more integration into the application layer, and that Sovereign AI is about proximity to local companies, not building another trillion-parameter model from scratch.

DigitalOean: Building the Inference Factory

We hosted Paddy Srinivasan, CEO of DigitalOcean, for a fireside chat moderated by Mark Zhang, Senior Research Associate of US Software Equity Research. Inference opportunities continue to gain traction with healthy demand from AI native customers and the use cases/workloads include 1. viable businesses that are injecting a heavy dose of intelligence into existing SaaS; 2. those that are reinventing a category with an AI-centric approach (e.g. next-gen Concur using AI native approach); and 3. agent-native companies that are emerging and taking an approach to reduce human involvement in certain domains. All three categories are intelligence hungry and require meaningful amount of inferencing. As a result, the business model is undergoing a fundamental shift from charging for inputs (e.g. GPU per dollar/hour) to charging for output, which is token and quality of tokens.

订阅研报添加微信 LCD12321Cad 04-17

DigitalOcean has built a five-layer inference stack – infrastructure, compute primitives, reinvented data layer, inference engine with extensive model support, and agent lifecycle capabilities (e.g. persistent memory and sandboxing) – creating a meaningful moat and attractive economics vs the neoclouds as well as the

Building AI Factories

We hosted William Falcon, Founder & CEO of Lightning AI; Nitin Agrawal, CFO of Coreweave; and Dmytro Dzulgakov, Co-Founder and CTO of Fireworks AI, for a panel on Building AI Factories. Panelists discussed compute needs are insatiable, with every customer unable to get enough compute through 2026, 2027, and even 2028. Agrawal described demand increasing across Coreweave's entire customer base, with the proliferation of inference use cases accelerating demand for both newest and prior generations of infrastructure. Falcon characterized the shift from 1:1 human-to-GPU usage to a world where a single business user building a million-agent deployment pulls 10x the endpoints, implying 100x growth in inference demand and a 1,000x need for compute.

订阅研报添加微信 LCD123321Cad 04-17

On bottlenecks, Agrawal described the shortage of power shell capacity as the key bottleneck, with Coreweave having secured 3 GW of contracted power. On depreciation cycles, Agrawal noted Hopper ASPs are maintaining price points within 10% of original contract at renewals, and Coreweave noted the inference explosion has revealed memory as a growing issue as workloads become more memory-intensive. The panel agreed on a multi-vendor silicon future, with Google, AMD, and AWS all expected to play significant roles alongside NVIDIA.

Replit: The Democratization of Software Creation

We hosted Amjad Masad, CEO of Replit, for a fireside chat moderated by Ronit Ghose, Global Head of the Future of Finance Team at the Citi Institute. Replit has grown its enterprise business 20x in the past year, with commercial usage now representing 40 to 50% of its customer base. Enterprise customers now include Meta, Adobe, PayPal, Visa, Plaid, Amex, SoFi, Zillow, and Sears Home Services, with accounts starting at $20,000 and scaling to millions. Masad emphasized that vibe coding is more serious than it sounds, with customers building mission-critical software with full specifications, test environments, and deployment pipelines, and that the next evolution moves toward intent-based directives where an agent executes autonomously over days or weeks.

订阅研报添加微信 LCD123321Cad 04-17

Masad framed AI as enabling a new class of accidental entrepreneur, citing an investment banker who built a hallucination-resistant deck tool overnight, launched a company, and raised at a $30M valuation within a month. Replit is model-agnostic by design, routing tasks to Claude for coding agents, GPT for math and data work, and Gemini for design, and currently ranks first on the Design Arena benchmark despite building on others' models. Where foundation model companies position themselves as utilities selling tokens, Replit's enterprise sales motion operates more like transformation consulting, focused on outcomes rather than consumption.

Keynote: Google Cloud

订阅研报添加微信 LCD123321Cad 04-17

We hosted Francis deSouza, COO of Google Cloud, for a keynote, moderated by Citi CTO David Griffiths. DeSouza, who is also leading the integration of Wiz into Google and driving Google's AI and agent security model, discussed the evolving cybersecurity threat landscape, the strategic rationale for the Wiz acquisition, Google's model strategy, and real-world AI ROI across Google's own operations.

AI-driven transformation of cybersecurity. DeSouza described a generational refactoring of the IT stack, with cybersecurity being completely transformed by AI across two dimensions: who is in the threat landscape and what attack techniques are being used. He noted new types of adversaries including criminal organizations and advanced entities leveraging fully agentic attacks and stated that the time for

a criminal organization to get into an enterprise environment has dropped from eight hours to 20 seconds. Code attacks went from a small fraction of incidents at the beginning of 2025 to the largest vector of attacks by year-end, per deSouza. He also cited data poisoning, drift, and distillation attacks as emerging threats requiring agentic infrastructure to be protected on its own, all playing out against a regulatory environment that is more unforgiving than ever.

Google's competitive advantage in AI security. DeSouza argued that Google is the only security company with a full AI stack, co-engineering its security portfolio with the model from day one. He noted that any other security company can only build security products after a model is released — a six-month to one-year lag — meaning the adversary has access to the model on day one while the defender could be two generations behind. Google's approach also enables power efficiency gains through co-engineering of the stack. On the Wiz acquisition, deSouza explained that as customers have migrated data to the cloud, the center of gravity for security has moved from endpoint to datacenter to cloud, and Wiz — built from the ground up to be multi-network — is well-positioned to be the provider of AI security. Google announced an AI application offering code-to-cloud-to-runtime protection for AI infrastructure, and at RSA announced a suite of agents — Red Agent, Blue Agent, and Green Agent — to serve as part of customers' security teams, with the vision that most cybersecurity work will be done by agents.

Model strategy and Gemini. DeSouza stated that Gemini was always built from the ground up to be multimodal and will continue to be so, with specific models for security. He noted that while Google initially anticipated many more specialized models, the core model has proven so powerful that the general-purpose Gemini model will be sufficient for the vast majority of cases, with continued support for specific domains such as mathematical use cases and protein coding — though fewer than originally anticipated. He acknowledged that competition for scarce resources between DeepMind, Search, and Google Cloud makes it difficult to bring all projects to market.

Internal AI ROI and data advantage. DeSouza cited Google's own security operations stack as the area of biggest return, noting that AI and agents reduced attack detection time by over 90% — already measured in minutes — and that this capability is being productized into the Red/Blue/Green agent offering. Internally, Google uses agentic workflows for website-driven leads, achieving a 15% increase in conversion rates for sales teams due to warmer leads. The Chrome team has moved to using AI models to generate collateral in over 15 languages across over 50 countries. Google's treasury team used Gemini to improve the accuracy of balance predictions, generating tens of millions of dollars in improvements, starting in the US and now rolling out globally. DeSouza highlighted Google's data advantage — protecting 4B devices and 1.5B inboxes — as a key differentiating factor, noting that Google sees the kill chains and tools for the most sophisticated attacks in the world, which directly informs the defenses it can build.

Quantum and encryption risk. DeSouza confirmed that quantum will be able to break some existing encryption standards and that organizations need to move to encryption standards that cannot be broken by quantum. He warned of the risk of adversaries harvesting data now and waiting until 2029 and stated that CISOs need to be implementing more quantum-resistant protocols today.

Anthropic Keynote: Enterprise AI at an Inflection Point

We hosted Scott White, Head of Enterprise Product at Anthropic, for a keynote session at the 2026 Citi AI Summit moderated by Heath Terry, Global Head of

Technology and Communications Investment Research. The session covered the state of enterprise AI adoption, the evolution from productivity tools to revenue-generating workflows, Anthropic's approach to model safety and capability development, its multi-cloud strategy, and the organizational and integration challenges that define large-scale enterprise deployment.

2026 is the year enterprise AI evolves beyond pilots. White drew a clear arc of progression: AI began with information retrieval, moved to task assistance, and has now reached full workflow delegation — where Claude is being used to run entire work streams end-to-end. Financial services, life sciences, and legal are the sectors seeing the fastest and deepest adoption, and the defining characteristic of this moment is scale: deployments that were experiments a year ago are now core to how work gets done. The cultural and organizational shift required to capture that value, White argued, remains the primary obstacle for most enterprises.

Opportunity in revenue creation, not just cost reduction. White pushed back on the framing of AI as primarily a cost efficiency tool, pointing to a beverage manufacturer that built a demand analysis capability on top of Claude and has since launched an entirely new analytics service line. An investment banking analyst using Claude for DCF modeling is not just faster — they are building more comprehensive models than were previously possible. The implication for enterprises is significant: the question is no longer how to do existing work cheaper, but what new work AI makes possible for the first time. White urged a "deploy, measure, manage" discipline rather than passive rollout, and reframed the cost conversation around a single metric: cost per outcome, not cost per token.

Mythos represents a capability step-change — and a deliberate safety pause. Anthropic's most capable model to date, Mythos, is currently in limited enterprise preview through the Project Glasswing consortium rather than broad release. White was candid about why: greater capability inherently carries greater safety risk, particularly around cybersecurity and CBRN concerns, and Anthropic chose not to release it broadly until those risks were fully understood. The model will be larger and more expensive than its predecessors, with raw token costs moving from $75/M tokens for Opus 4.0 to $125/M for Mythos — but White's framing was consistent: the right measure is what outcomes those tokens produce, not what they cost.

Safety is a strategy, not a constraint. Anthropic's mission-driven approach to safety is increasingly a commercial differentiator as enterprises grapple with data leakage, regulatory compliance, and reputational risk at scale. White cited the deliberate delay of browser use capabilities until prompt injection risks were fully addressed as an example of the rigor that builds enterprise trust over time. The Model Context Protocol (MCP) — described as a USB-C for AI integration — reflects the same philosophy applied to enterprise complexity: a standardized integration layer that works across the heterogeneous reality of large organizations, which White characterized as more like countries than uniform actors, with internal factions, legacy systems, and competing priorities that shape how AI actually gets adopted. Anthropic's multi-cloud availability across AWS, Google Vertex, and Microsoft Azure reinforces this posture — positioning Claude as a horizontal platform that meets enterprises where they are rather than competing to own the full stack.

Allocating the AI Future: Where Growth Equity Is Going Next

We hosted Paul Abrahimzadeh, Managing Partner of 1789 Ventures; Abhishek Shukla, Managing Director of Prosperity7 Ventures (backed by Saudi Aramco);

Alexander Harstrick, Managing Partner of J2 Ventures; and Emma Norchet of T. Rowe Price, for a panel on where growth equity is going next in AI. Norchet described backing model companies including OpenAI and Anthropic, participating in $170B and $350B rounds, as well as infrastructure companies including Together AI and others, with increasing time spent on the application layer, citing Harvey AI as a recent investment and a Series B in Physical Intelligence. Abrahimzadeh highlighted Reflection AI, noting its model release schedule for Q3 and a 100,000 Blackwell GPU allocation from NVIDIA as a potential game changer for open vs. closed weights.

订阅研报添加微信 LGD123321Cad 04-17

On sovereign AI, Abrahimzadeh stated he is a big believer in open source for enterprises wanting more control over the data layer, noting that many Silicon Valley applications are built on top of Chinese models, which carries data risks. Harstrick identified physical AI and telecom infrastructure as the next physical manifestation of AI, noting the literal bottleneck is in physical buildings, with companies like American Tower and Crown Castle that have not done a refresh. On infrastructure, Abrahimzadeh highlighted optical interconnects and an investment in IR Labs, while Harstrick cited Aalyria Technologies in pre-space optics.

Securing AI Panel Perspectives

We hosted a panel on AI security moderated by Nick Sands, Director of Venture Investing at Citi Ventures, featuring Vibhav Sreekanti, Co-Founder and CTO of Prophet Security; Manoj Saxena, Chairman of Trustwise; and Bo Li, CEO of Virtue AI. The panel drew a sharp distinction between two rapidly diverging markets: using AI to improve security, and securing AI systems themselves, with the second identified as the most underappreciated by the investment community. Prophet Security expects that Mythos-class capabilities will be broadly available, including in open-source form, within six to nine months, at which point the cost of sophisticated cyberattacks drops to near zero, with models capable of discovering novel vulnerabilities, chaining zero-day exploits autonomously, and operating "live hacking" campaigns over weeks. Trustwise was direct: an explosion of breaches is expected over the next 12 months.

订阅研报添加微信 LGD123321Cad 04-17

The panel argued that the entire security architecture must evolve from defending perimeters to controlling AI behavior from the inside out, in real time, with the critical missing infrastructure being runtime policy enforcement operating at millisecond speed at the kernel level. Virtue AI pushed further, arguing that using models to protect models is fundamentally unreliable: security requires determinism and guarantees, not probabilistic outputs, and the company is developing a "Shield Agent" built on logical inference rather than LLM judgment. For investors, the shift from detection to control favors new entrants over incumbents built around perimeter defense, and the ability to provide deterministic enforcement will be the defining technical differentiator in regulated industries.

订阅研报添加微信 LGD123321Cad 04-17

Intelligent Finance Panel

We hosted Gabe Stengel, CEO of Rogo; Rafael Loureiro, CEO of Wealth.com; Patrick Vergara, CEO of Norm AI; Parker Ence, CEO of Jump; and George Sivulka, CEO of Hebbia, for a panel on Intelligent Finance on Day One of the conference. Stengel identified three sources of differentiation for Rogo: model agnosticism, the extreme difficulty of creating last-mile outputs for end users, and cross-border deployment models. Loureiro noted that wealth management requires deterministic answers, as 90 to 95% accuracy is not good enough, while Vergara stated that Norm AI's differentiation has shifted from reliability to the proprietary data and nuance it has built in complex legal workflows. Sivulka acknowledged that Claude is likely to eat the wide majority of AI for finance companies, and that

Hebbia's differentiation lies in a separate deployment model that reinvents workflows from the ground up, focusing on institutional intelligence rather than individual productivity tools.

On the regulatory frontier, Sivulka described deterministic agents being used by asset managers to look at all private credit agreements in a portfolio, assess business risk, and originate off-market opportunities, finding positions that would not have been identifiable before given credit documents running to thousands of more capable, available, and usable. While AI consumption participants described that a super-majority of AI consump alls at the junior banker level, and that leading firms are mapping every process and line of business to find what is bottlenecked by people versus technology. On the long-term outlook, Stengel predicted a proliferation of high finance services for every individual within 10 years, analogous to how Rocket Mortgage and other firms have democratized mortgage and investing.

订阅研报请加微信 UCD123216 or1 01.17

Decagon: Redefining Customer Service

We hosted Jesse Zhang, CEO of Decagon, for a fireside chat moderated by Charlie Black, Head of Technology ECM and West Coast ECM. Decagon deploys a unified AI agent across phone, chat, and email with a single memory layer, replacing the fragmented, siloed chatbot experience that has defined enterprise customer service for the past decade. Agents are built and maintained entirely through natural language in-product, enabling non-technical teams to iterate without engineering resources, and the company measures success on two metrics: automation rate and customer satisfaction scores. Zhang highlighted that getting an AI agent live is 20% of the work; the other 80% is ongoing maintenance, iteration, and procedure management as business conditions change, which is the core reason enterprises are unlikely to build and sustain this capability in-house.

订阅研报添加微信 UCD123216 or1 01.17

Decagon is not a model company by design, but voice AI forced its hand: frontier models are too slow for the sub-second response times that voice agents require, so Decagon post-trains smaller, specialized open-source models optimized purely for customer service contexts. Zhang argued that value in AI accrues at the point of customer interaction, not in the underlying commodity, positioning the application layer as the most defensible place to compete. Decagon's enterprise customer base spans banks, airlines including Delta, telecom companies, and car rental companies including Avis, reflecting the breadth of the opportunity across a global call center industry estimated at $400B, and Zhang expects the market to consolidate to two or three winners, with execution, not technology, as the differentiating factor.

Vertical Intelligence Panel

We hosted Ari Tuchman, CEO of Quantifind; Sanjay Dhawan, CEO of SymphonyAI; Mike Gadd, CEO of Kindred; and Wolfgang Boecking, CEO of Vitelis, for a panel on Vertical Intelligence, moderated by Tyler Radke, Managing Director, Co-Head of US Software Equity Research. Tuchman described the enterprise AI stack as consisting of a data level, a risk and signal extraction layer, an AI middle layer that fuses signals and builds network graphs, and an agentic layer, arguing that as foundation layers get better, the middle layer becomes more important. Dhawan noted that SymphonyAI builds its own small reasoning models specific to use cases such as retail merchandising, asset performance, and financial crime, with a typical Tier 1 bank spending $200M to $500M per year on financial crime operations, and SymphonyAI's pricing model shifting to 80% of savings going to the customer and 20% to SymphonyAI.

订阅研报请加微信 UCD123216 or1 01.17

On ROI, Tuchman stated that a year ago the focus was cost savings, but now it is revenue, citing a large crypto exchange in Europe where Quantifind eliminated 98% to 99% of false positives, netting 10,000 new customers per month. Dhawan noted a recent competitive win at Visa where AI enabled real-time dynamic solutions to onboard customers and monitor risks, replacing static solutions to dynamic problems, and cited work with Walmart and Albertsons to drive improved merchandising performance. Boecking described finding $80M in additional pipeline for one telco customer by identifying competitive gaps across a product catalog of up to 200 products that no human could process in full.

订阅研报添加微信 LGD123321Cad 04.17

Salesforce: Agents, Outcomes and the Enterprise

We hosted Mark Sullivan, President of Regulated Industries at Salesforce, for a fireside chat moderated by David Livingstone, Chief Client Officer. Sullivan framed Salesforce's transformation as a fundamental expansion of scope, from managing customer data to actively participating in business systems and driving measurable results, and pushed back on the SaaSpocalypse narrative, arguing that an outcome-centric model benefits Salesforce rather than threatens it. Salesforce's path to its current pricing model took three iterations: seat licenses plus credits created friction, pure credit consumption transferred metering burden to ill-equipped customers, and the current model — unlimited agentic enterprise agreements where Salesforce absorbs the consumption risk — has become the dominant buying structure.

Life sciences has been the most aggressive adopter, with Salesforce completing 160-plus transactions with meaningful penetration into the top ten pharma companies. Financial services started from caution, as liability concerns mean 90% accuracy is not an acceptable threshold, but is beginning to graduate toward revenue-generating applications. Sullivan was clear that Salesforce is not competing in the LLM race, with models commoditizing and Salesforce's architecture deliberately model-agnostic, and that the real strategic positioning is in deterministic orchestration: the ability to register, manage, and coordinate agents reliably across complex enterprise environments. With 5-to-10-year contracts and three product releases per year, the upsell runway is significant, though customer readiness, not product capability, is the gating factor.

订阅研报添加微信 LGD123321Cad 04.17

The Rise of Voice AI: Real-Time Intelligence in Action

We hosted Nikola Mrksic, CEO of PolyAI; Scott Stephenson, CEO of Deepgram; Ping Wu, CEO of Cresta; and Raj Koneru, CEO of Kore.AI, for a panel on the Rise of Voice AI, moderated by Ron Josey, Managing Director & Senior Analyst, Internet Equity Research. The panel was aligned that voice AI is technically hard because it is real time: voice is higher bandwidth than chat, people are very picky about voice quality, most companies with voice have legacy systems making it harder for agents to take actions, and voice is usually the final point of escalation for customers, meaning people do not experiment easily. Mrksic identified trustworthiness as the number-one barrier to enterprise adoption, noting that for five customers, doing the work of 5,000 FTEs took time, and that he does not expect an Anthropic-equivalent moment for enterprise voice AI adoption this year. Stephenson cited CVS as a customer advised to first tackle prescription readiness inquiries, accounting for 30% of all calls, rather than attempting everything at once.

订阅研报添加微信 LGD123321Cad 04.17

On the long-term market, Stephenson stated that while some believe the voice AI market is $50B to $200B, he believes it will be a $1T market, with 8 billion people and an expectation of 50 billion voice agents. Koneru predicted that in three to four years, every application will have voice AI, and Wu noted that SMBs will be the

biggest beneficiaries. The panel agreed that LLMs help but are not the answer to voice AI on their own, and that what is needed is encapsulating the right control flow and providing application logic that can be trusted, with the best AI agent trained on humans and designed so that handoffs do not require customers to start from scratch.

Pegasystems: Determinism Over Autonomy

订阅研报添加微信 LCD123321Cad 04-17

We hosted Alan Trefler, Founder and CEO of Pegasystems, for a fireside chat moderated by Steve Enders, Vice President of US Software Equity Research. Trefler made a deliberate choice to reject the dominant approach of creating agents, handing them a prompt, and letting them execute independently, instead building Blueprint, a system that harnesses GenAI to identify best practices and design optimal workflows while leaving execution to Pegasystems' deterministic workflow engine. Before Blueprint, Pega engagements began with sales teams spending up to two weeks at whiteboards; today, customers upload business documents and Blueprint designs a tailored workflow using vector databases and accumulated knowledge, compressing timelines from months to days or weeks.

Trefler was candid that customer adoption of fully agentic capabilities remains in early stages, held back by a simple reality: most businesses need operations to work consistently regardless of when or how they are triggered, and in many cases there are legal and regulatory implications to variability. For high-volume processing areas like credit card chargebacks, Pega does implement agentic automation that skips human steps entirely, but within tightly defined and governed parameters. Trefler's core argument to investors was that most businesses are, at their core, a collection of processes, and a workflow engine is how those processes get executed reliably at scale, with vibe-coded applications built without a mature workflow engine underneath likely to prove unsustainable over time as complexity compounds.

订阅研报添加微信 LCD123321Cad 04-17

The Intelligent Enterprise

We hosted Nicolas Kopp, CEO of Rillet; John Glasgow, CEO of Campfire; Debanjan Saha, CEO of DataRobot; and Dan Springer, CEO of Ironclad, for a panel on the Intelligent Enterprise, moderated by Tyler Radke, Managing Director, Co-Head of US Software Equity Research. Kopp stated that for some workflows, AI is quite existential, particularly for reporting, analytics, and point-solution product management components of the stack, where vibe-coding can produce a solution in one to two days with little to no maintenance. Glasgow stated that moats are starting to chip away for large incumbents, and that the old model of seat-based pricing with a year of implementation and a large maintenance team will go away. Springer argued that the short-term impact of AI on SaaS is overestimated while the long-term impact is underestimated, analogous to the cloud transition, and that some companies will make the transformation and some will not.

订阅研报添加微信 LCD123321Cad 04-17

On headcount optimization, Glasgow cited a customer that achieved a 40% reduction in its accounting team after going live with Campfire, and noted that accounting is seeing a revival as the nature of the work changes. Glasgow noted that tokens are eating budget more from human capital, a materially larger budget than IT, than from IT spend, and predicted a progression from spot pricing to large prepayments, as seen with AWS in its early days, with a cost optimization wave coming once marginal advancements flatten and open-source models proliferate. Saha identified services, SDRs, and BDR-type roles as the areas most impacted by AI-driven headcount optimization.

Perplexity: Search, Trust, and the Anti-Ad Model

We hosted Daniel Kwak, Senior Vice President of Finance at Perplexity, for a fireside chat moderated by Heath Terry, Global Head of Technology and Communications Investment Research. The session covered Perplexity's product philosophy around the browser as the primary AI interface, the company's move into specialized professional workflows, and its deliberate decision to walk away from advertising — a strategic choice with significant implications for how AI-differentiation and a conduit for mass adoption, analogous with its System/360 mainframe franchise launch 50+ years systems deployed globally and ~200 commercial partners (

订阅研报添加微信 LCD123321Cad 04-17

The browser is the operating system, a core thesis. Perplexity's foundational product thesis is that the browser has effectively become the operating system for knowledge work, making it the critical surface for AI-powered information retrieval. The company stays unusually close to actual user query patterns — using real behavior rather than top-down feature planning to drive product decisions — and has built a multi-model architecture where each subtask within a query routes to a different model and continues running in the background until completion. This orchestration approach, optimizing for quality across different task types within a single query, reflects the same multi-model reality that surfaced across multiple sessions at the conference: no serious workload runs on a single model.

The next frontier is specialized professional workflows. Kwak leads development of Perplexity's professional finance product within Computer, the company's agentic product, and was clear that the direction is deliberately vertical rather than horizontal. The focus is on deeply understanding specific workflows before building automation — a significant investment of time and effort that precedes any one-click feature — with division-specific automations emerging as a major area of development, recognizing that workflow needs vary materially across teams within the same organization. The move positions Perplexity as a potential competitor to vertical AI workflow tools like Hebbia, though arriving from a search-first starting point rather than a workflow-first one.

订阅研报添加微信 LCD123321Cad 04-17

Walking away from advertising. Perplexity's decision to abandon impression-based advertising was framed not as a business model choice but as a philosophical one: advertising is structurally incompatible with maximal helpfulness, and the tension between serving ads and serving users cannot be resolved in the user's favor. Kwak put it plainly — "I can't get a post-earnings summary that also tries to sell me something" — and the company has concluded that trust, once compromised by advertising incentives, is not recoverable. The implicit signal: Perplexity sees its future in enterprise and subscription monetization, validating a path for AI-powered information retrieval that does not depend on the ad-supported model that has defined search for two decades and that Google's business is built upon.

Keynote: Cursor

订阅研报添加微信 LCD123321Cad 04-17

We hosted Michael Truell, CEO and Co-founder of Cursor, for a keynote, moderated by Tim Ryan, Citi's Head of Technology and Business Enablement. Cursor, founded in January 2022, builds AI-powered software development tools and is now used by over 60% of the Fortune 500, including PwC, Starbucks, and financial services customers. The session covered the evolution of AI coding from assistive tool to autonomous agent, Cursor's competitive positioning, the future of the software development lifecycle, and what the future engineer looks like.

The arc of AI coding and the agent transition. Truell described an amazingly fast and disruptive arc in the AI coding market. In 2022, Cursor was adopted by individual engineers at $10 per person per month, delivering 5–10% speed

improvements. He stated that the company is only about 20% of the way along the curve toward AI coding becoming critical infrastructure, with tech leaders expected to allocate tens of percentage points of R&D spend to AI labor. Currently, 80% of code among Cursor's enterprise customers is being assisted by AI, with users running 2–3 agents at a time. Truell described the next development as coding agents that can work like full colleagues — running for hours and days with their own compute — and stated that over the next 12 months, local agents will shift to backplanes creating an entirely new TAM inside the rack, and colleague. Switching (OCS) for GPU failover, with significant growth e Financially, Jamin pointed to $1.5B in quarterly $2B in quarterly

订阅研报请加微信 ECD123321CAI 研报

but Cursor made an explicit trade-off to work in a domain with many competitors because the market is large — $1T is spent on software, and the shift from all labor to AI represents a market opportunity that will rival that of cloud. He noted that the competitive landscape has consolidated, with Microsoft as the first competitor followed by many startups, and that Microsoft has struggled to make its tools better and faster. Over 60% of the Fortune 500 pays Cursor today. He noted that people are thinking about the AI coding space strategically like the cloud space — wanting access to the cheapest models and avoiding lock-in to one provider — and that over the past month, models from a competing LLM provider have become 2–4x cheaper for coding tasks than the previously dominant player.

The future of the SDLC and the engineer's role. Truell described two main axes for Cursor's product: autonomy (from assistive to autonomous to remotely operated agents) and a second dimension he is focused on around verification. He expects the ratio of generative to verification code to invert over the next year, with most tokens spent on verification rather than generation. He noted that PR review pages are getting antiquated and that the UI and surface for reviewing code is starting to change. On the future engineer, Truell stated that being a steward of the codebase will remain important for a long time, and that over time the role will evolve into a blended product-engineering-design role focused on how the product should look and what customers want. He also noted that non-coding positions are increasingly using coding tools, with proliferation of internal tools as accessibility increases and costs come down.

订阅研报请加微信 ECD123321CAI 研报

ROI measurement and the next phase of deployment. Truell described the current moment as interesting for leaders, noting that AI adoption requires changing how teams work and changing behaviors, and that there are very few standard metrics for measuring efficacy. Cursor is working to give customers visibility into token spending with admin controls, and is providing access to statistics such as how many bugs came from AI agents and how much code was produced by AI. He described working on a holy grail ROI metric and wanting to flag low-ROI use cases of spend. On Cursor's own internal use, Truell described the next phase of AI coding deployment as Factory — building a factory of software engineering where coding agents go end to end, with bug tracking agents firing automatically when a bug comes in. He announced Cursor 3, the next version of Cursor, entirely focused on coding agents for both local development and cloud agents.

订阅研报请加微信 ECD123321CAI 研报

Day Two Takeaways

THURSDAY, APRIL 16

OpenAI Keynote: From Hype to Habit

We hosted Ashley Kramer, Chief Revenue and Marketing Officer at OpenAI, for a keynote conversation moderated by Vis Raghavan. The session covered OpenAI's enterprise growth trajectory, its three core competitive differentiators, the evolution from chat to reasoning to agentic AI, and the company's vision for deep enterprise partnerships that go beyond automation to full business transformation.

订阅研报添加微信 LCD123321Cad 04-17

Enterprise adoption has reached a pivotal moment. Kramer opened by noting that the conversation has fundamentally shifted — the key question is how to move from pilot to production at scale. OpenAI was initially perceived as a consumer company, even though its enterprise offerings launched concurrently, and that perception is now correcting: enterprise represents approximately 40% of the business today, with a stated target of reaching a 50/50 consumer-enterprise split by year end. Kramer's own signal for measuring the pace of change is direct customer engagement — she meets with more than 20 enterprise customers per week to understand what is working, what is not, and whether the ROI enterprises expected when they began their AI journeys is materializing. The metric that matters, she was clear, is outcomes.

The $122B funding round is fundamentally about compute, reliability, and the ability to scale. Kramer framed the capital raise not as a valuation story but as an infrastructure imperative: greater compute capacity translates directly into system reliability, and reliability is the prerequisite for enterprise trust at scale. Massive scale also enables faster model iteration at lower marginal cost, creating a compounding advantage that sustains innovation velocity over time. The research-led foundation that underpins OpenAI's model development is what Kramer identified as the long-term differentiator — an advantage that compounds as the pace of capability development accelerates.

订阅研报添加微信 LCD123321Cad 04-17

The evolution from chat to agentic AI defines the current moment — and the value gap in between. Kramer traced a clear arc: ChatGPT launched in November 2022 as a way to understand how the general public would use AI, reached 100M users in two months, and demonstrated an appetite that reshaped the industry. The first wave was conversational AI; the second was reasoning; the third — and current — wave is agentic AI, where systems do work on behalf of users rather than simply informing them. To help enterprises bridge the gap between model capability and realized impact, OpenAI launched its Frontier Platform (February 2026) — a context layer that sits above existing systems of record to enable reliable agent development, monitor agent behavior in production, and help organizations adopt AI in the right direction at the right pace.

订阅研报添加微信 LCD123321Cad 04-17

Three differentiators set OpenAI apart, and the partnership model is where they converge. Kramer identified compute capacity, consumer-scale insight, and partnership depth as OpenAI's core advantages. The 900M+ ChatGPT WAUs create a familiarity and comfort level that translates directly into enterprise adoption, giving OpenAI a consumer-to-enterprise flywheel difficult to replicate. The Target partnership (November 2025) illustrates the breadth of what deep engagement looks like in practice — spanning consumer-facing outfit recommendations, in-store manager enablement, and trend AI that compresses fashion decision cycles. BBVA (December 2025) represents the long-term version of this model: 120,000 employees across 25 countries, powered by Codex and

ChatGPT Enterprise, with joint investment, shared milestones, and a stated goal of becoming a fully AI-native bank.

On froth, safety, and societal impact. Kramer stated that AI has moved beyond hype and into habit — durable value emerges when AI becomes embedded in everyday workflows, and that is where the focus belongs. On safety, Kramer was equally clear: OpenAI's commitment is to release models only when they are safe, secure, and ready, and she acknowledged that not all players in the market adhere to the same bar — a point that carries weight as frontier model capabilities accelerate. On workforce impact, the message was consistent with what was heard across the conference: AI should eliminate the mundane to make room for creative problem solving, and the human elements of business — trust, relationships, judgment — remain irreplaceable. A job certification program developed with Walmart HR reflects OpenAI's active effort to help organizations and their employees embrace that transition rather than resist it.

订阅研报添加微信 [illegible]

Fireside Chat with Adobe

We hosted Anil Chakravarthy, President of Customer Experience Orchestration at Adobe, for a fireside chat focused on Adobe's AI strategy, monetization approach, and long-term differentiation as generative AI reshapes content creation and marketing workflows. Adobe is deliberately model-agnostic, now supporting 30+ third-party models, and does not view itself as competing at the model layer, instead positioning as an orchestration and workflow platform enabling customers to deploy the right model for the right task. Adobe outlined sharp differences in how customers are prioritizing AI: regulated industries focus on speed with compliance and commercial safety; consumer products emphasize cost reduction and efficiency; fashion and retail prioritize speed and content velocity around campaign iteration and time-to-market.

订阅研报添加微信 [illegible]

Management acknowledged that content creation itself is commoditizing as GenAI lowers barriers, but argued that Adobe's moat lies in making content commercially useful, with Brand Intelligence, first-party customer data, and complex end-to-end creative and marketing workflows as key areas of durable differentiation. Chakravarthy discussed Adobe's longer-term vision of consolidating a highly fragmented marketing stack into a super app spanning the content supply chain, and pushed back on the view that generative AI reduces the value of software, arguing instead that AI expands Adobe's TAM by accelerating workflows, increasing complexity, and enabling new use cases.

Architecting Data Panel

We hosted a panel on Architecting Data for AI moderated by Fatima Boolani, featuring Vasek Mlnejnsky CEO of E2B; Matt Fitzpatrick CEO of Invisible Technologies; and Venkat Rangapuram CEO of Centific. Rangapuram was direct: 60% of AI failures are driven by data issues, not model limitations. Fitzpatrick added that data harmonization is the single biggest initial gap, as the average large enterprise runs 16 legacy ERP, HR, and CRM systems that were never designed to talk to each other, let alone feed an AI agent, with governance and the shift from answering questions to taking actions as the second and third bottlenecks. E2B's numbers illustrate the pace of change: six months ago the platform was running 20 to 30M agents per month; today it is over 100M and growing.

订阅研报添加微信 [illegible]

On build vs. buy, Fitzpatrick cited data showing that the success rate on buying AI solutions is currently double that of building in-house, attributing the gap to internal builds becoming capability exercises rather than business initiatives without clear outcome measurement. Rangapuram offered the strategic frame:

build what encodes the competitive moat, buy everything underneath it, and recognize that inference will commoditize within three years, pushing value further up the stack. The panel's closing point was consistent with what was heard across the conference: companies that position themselves 6 to 12 months ahead of current model capabilities, remain model-agnostic, and focus on becoming more important as models improve will be the ones that compound rather than commoditize.

Fireside Chat with Applied Intuition

We hosted Peter Ludwig, Co-Founder and CTO of Applied Intuition, for a fireside chat covering the company's physical AI thesis, industry deployments, technology stack, and 2030–2035 milestones. Ludwig defined physical AI as the application of technology to things that move, and said Applied Intuition's thesis since 2017 has been that anything that moves will become autonomous. In trucking, the company is doing Level 4 autonomous trucking in Japan with Isuzu, noting that Japan's demographic problem makes technology the only solution and that Japan will be a lighthouse for the rest of the world. In defense, Ludwig said drones are the hot topic and that the company has been investing in low-cost collaborative drones for years, with tech and market timing now ripe for real growth.

Ludwig described Applied Intuition's work across three buckets: engineering and AI tools and infrastructure; operating systems technology for moving machines with safety-critical and low-latency requirements; and the actual AI systems and autonomy technology controlling vehicles across cars, trucks, agriculture, mining, construction, defense, and drones. By 2030, Ludwig said driverless applications of on-road technology will be widespread, and collaborative drone technology will see full operational deployment in defense. Home robotics will have a handful of interesting use cases by 2030 but will not be as mature as current marketing suggests, with 2035 being a more likely timeframe. On US government policy, Ludwig said the industry's biggest ask is a national framework, as the industry cares about rules being consistent more than about what the specific rules are.

Infrastructure Foundations of AI Panel

We hosted a panel on the Infrastructure Foundations of AI moderated by Mike Rollins at the 2026 Citi AI Summit, featuring J.J. Kardwell CEO of Vultr; Piotr Tomasik COO and Co-founder of Tensorwave; Nate Walkingshaw Founder and CEO of Torus; and Mohsen Moazami Vice-Chair of DDN. Moazami drew a direct line between Anthropic's $30B ARR trajectory and the amount of compute available to it, arguing that as long as the correlation between infrastructure capacity and revenue generation holds, every base case forecast will prove to be on the low side. Kardwell put a number on the supply-demand rebalancing timeline: 2029 at the earliest, even assuming no additional supply constraints emerge, noting that memory bottlenecks were not anticipated and power may be next.

Walkingshaw described a structural grid inefficiency not widely understood: once overhead capacity is purchased on the grid, it is owned permanently, like a water right, creating congestion that drives up costs and limits flexibility. Torus exists to identify and unlock that capacity through small modular power plants that can generate, store, and arbitrage power at the edge, skipping congested transmission infrastructure entirely. Moazami offered a striking data point: 37 companies on the planet are currently responsible for 98% of token production and consumption, all AI-native companies or coding environments, with traditional large enterprises held back not by lack of interest but by lack of trust in what happens to their proprietary data inside AI systems. The panel's closing caution was directed at the

legislative and regulatory environment: the industry needs to do a better job of bringing the public along, or political friction could become the constraint that supply chains and power grids are not.

Fireside Chat with Mercor

We hosted Adarsh Hiremath, Co-CEO and Co-Founder, Mercor for a fireside chat, moderated by Heath Terry, Global Head, Technology & Communications Investment Research, covering the company's business model, customer base, benchmarking strategy, and the role of human and synthetic data in AI model development. Mercor has grown from $1 million in revenue run rate to $1 billion recently, helping AI labs with training, deploying, and evaluating AI agents. Hiremath described Mercor's work as fundamentally high-skilled demand, requiring the best investment bankers, mathematicians, and software engineers to teach models how to perform domain-specific tasks, with lab budgets for data having gone from $100 million for specific datasets to orders of magnitude in the billions in a very short period, with no other area besides compute seeing budgets increase so rapidly.

Hiremath argued that in an enterprise context, human intelligence remains critically important, as a model may be able to create slides but cannot align them with an organization's best practices or a partner's understanding of quality. On synthetic data, he said it is complementary to human data rather than competing with it, and that for agentic use cases, a process called "world building" — giving files, tools, and context for the agent to succeed — may require synthetically engineered data. He said enterprises will ultimately be a larger business for Mercor than the labs, because as labor budgets shift to inference, agents will be applied throughout the entirety of the economy, and the number one barrier to enterprise adoption remains evaluations.

Revolutionizing Workflows

We hosted a panel on revolutionizing enterprise workflows, featuring Dheeraj Pandey CEO of DevRev; Wiley Jones CEO of Doss; and Keyvan Mohajer Co-Founder and CEO of SoundHound. Mohajer drew a parallel to the mobile transition, a battle of operating systems followed by a decade of value creation at the application layer, while Jones pushed back, arguing the industrial revolution is the more apt comparison: electricity and compute are the new constraints, and the scale of disruption to knowledge work is closer to what mechanization did to physical labor. Pandey framed the enterprise opportunity concretely, a $4T global knowledge work payroll market that agents are beginning to address, starting on the desktop and moving into the cloud.

Jones was direct: the rate limiter is already human comprehension, and the single most important factor in enterprise AI adoption is an authoritative, top-down mandate from inside the company. Pandey argued that personalization is the path to grassroots adoption, as the more an AI system knows about an individual employee, the more embedded it becomes. Mohajer closed with a point that cut against the prevailing conference narrative: small AI models are underhyped, and the next constraint will not be model capability but human error, the systematic failure to optimize decisions even when the tools to do so are available. The panel's shared conclusion was that the winners of this transition will be the organizations and individuals who learn to use AI effectively, not those who assume the technology will do the work for them.

Fireside Chat with Astera Labs

We hosted Jitendra Mohan, CEO of Astera Labs, for a fireside chat covering the company's founding thesis, product portfolio, competitive positioning, and outlook for AI connectivity infrastructure. Astera Labs was noted as the second-fastest semiconductor company to scale from inception to IPO, second only to NVIDIA, and the fastest company to scale to $1 billion in revenue. Mohan said the company was founded in late 2017 on the premise that AI connectivity would become a bottleneck, and its Scorpio PCIe switch product, introduced in October 2024 in partnership with NVIDIA and hyperscalers, crossed $100 million in revenue in three quarters, the fastest ramp the company has seen. On copper vs. optical, Mohan said copper has not run out of steam and will stay at the rack level providing the lowest cost and lowest power solution, while optical transceivers are used for longer distances at scale-out, with CPO solutions deployable for scale-up in 2028 at the earliest.

Mohan described a $20 billion TAM between memory-based protocols and ethernet-based protocols, and said the Cosmos software platform provides the customization, optimization, and diagnostic capabilities needed across different customer requirements. Talent was described as the single biggest challenge, as the company is limited by engineering capability rather than customer demand, having grown to approximately 1,000 employees from around 270 to 280 at IPO. Looking five years out, Mohan said customers will demand more tailored and bespoke solutions, physical AI will drive the need for more compute and connectivity, and optics presents a tremendous opportunity for Astera, pushing back on the "terminal value of copper is zero" narrative.

订阅研报添加微信 LCD123321Cad 04-17

AI in Healthcare

We hosted a panel on AI in Healthcare moderated by Heath Terry, featuring Ali Madani Co-Founder CEO of Profluent; Simon Burns CEO of Vial; and Ken Kaufman CEO of Sirona. Burns described Vial's mission as solving the bottlenecks in clinical trial execution at remarkably low cost per program, with the timeline from molecular design to entering the clinic shrinking from two to three years to closer to 12 to 18 months. Madani framed Profluent's approach as treating biology like a programming language, using the same transformer architecture that underpins LLMs but trained on billions of years of protein evolution, to design molecules from scratch. The most energizing proof point cited was Baby KJ (May 2025), a rare pediatric disease case in which a novel treatment was identified, designed, and administered within six months, moving AI-designed medicine from theoretical to tangible.

Kaufman noted that $6B has been spent on radiology AI, yet financial returns have been negligible, not because the algorithms don't work, but because the delivery infrastructure is broken, with 7,000 on-premise radiology systems sitting in hospital basements running on technology not meaningfully updated in decades. Sirona's response is to be the first cloud-based radiology platform, taking the 70% of healthcare data that is radiology data and moving it to an environment where new software can be delivered every two weeks. The panel was candid that physician adoption of AI is not happening as enthusiastically as the technology would warrant, with the wins gaining traction being passive, such as ambient listening tools that populate medical records without adding to physician burden, and that a new generation of physicians, more native to AI tools, will be required before adoption reaches its potential.

Bringing AI Apps into Production

We hosted Navrina Singh, Founder and CEO, Credo AI; Jeffrey Akiki Co-Founder & COO, Not Diamond; and Jason Warner, Co-CEO and Co-Founder, Poolside for a panel, moderated by Mark Zhang, Senior Research Associate, US Software Equity Research, on bringing AI applications into production. Singh said what really matters is getting a clear view and confidence into the systems being deployed, which requires understanding holistic risks — reputational, regulatory, and operational — as well as continuous runtime monitoring. Akiki said governance does not get enough attention during the pilot phase, and that cost is a major issue, as people leverage state-of-the-art models during pilots but at scale this becomes extremely costly. Not Diamond's dynamic routing logic allows customers to define their priorities between cost, latency, and accuracy, with potential savings of more than 90% on inference costs in some cases.

订阅研报添加微信 LGD13321Gal 04-17

Warner said the things that matter over the next 10 years are energy, compute, and power infrastructure, and that enterprises that last will take advantage of proprietary data and build custom domain models. Singh noted that shadow AI reveals which applications and models make sense to bring into the enterprise, and that she has seen cases where customer retention has gone up with governance in place. The panel's shared conclusion was that the only durable companies will be those that can make apps that work for everyone, built on trust, which requires continuously measuring whether the objectives set for a system are being met within the specific use case, context, and industry.

Glean Fireside: The Connective Tissue of AI

订阅研报添加微信 LGD13321Gal 04-17

We hosted Tony Gentilcore, Co-Founder of Glean for a fireside chat moderated by Heath Terry. The session covered Glean's evolution from enterprise search to a horizontal AI platform, the foundational role of data governance and the enterprise graph in enabling effective AI deployment, the economics of token consumption at scale, and the company's growth trajectory and model-neutral strategy. What began seven years ago as high-quality enterprise search has evolved into what Gentilcore described as the connective tissue for enterprise AI operations, a platform that understands not just what data exists, but what is canonical, what is deprecated, and what matters to which employees across a global organization. Glean's model-neutral architecture, spanning 15+ LLMs, allows enterprises to optimize across cost, performance, and capability on any given dimension and to benefit from model competition rather than be exposed to it, with Gentilcore's view that model specialization will deepen over time, making the ability to route intelligently across providers increasingly valuable.

Gentilcore highlighted one engineer at Glean running $12,000 per month in personal token usage, a signal that AI spend is no longer an IT line item but closer in nature to headcount, requiring the same kind of allocation discipline and ROI accountability that hiring decisions demand. He illustrated the ROI framing with Dell's 26,000-person sales team (May 2025), which saw a 50% reduction in support requests after deploying Glean. Beyond search and retrieval, Glean is building an activity graph, a dynamic map of how individuals and teams work, what they prioritize, and what actions they take, that gives AI agents the contextual grounding to operate effectively rather than just retrieve information. Glean doubled its ARR in nine months and is targeting another 2x, and on security, Gentilcore was direct: with prompt injection, PII exposure, and data exfiltration risks accelerating, he expects major incidents in the next three to six months.

The VC Perspective on AI: What's Next?

We hosted a venture capital panel moderated by Todd Speece, Head of VC Coverage at Citi, and Heath Terry, featuring Ganesh Bell, Managing Director, Insight Partners; Lucas Swisher, General Partner, Coatue Management; Chetan Puttagunta, General Partner, Benchmark; Chris Emanuel, Head of Technology Investment Group, GIC. The session brought together investors with front-row seats to the most consequential platform shift in a generation — collectively invested across frontier model labs, application companies, and infrastructure — to discuss what is driving the current acceleration, where durable value is being created, and what comes next.

The defining shift is from AI that answers questions to AI that does work. Swisher framed it plainly: the transition from copilot to agent is the biggest shift underway, and the usage data is confirming it in real time. Harvey, a legal AI portfolio company, shared charts at a recent board meeting showing minutes-per-day usage going parabolic — a linear curve that has turned vertical as AI moves from research assistance to handling real workflows end-to-end. Emanuel added the labor economics dimension: AI makes labor variable rather than fixed, and that shift is showing up in adoption curves that look unlike anything prior platform transitions produced. Anthropic's revenue trajectory was cited as the starkest illustration — growing from $9.4B at year-end to $19B in February to over $30B in March, the equivalent of adding four Databricks-scale customers in three to four months.

Manus was the clearest early signal that consumer demand for agents is insatiable. Puttagunta described investing in Manus in March 2025, before its public launch, on the basis that the product was materially ahead of the market in agent-based thinking — pushing further into real task completion than anything else available at the time. The launch uptake was extraordinary: zero to approximately $100M run-rate in roughly eight months, driven almost entirely by consumer demand. When enterprises asked to purchase licenses, Manus refused to invoice them — forcing credit card signups — a deliberate choice that proved the point about consumer willingness to pay meaningfully for software that does real work. Post-acquisition inside Meta, Manus has become significantly more sophisticated. Puttagunta's conclusion was direct: the $1T in global software spend is genuinely up for grabs, and SoR are proving to be far less durable as lock-in mechanisms than previously assumed.

Durability in the AI stack comes from a short list of structural advantages. Swisher laid out the framework: durable defensibility at the application layer requires control over at least one of four things but ideally all — the chip layer, the language model layer, a large proprietary user base, or proprietary data. Possessing all four is rare and highly attractive, which is why frontier model companies are vertically integrating into chips. Traditional SaaS metrics still matter — are customers consuming more tokens year over year, does usage compound — but they are necessary rather than sufficient conditions. Bell added a structural observation: the GPU represents an entirely new architecture to build against, and every frontier lab is building custom tooling and middleware that creates new surface area for startups. The tension between a new stack emerging around agents and workflows and the compounding improvements in frontier foundation models is real, and every actor in the ecosystem is optimizing for a different constraint — cost, latency, capability, or control.

Security requires a complete rewrite, and the window to act is narrowing. Puttagunta was the most direct voice on this: **the value of code itself is trending**

1,500 employees globally. Steele described two product lines: a software product line delivering mission autonomy, enabling systems to coordinate, swarm, and team, which has been flown in over 26 aircraft; and a Group 3 drone deployed with the Coast Guard to help identify drug cartels at sea. Shield AI has been deployed in Ukraine for approximately 2.5 years with roughly 30 employees on the ground, operating in GPS-denied and communications-denied environments, providing target coordinates in the battlefield and coordinating Group 3 drones and one-way attack drones with the Ukrainian military.

订阅研报添加微信 LCD123321Cad 04-17

Steele described the government's decision to separate the software decision from the aircraft decision in the Collaborative Combat Aircraft (CCA) program as a progressive move reflecting the government's desire for an open architecture and no single-vendor dependency for software. Shield AI raised $2 billion several weeks ago, doubling its valuation, and used proceeds to acquire Echelon, a leader in manned simulation and training, with the goal of bringing unmanned and manned simulation together under the Hivemind platform. Steele announced a next-generation autonomous airplane with the same munitions as seen on the F-35, runway-independent, with a first flight planned for fall 2026 and production expected in the 2029 to 2030 timeframe, describing the product as a potential game-changer that could generate billions in revenue over time.

Lightmatter Fireside: The Next 1,000X

We hosted Nick Harris (CEO) and Simona Jankowski (CFO) of Lightmatter for a fireside chat moderated by Atif Malik. Lightmatter's foundational thesis is that linking chips together efficiently, not making individual chips faster, is the defining infrastructure challenge of the AI era. The core problem is geometric: compute scales with chip area, but I/O is constrained to the chip perimeter, creating a bandwidth bottleneck that worsens as chips get larger. Lightmatter's photonic interposer technology resolves this asymmetry by enabling I/O to scale with area rather than perimeter, with the M1000 platform demonstrating 114 terabits per second at two picojoules per bit, 10x more energy efficient than pluggables and 70x faster than current NVIDIA CPO.

订阅研报添加微信 LCD123321Cad 04-17

Lightmatter's roadmap runs from NPO today to CPO and photonic interposers through 2028, with a Qualcomm partnership (March 2026) demonstrating fourth-generation OCI technology at 1.6 terabits per second per fiber already in silicon. Critically, Lightmatter has also developed a guided laser product line called VLSP using very large scale photonics on 300mm CMOS fabs, a silicon photonics approach that sidesteps the indium phosphide supply bottleneck entirely. Harris was direct about the competitive landscape: TSMC is highly selective about the partners it works with, and Lightmatter's partnerships with TSMC, Global Foundries, Tower, ASE, and Amkor, combined with 18-to-24-month chip design cycles that lock in customer relationships at the architecture phase, create a structural lead that is difficult to close. With 90% or more of the GPU market flowing through TSMC, Lightmatter's position in that ecosystem is a meaningful competitive moat.

订阅研报添加微信 LCD123321Cad 04-17

Physical AI

We hosted Jon Runyan, COO and Co-Founder, Armada; Ali Agha, CEO, FieldAI; and Steph Tryphonas, Co-Founder, VP of Business Operations & Customer Success, Collaborative Robotics, for a panel on Physical AI. Agha said physical AI means more general-purpose solutions, as robotics used to be single-purpose for decades but now robots are doing more complex tasks, benefiting from digital AI advances, though hallucinations carry greater consequences in the physical world. Tryphonas described Collaborative Robotics' robot-as-a-service model for material

On product evolution, Ghelberg said Harvey's first product was an assistant and copilot that is now becoming increasingly agentic, with new models in the last six months enabling long-horizon agents and pulling forward parts of the product roadmap. The company launched Shared Spaces to enable collaboration across counterparties, incorporating agents into the process instead of email attachments, and is investing in post-sales legal engineers, formerly lawyers, to train, onboard, and do change management. On the future of the profession, Ghelberg said the human elements — trust, strategic advice and counsel, and understanding certain situations in ways technology cannot — will remain and become more important earlier in careers, and that he believes there will ultimately be more lawyers and more legal output.

订阅研报添加微信 LCD123321Cad 04-17

Embedding AI Into Workflows: Customer Perspectives

We hosted a panel on Embedding AI into Enterprise Workflows moderated by Vibhor Rastogi, Managing Director, Venture Investing, Citi Ventures, featuring Gary Kotovets, Chief Data, Analytics and AI Officer, Dun & Bradstreet; Ishan Mukherjee, CEO, Rox. The session centered on a theme that cut across every layer of the AI stack discussed at the conference: before agentic workflows can function reliably at enterprise scale, the foundational questions of identity, data integrity, and trust must be resolved, and most organizations are not yet there. Kotovets described a survey of CDOs in which the entire conversation, across 20 CDOs and 40 to 50 CIOs, was consumed by a single topic: data governance, with the number one question being how to get the data straight before standing up agentic workflows. Dun & Bradstreet's response is to serve as that reference layer, with 6T data points organized around a common entity resolution framework and pre-calculated scores and analytics that allow agents to make deterministic decisions without burning compute on probabilistic resolution at runtime.

‹Φ—xb▪mûR_®Oa LCD123321Cad 04-17

The panel converged on a sequencing point with significant implications for enterprise AI deployment timelines: business identity and human identity must be resolved first, before agents can be trusted to act autonomously across workflows. The call center was cited as a microcosm of the broader challenge, as significant effort is going into letting agents resolve issues without human intervention, but fraud risk and identity uncertainty mean that automation is still very much a work in progress in high-stakes customer interaction environments. Kotovets described a deliberate approach of validating models across half a dozen pipelines and use cases, a mix of open and closed source, before committing to any single provider, with 40% of engineering code now written using AI coding tools. The panel's broader point echoed across the conference: the companies that will capture durable value in the agentic era are not those that move fastest to deploy agents, but those that invest first in the trust, identity, and data infrastructure that makes agent outputs reliable enough to act on.

From Model to Motion: The Skild AI Approach

‹Φ—xb▪mûR_®Oa LCD123321Cad 04-17

We hosted Deepak Pathak, Co-Founder and CEO of Skild AI, for a fireside chat, moderated by Heath Terry, Global Head, Technology & Communications Investment Research, covering the company's approach to building a general-purpose brain for robotics, its technology stack, data strategy, and partnerships. Skild AI is building what Pathak described as the equivalent of ChatGPT or Claude for robotics, a general-purpose, omni-bodied brain for any robot form factor and any task. Pathak said the biggest missing piece in robotics has been data, and that the industry has historically created ultra-specialized systems, robots that are more like microwaves, designed to perform the exact same tasks every time. He argued that a general model serves as the foundation for specialized tasks, providing the common sense for robotics that has been missing,

and pushed back on the perception that China is ahead in robotics, saying it is primarily a brain problem, and that the ultimate metric of success is tracking live deployments, by which measure the number is essentially zero for everybody.

Pathak described Skild AI's pre-training approach as using two scalable data sources, human videos and simulation, before deploying and fine-tuning, similar to LLM-style pre-training and post-training. The biggest dataset for robotics is approximately one million data points, compared to the hundreds of trillions at

订阅研报请加微信 LCD123321Cai 备注

We hosted Alan Trefler, Founder and CEO of Pegasystems ies, making synthetic **moderated by Steve Enders, Vice President of US Softwa** simulation data essential. Skild AI has reduced the data required for a task from 17 hours to four hours, and robots no longer need to be pre-programmed and can handle disturbances. On partnerships, Pathak noted that ABB has half a million pre-programmed robots in the field, and that by providing the model as the brain while ABB provides hardware and tools, Skild AI can access that installed base, with the goal of entering the warehouse market and creating a fully autonomous warehouse of the future.

Broadcom Keynote: Custom at Scale

We hosted Charlie Kawwas, President of Semiconductor Solutions at Broadcom, for a keynote conversation moderated by Mark Keene. The session covered Broadcom's XPU platform strategy, the company's line of sight to $100B in AI revenue, the competitive dynamics between custom silicon and GPUs, supply chain differentiation, and the role of Ethernet as the networking backbone of AI infrastructure at scale.

订阅研报请加微信 LCD123321Cai 备注

The AI infrastructure wave is larger than anticipated. Kawwas opened with a striking data point: hyperscalers will spend $665B in IT infrastructure this year, up from less than $250B last year, with at least half directed at AI according to public market data. Adding tier-two and neo-cloud spend pushes the total toward $800B. Broadcom has scaled its semiconductor revenue from $18B to $50B on the back of this wave, with a stated path to $100B next year — a trajectory Kawwas grounded not in projection but in the visibility of existing customer commitments and platform roadmaps.

XPUs are not ASICs — they are a platform, and they are on a path to surpass GPUs. Kawwas was precise about the distinction: GPUs are general compute platforms that made sense for the initial training phase of AI, but they carry a general compute tax in power and total cost of ownership that becomes increasingly difficult to justify at scale. Broadcom's XPU platform, built over the past two years, delivers the same workload at materially lower power and cost. Meta will produce four XPUs in two years according to Meta's Blog (March 2026); Broadcom currently has six XPU customers and is selective about adding more, requiring at least two to three generations of commitment given the several-hundred-million-dollar development cost per platform. Kawwas's forecasts that the XPU unit market will be very close to the GPU market next year and will surpass it by 2028.

订阅研报请加微信 LCD123321Cai 备注

The economics of the XPU model are structurally advantaged over building alone. Kawwas explained why Broadcom can build custom silicon at lower R&D spend than GPU players: NRE payments are collected across each technology feature from multiple customers, and a central engineering organization sits across all divisions, creating reusability that individual hyperscalers building in-house cannot replicate. The NRE alone on a 3-nanometer platform exceeds hundreds of millions, which is why only half a dozen to ten companies in the world have the scale to pursue custom compute at this level. Kawwas noted the dynamic: the

faster NVIDIA runs, the more its customers need Broadcom — NVIDIA itself is a Broadcom customer, as is AMD — making the XPU platform a beneficiary of GPU adoption rather than purely a competitor to it.

Supply chain self-sufficiency and Ethernet networking are Broadcom's two most underappreciated advantages. On supply, Kawwas described a substrate factory going into production in September — substrates are in massive shortage, and Broadcom is the only company with its own substrate manufacturing capability. The philosophy is simple: if it is strategic, build it. On networking, Kawwas was emphatic that an open networking ecosystem is not optional at AI scale — without it, customers hit a data wall that prevents scaling regardless of compute capacity. Ethernet, he argued, has decisively won over InfiniBand for scale-out and scale-across architectures, and will extend into scale-up as well. Broadcom is also the only company currently shipping CPOs in production, and can support copper connectivity up to 400 gigabits — positioning it across the full spectrum of interconnect solutions as the industry transitions from pluggables to CPO.

订阅研报添加微信 LCD123321Cad 04-17

订阅研报添加微信 LCD123321Cad 04-17

订阅研报添加微信 LCD123321Cad 04-17

Key Stocks Mentioned:

Adobe Inc. (ADBE.O; US$248.16; 2; 16 Apr 26; 16:00) | Akamai Technologies, Inc. (AKAM.O; US$96.81; 2; 16 Apr 26; 16:00) | Alphabet Inc (GOOGL.O; US$336.02; 1; 16 Apr 26; 16:00) | Amazon.com, Inc. (AMZN.O; US$249.7; 1; 16 Apr 26; 16:00) | AT&T Inc (T.N; US$26.4; 1; 16 Apr 26; 16:00) | Broadcom Inc (AVGO.O; US$398.47; 1; 16 Apr 26; 16:00) | Cloudflare Inc (NET.N; US$197.38; 1H; 16 Apr 26; 16:00) | Coherent Corp (COHR.K; US$328.0; 1; 16 Apr 26; 16:00) | Crowdstrike Holdings Inc (CRWD.O; US$418.2; 1; 16 Apr 26; 16:00) | Datadog (DDOG.O; US$123.47; 1; 16 Apr 26; 16:00) | Digital Realty Trust Inc (DLR.N; US$199.05; 1; 16 Apr 26; 16:00) | DigitalOcean Holdings, Inc. (DOCN.N; US$88.43; 1; 16 Apr 26; 16:00) | Dynatrace (DT.N; US$35.62; 1; 16 Apr 26; 16:00) | Equinix Inc. (EQIX.O; US$1070.9; 1; 16 Apr 26; 16:00) | Lumen (LUMN.N; US$8.51; 2H; 16 Apr 26; 16:00) | Lumentum Holdings Inc (LITE.O; US$891.22; 1; 16 Apr 26; 16:00) | Palo Alto Networks, Inc. (PANW.O; US$166.97; 1; 16 Apr 26; 16:00) | Rubrik, Inc. (RBRK.N; US$51.94; 1H; 16 Apr 26; 16:00) | Salesforce Inc (CRM.N; US$181.22; 2; 16 Apr 26; 16:00) | Uniti Group Inc (UNIT.O; US$11.06; 2H; 16 Apr 26; 16:00) | Verizon Communications Inc (VZ.N; US$46.78; 1; 16 Apr 26; 16:00)

订阅研报添加微信 LCD123321Cad 04-17

订阅研报添加微信 LCD123321Cad 04-17