

Revisiting Moats IV: Shifting Center of Gravity

15 April 2026 | 12:47AM EDT

We met with 30 private companies and 10 public companies as part of our Private Company Software & Internet Conference (4/8-4/9) and our Annual Silicon Valley Bus Trip (4/13-4/14), alongside conversations with several VC thought leaders. Public companies we met with: ADSK, CRM, CRWD, GWRE, IOT, OKTA, RBRK, WDAY, and ZS; as well as pre-quiet check ins with PANW & SNOW.

We are more convicted in the Software TAM: companies are moving to sell units of labor or productivity, which allow them to capture more budget; and AI Natives are selling into the white space between traditional SaaS silos. As the frontier labs set new highs on benchmarking and associated inference costs, the industry is countering by redistributing value throughout the stack. This was most evident to us in the discussions around using open source distilled models and the focus on roadmaps that are either independent from the frontier labs because they are specialized, or weighted to proprietary SLMs. This dynamic is one of the more compelling reasons why enterprises will likely engage with an ecosystem of software providers that are independent of any one model provider: the enterprise benefits from much more efficient routing and lower inference costs, and is not exposed to step function changes in price if any one model provider stops subsidizing tokens.

Of the public companies we met with, we are more positive on CRM (more thoughtful agentic strategy and upcoming launch of dev tools); CRWD (several modules seeing better demand); GWRE (AI accelerating cloud conversions and purely TAM expansive); IOT (ROI-driven demand, data / AI differentiation); RBRK (Agent Cloud / Identity broadening the story; expanding resilience needs)

Key Takeaways

- 1. We are more convicted on the Software TAM getting bigger.** Companies are increasingly positioning their AI workflows as selling a unit of labor or a unit of productivity, which allows them to tap into larger deal sizes and new budget allowances. (**Sierra** notes the atomic unit of productivity in AI is a process, not a person; **Salesforce** similarly introduced Agentic Work Units; **Workday** maps flex credits to “units of work”). This also allows platforms to anchor to value-based selling and decouples gross margin from underlying LLM costs, such that the ecosystem could still solve for 80%+ gross margins. In the meanwhile, **Kleiner**

Table of Contents

Archera	9
Automation Anywhere	9
Autodesk	10
Axonius	11
Baseten	11
Bright Machines	13
CrowdStrike	13
Decagon AI	14
Fivetran	15
Glean	16
Grafana	16
Guidewire	17
Hightouch	18
Huntress	19
Intercom	20
Island Technology	21
KnowBe4	21
Kong	22
Okta	22
Outreach	23
Physical Intelligence	24
Salesforce	25
Samsara	27
Serval	28
Shield AI	29
Sierra Technologies	30
Superhuman	31
Tekion	31

Traversal	32
Turing	33
Rubrik	34
Writer	35
Workday	36
Zscaler	37
Disclosure Appendix	38

订阅研报添加微信 LCD123321Cad 04-16

订阅研报添加微信 LCD123321Cad 04-16

订阅研报添加微信 LCD123321Cad 04-16

Perkins noted that even when solving for 60% gross margins, gross profit goes up by a meaningful factor. As **a16z** framed it: AI is not a zero sum game; it creates a meaningfully larger TAM. The best defense is a good offense; incumbents should go out and build with their advantages including access to capital and therefore compute.

2. **As the frontier labs continue to set new highs on benchmarking and their associated inference costs, the industry is responding by finding more ways to restore the balance of power and redistribute value throughout the stack.** This was most evident to us in the discussions around using open source distilled models (we heard that distilling a model costs just 2% of the cost of training the parent model), as well as the focus on roadmaps that are either entirely independent from the frontier labs because they are more specialized (CrowdStrike), or increasingly weighted to proprietary SLMs. **Superhuman** now routes 97% of its API calls through its proprietary LLMs, **Intercom** leverages a mix of 12 models with the majority of routing through proprietary models, and realized an 80% reduction in cost when it rebuilt its reranker on its proprietary model; **Writer** just released a self-evolving LLM – these smaller models run in the background and train themselves on personalized usage, isolated to the enterprise environment with full model sovereignty. After 5 training runs, Writer believes this model will perform on par with any frontier LLM, because it more than compensates for not being on the leading edge with domain specific knowledge. This balance of power argument is one of the more compelling reasons why enterprises will engage with an ecosystem of software providers that are independent of any particular model provider: the enterprise benefits from much more efficient routing and lower inference costs, and is not exposed to step function changes in price if any one model provider stops subsidizing tokens.
3. **Compute has emerged as the binding constraint in scaling AI.** As inference demand grows faster than available capacity, the key differentiator is no longer just model quality, but the ability to reliably access and finance compute in the most performant way. As a result, value is shifting towards how models are run in production (e.g. the agentic harness and model routing capabilities) and what vendors can commit to customers around cost, reliability, and performance.
 - **BaseTen commented that the supply/demand imbalance is materially worse than most industry participants perceive.** Companies are overrunning their initial budgets for inference by orders of magnitude (we heard one industry datapoint on inference costs in engineering now approaching about 10% headcount cost, but could be on track to be on par with headcount costs in the next several quarters based on current trajectories).
 - **Multi-cloud production is becoming a necessity, not a preference.** Our conversations reinforced that supply constraints, pricing volatility, and uneven access to leading-edge accelerators are forcing customers to run inference across hyperscalers and neoclouds simultaneously. These workloads are stickier than training workloads, therefore providing more durable revenue for neoclouds that have traditionally been more exposed to training.

4. **In Security, there is a compelling path to a more secure enterprise stack that pushes the industry toward higher ROI solutions.** Huntress acknowledged that the odds could shift away from the defenders over the next 12-18 months as the industry executes on a step function change toward a more comprehensive patching and hygiene upgrade cycle, given the urgency of hardening existing infrastructure ahead of new model releases (new model releases will in turn likely build in more thoughtful safety alignments before broad distribution – see ChatGPT 5.4 cyber release [here](#)). This cycle will deploy AI to fix most of the low hanging fruit in the stack, similar to the upgrade cycle in 2000 ahead of Y2K or the firewall capex cycle that established perimeter-based security in the early 2000s. While the next big patching cycle may not be discretely one time in nature, this first lift will likely be more substantial than subsequent lifts. The vendors that are part of leading edge LLM programs that grant preferential access to LLM innovation are likely to be viewed by customers as leading edge, and this will likely create further separation in momentum between these companies and the rest of the field. We then see two subsequent product cycles: in the Security Operations Center to account for higher volumes of alerts that will need to be more automated (**CrowdStrike, Palo Alto, Huntress**); and for identity for agents (**Okta, Microsoft**). In network security, **Zscaler** is already seeing an uptick in agentic traffic on their networks, the next question will be to what extent network security providers will be able to capture this upside via hybrid user x bandwidth monetization. Please see our 4/8 Security note post our meeting with CrowdStrike's SVP of Counter Adversary [here](#).
5. **In Observability, Grafana and Traversal illustrated two competing dynamics: a) observability budgets are accelerating; b) customers are increasingly focused on total cost of ownership and more open architectures.** Grafana noted an increase in demand (their cloud business is accelerating, as is Datadog's) as a function of more microservices and greater data labeling; AI can help SREs (site reliability engineers) extract more context out of longer duration datasets. However, we expect to see a push towards more open standards that in turn reduce the cost of processing and storing telemetry (Grafana noted their adaptive trace technology allows customers to drop 90% of the traces sent to infrastructure monitoring), and we continue to assess whether newer generation companies like Traversal are well positioned to build AI SREs on top of existing incumbents. This may drive a shift away from dashboards toward search-driven interfaces, where the system surfaces the relevant signal rather than requiring a human to review dashboards.
6. **AI Natives are not competing head-on with SaaS incumbents, they are winning by occupying the white space between them.** As software value shifts from systems of record to systems of action, the center of gravity moves away from databases and dashboards toward agents that resolve problems end-to-end. Incumbent SaaS products were built around discrete functional silos while AI Natives are built around outcomes. AI Natives are winning with packaged solutions that go live in weeks, convert proofs of concept at high rates, and monetize on business impact rather than seat count or feature breadth.

- **This advantage compounds because AI Natives naturally span multiple tools.** Real enterprise work does not live inside a single application, it cuts

across CRM, OMS, ticketing, logistics, payments, and bespoke infrastructure. Agents are not confined into a single system. Orchestration across systems, models, and workflows becomes the new control layer, replacing suite depth with cross-stack coherence. Where incumbents are constrained by legacy architectures and module-centric roadmaps, AI Natives are defined by their ability to integrate, adapt, and execute across the full enterprise stack.

- **This dynamic may weaken two of SaaS' most-cited moats: domain experience & data.** Our conversation with Sierra and Serval highlighted the domain-experience moat is less durable than advertised, as incumbents rarely have the right to cross-functional domain knowledge. Organizational silos, permissions, and legacy design limit their view of the end-to-end process. At the same time, the data moat in enterprise software is structurally overvalued as the data belongs to the customer, not the vendor, with SaaS companies acting as stewards rather than owners. As a result, defensibility migrates away from the vendor that houses the data, and toward the vendor that can safely, reliably, and repeatedly turn that data into action across systems.

7. Applied AI innovation may be at an inflection point that parallels the GenAI inflection point in 2022. While meaningful commercialization may still be 5-10 years out, there are meaning advances happening in Applied AI as progress in foundation models, simulation, and edge autonomy collapse the gap between digital intelligence and real world action, creating durable software-defined platforms across logistics, labor automation, and defense. The critical shift is that intelligence is no longer confined to perception or decision support, but embedded directly into machines that act, learn, and improve in continuously changing environments. Across verticals, winners are those with sustained real world deployments that generate proprietary data flywheels, enabling compounding performance improvements that are structurally difficult to replicate with simulation or lab only approaches.

- **Robotics:** Locus Robotics and Physical Intelligence discussed how model breakthroughs are expanding automation. Locus said its RaaS model lowers adoption friction and maximizes real world data capture across ~17k deployed robots, while the Array platform extends automation into more autonomous, higher value workflows. In parallel, Physical Intelligence is pushing beyond task specific robots toward generalizable foundation models trained on proprietary real world interaction data, unlocking broader labor substitution over time. Physical Intelligence highlighted the technical challenges of collecting and simulating real world data such as the texture of a glass and the properties of the water within it; their focus is driving the robots through as many real world use cases as possible to collect this compounding pool of physical data.
- **Defense: Shield AI's** Hivemind exemplifies the shift toward software centric defense platforms that can reason, coordinate, and execute missions autonomously at the edge, supported by simulation-led validation and embedded assurance. This model accelerates time to deployment, supports swarm level behaviors, and allows venture backed innovators to out iterate

traditional primes despite slower procurement cycles.

8. Company building in 2026 looks very different to company building in 2022. The importance of moving quickly was obvious to us in all of our conversations – and there is separation between public and private vendors that understand the need to innovate at a much faster pace today vs. those that have yet to fully appreciate or set up their org structures for this pace of innovation. **General Catalyst** noted that across their B2B SaaS portfolio, companies are now allocating savings from engineering headcount (e.g. hiring fewer junior hires) to S&M (budgets 30-40% higher, accruing to performance marketing channels and AI SDRs), and that in a market opportunity that is potentially so much larger than before (given labor budgets > software budgets), companies should not be solving for FCF or high LTV:CAC today. Instead, as long as the ROI on each customer cohort is strong, companies should consider using S&M as a moat: once you land a customer, you can then solve for cross sell and higher lifetime value than in prior software evolutions. Companies are also rethinking traditional org structures: **Serval** noted that the traditional notion of customer success is ineffective, and instead aligns incentives such that the deployment strategist/FDE is also responsible for customer success; and that in many cases it is not hiring for specific roles (as the roles are being redefined real time), but rather, specific skillsets. Serval also looks for creative ways to solve the process management problem, such as hackathons that allow the AI native company to paint a picture of what ideal workflows should look like, rather than working within existing workflow constraints.

Please also see our other “Revisiting Moats” reports:


- [Revisiting Moats Part I: Exploring AI Steel Men Arguments](#)
- [Revisiting Moats II: Datapoints on General Intelligence Tools](#)
- [Revisiting Moats III: Cybersecurity](#)

We summarize company specific highlights below.

订阅研报添加微信 LCD123321Cad 04-16

Exhibit 1: With competition elevated from two new cohorts in the Software app layer (AI Natives and Frontier LLMs), the question remains which cohort can best address or compensate for their disadvantages

GS illustration of strengths and weaknesses across the 3 cohorts of competition



Source: Goldman Sachs Global Investment Research

订阅研报添加微信 LCD123321Cad 04-16

订阅研报添加微信 LCD123321Cad 04-16

We hosted Archera's CEO and Co-Founder, Aran Khanna. Archera helps businesses manage and optimize cloud spend through AI-driven forecasting and automation, FinOps insights, and insured commitments. Through offerings such as Insured Commitments, Guaranteed Savings Plans (GSPs), and Guaranteed Reserved Instances (GRIs), Archera enables businesses to secure cloud discounts with flexible terms, reducing costs risks associated with traditional long-term cloud commitments.

- **AI has caused a combination of supply constraints and increased uncertainty that has increased the complexity of purchasing cloud compute.** This has made customers hesitant to commit to longer term contracts but also increased the leverage CSPs have. This has resulted in hyperscalers raising prices on certain high-end GPUs (H200 & GB200) and limiting commitment availability, forcing customers into bidding on spot-markets depending on supply-demand conditions. Overall the massive increase in cloud computing spend has increased the market opportunity for Archera.
- **Multi-cloud is no longer optional experimentation but a permanent production reality.** Archera highlighted a clear shift from customers running on one or two clouds to operating production workloads across three or four providers; driven by GPU supply constraints, volatile pricing, and hyperscalers' reluctance to lock in pricing on leading-edge accelerators (e.g., H200, Blackwell). Archera is seeing inference/production workloads being run on neoclouds/Oracle due to supply constraints.
- **Archera's differentiation increasingly resides in its underwriting data, not just its tooling.** The company's ability to guarantee commitments is informed by granular consumption data across almost 2k customers with ~\$5bn in cloud commitment spend, enabling high-confidence claims management. As agentic workflows drive step-function increases in compute intensity per user, Archera sees underwriting-backed commitment guarantees becoming an important factor for ISVs and digital natives that want to scale aggressively without locking in provider- or SKU-specific risk.

Automation Anywhere

We hosted Automation Anywhere's Co-Founder & CEO, Mihir Shukla. Automation Anywhere is a provider of AI-powered automation, providing an enterprise-grade platform that combines traditional Robotic Process Automation (RPA) with advanced generative AI and agentic capabilities. Automation Anywhere enables organizations to deploy intelligent digital workers and autonomous AI agents that can reason, make decisions, and execute complex business processes across fragmented legacy systems and modern cloud applications.

- **SaaS disruption (the agent layer eats the UI layer).** Mr. Shukla makes the claim that AI will effectively re-bundle SaaS. If agents can navigate, reconcile, and execute

across systems, the value migrates away from human-facing workflows toward systems-of-record + APIs - while an orchestration layer coordinates everything above it. In that world, enterprises don't want hundreds of AI applications and instead will standardize on 1-2 platforms to run work across the sprawl, and that standardization shows up bluntly as seat-count compression (not just efficiency) - especially in categories like ITSM where a large share of tickets can be resolved without a person ever opening the tool.

- **AI as labor arbitrage + throughput unlock.** Mr. Shukla noted the practical unit of value should not be feature automation, but instead cross-application time reclamation. Today workers spend a minority of time in their primary system, so optimizing that core application still only impacts a slice of the day. The bigger unlock is stitching the 8-18 applications that actually constitute the job into one continuous, partially autonomous flow - turning coordination labor into machine labor. That reframes ROI as a mix of cost take-out (fewer manual touches, fewer exception handlers) and capacity creation (same team, more throughput), with risk reduction where errors are expensive or regulated.
- **Long approval cycles are a main bottleneck to AI adoption.** Companies can stand up meaningful automation fast but approvals remain a lengthy process. Mr. Shukla noted Automation Anywhere could automate the bulk of ITSM tickets in eight weeks, though receiving the internal approvals can take nine months or longer. The gating factor isn't model capability, it's governance cadence. Additionally, because multi-step automations mix probabilistic and deterministic steps, enterprises are conservative, and teams demand guardrails, auditability, and deterministic outcomes even when some steps are probabilistic. Ultimately, Mr. Shukla believes the winners won't just have better agents, they'll have faster approval machinery.

Autodesk

We hosted Simon Mays-Smith (Head of IR) of Autodesk.

- **Management discussed a series of commercial changes aimed at simplifying the go to market model and improving customer visibility.** These included recent adjustments to sales coverage, continued evolution of channel relationships, and an increasing emphasis on direct billing and e-store motions. Management noted that moving closer to the end-customer improves visibility into usage, purchasing patterns, and renewal behavior, while also enabling greater automation across renewals, co-terming, and lower-touch transactions. They acknowledged that these transitions can introduce short-term noise as relationships and buying motions reset but framed the changes as important steps toward a more scalable and transparent commercial model.
- **AI value is emerging along a continuum, with early traction in task level automation and productivity.** Management described AI adoption as progressing from narrow task automation toward broader workflows over time. They pointed to concrete examples such as constraint automation, where customer acceptance rates

have improved meaningfully, delivering tangible throughput benefits without requiring customers to re-architect end-to-end processes. As AI usage expands and compute intensity increases, Management emphasized the importance of aligning monetization to value delivered, noting that consumption-based models already represent a portion of the business and are a natural fit for certain AI-driven use cases.

- **Autodesk emphasized proprietary models and infrastructure as central to delivering AI in real world, deterministic workflows.** Management noted that AECO and Manufacturing customers generally lack the data access, foundation models, and inference infrastructure required to build advanced AI capabilities independently. As a result, Autodesk has invested in proprietary domain models and its own inference stack, which Management cited as materially more cost-efficient than standard cloud alternatives. This infrastructure is designed to support AI use cases that must operate within structural, safety, and regulatory constraints, combining probabilistic inference with deterministic outcomes. Management framed this as critical to scaling AI responsibly in physical, regulated environments.

Axonius

We hosted Axonius' Joe Diamond (President & Interim CEO) and Wendy Zheng (CFO). Axonius provides aggregates and normalizes data across an organization's security and IT tools to create a unified view of all assets, including devices, users, and applications. The platform is used to identify gaps in coverage, prioritize exposures, and automate response actions across the existing stack.

- **The core problem remains limited visibility across assets.** Most enterprises still lack a clean, real-time view of their asset base, and AI is driving a step-function increase in complexity through non-human identities, agents, and ephemeral infrastructure, making visibility and control more critical than ever.
- **Axonius is positioning itself as the independent data layer across a fragmented stack,** with its differentiation tied to aggregating and reconciling data across systems into a single source of truth and enabling action back into those systems; recent M&A, including ServiceNow's acquisition of Armis, reinforces the importance of asset visibility and exposure management as a core control point.
- **From here, the company is focused on expanding beyond visibility into orchestration and automation,** including enabling more autonomous workflows and AI-driven actions, while broadening coverage across additional asset classes (e.g., identities, SaaS, OT/IoT) and pushing further into both large enterprise and mid-market adoption.

Baseten

We spoke with Tuhin Srivastava (CEO & Co-Founder) of Baseten. Baseten is a production inference platform purpose built to deploy, run, and scale large AI models

reliably and cost efficiently, with deep investment at the infrastructure and runtime layer to serve high-throughput, low-latency workloads in production.

- **Acute compute scarcity reinforces the strategic value of inferencing platforms.** Management emphasized that supply constraints are materially worse than what is commonly understood from the outside. Scaling inference capacity today is constrained by OEM-delivered AI systems (e.g., from Dell and Supermicro), with 9–12 month lead times and capacity increasingly tied to multi-year contracts with significant upfront prepayment. Baseten currently operates inference clusters ranging from ~40K–80K GPUs (primarily Hopper and Blackwell) and noted that securing incremental capacity has become exponentially more capital-intensive, citing examples where ~5K additional GPUs imply ~\$100–125M of annualized commitments – roughly an order of magnitude higher than comparable capacity a year ago. As models improve and use cases become more agentic, Management noted inference demand continues to grow faster than supply, widening the gap between what customers want to deploy and what infrastructure can support. A meaningful portion of Management’s focus is now on financing and structuring this capacity responsibly, underscoring how inference has become both the bottleneck and the control point in the AI stack.
- **Inference economics are increasingly determined at the runtime layer as production workloads scale.** Management spent considerable time discussing how cost, latency, and reliability in production are increasingly driven by inference-time architecture rather than model choice alone. As models grow larger and workloads become more agentic, bottlenecks shift toward memory management, caching, scheduling, and orchestration, with meaningful cost dispersion even when serving the same underlying model. One high-signal point was that verifying outputs is materially cheaper than generating them, enabling architectures where larger models are paired with smaller routing or verification models to improve throughput and reduce cost per request. The broader takeaway is that control over the inference runtime is becoming a primary lever for sustaining unit economics at scale, particularly as customers deploy a mix of proprietary, custom, and open-source models.
- **Management sees exponential growth ahead as inference becomes a core control layer in the AI stack.** Management estimated current annualized inference spend across deployed models at roughly ~\$1.5–\$2B, which they view as an early snapshot relative to the scale of AI usage they expect over time. As AI becomes embedded across software workflows, search, support, and increasingly agentic systems, inference demand is expected to expand rapidly alongside usage intensity. In that environment, Management believes platforms that govern latency, availability, and cost across heterogeneous model stacks take on greater strategic importance. Management highlighted that production inference places fundamentally different demands than training workloads – favoring always-on, latency-sensitive, and distributed architectures over centralized, batch-oriented systems – reinforcing why inference-native platforms like Baseten are positioned to capture durable value as AI moves deeper into revenue-critical use cases.

We hosted Bright Machines' CEO Sviat Dulianinov. Bright Machines specializes in software-defined manufacturing. Founded in 2018, the company aims to modernize the electronics assembly industry by replacing traditional, rigid manual labor and fixed automation with flexible, AI-powered robotic systems. As of April 2026, Bright Machines has positioned itself as a critical infrastructure provider for the AI backbone, specifically focusing on the automated assembly of high-value AI hardware, such as data center servers, GPUs, and networking equipment.

- **AI infrastructure is a throughput problem, not a labor story (yet).** In AI infrastructure, the bottleneck is time-to-capacity, cutting time-to-market by even a third can translate into enormous downstream value, and the company's pitch is that robotics + software can ship faster with higher quality while reducing mechanical failure modes that force costly swaps in the field. Labor substitution is in the background, the near-term value is throughput unlock + quality under load.
- **Localized manufacturing becoming increasingly common.** Mr. Dulianinov highlighted a shift away from centralized, labor-dependent manufacturing toward distributed, localized production - micro-factories deployed near demand (including sovereign datacenter builds), reducing dependency on long global supply chains and shrinking lead times by moving assembly closer to where capacity gets installed. By building where the customer is, you can iterate faster, and treat production capacity as something you can replicate and relocate.
- **The data moat is traceability → learning → better builds, not just automation.** Bright Machines is accumulating a specialized dataset around how complex electronics are actually assembled, full traceability down to what went into a unit and how it was built, then layering AI on top. The compounding advantage isn't a single robot, it's the closed loop between production data, defect detection, and continuously improving process control in a narrow, high-value domain.

We hosted CrowdStrike's CFO Burt Podbere.

- **CrowdStrike views advances in frontier models as a meaningful tailwind to its business** given its strategic positioning to protect, with the potential for AI agentic security to be bigger in scale than CrowdStrike's cloud business (\$800mn ARR today, growing 35%). In addition to heightened focus on vulnerability management, CrowdStrike noted that AI gives attackers more speed. Customers are asking for visibility (EDR), control (guardrails, regulatory) and auditability (an understanding of what employees are using AI for, to in turn control costs). In response, CrowdStrike has compounded its core IP with technology acquisitions to be able to stop breaches across multiple planes, including new AI planes. Specific examples where CrowdStrike is benefiting from increased business momentum includes Pangea (visibility into AI agent attacks, along with CrowdStrike's threat intelligence), as well

as its identity module (ITDR), just-in-time PAM, Shield (for SaaS applications), Seraphic for browser protection, and SGNL for just in time authentication. CrowdStrike noted LLMs are incentivized to work with security companies (rather than build their own security products) to in turn lower friction points to their own adoption in a broader general purpose AI TAM.

- **On budget, CrowdStrike noted that business users will get involved to release budget for IT to enable AI to make knowledge workers more productive, with some trade off between hiring plan and budget for inference and AI tools,** given AI is a strategic imperative. At the same time, Flex is at the center of these conversations to help customers maximize feature usage; with CrowdStrike being positively surprised by “reflex” activity in particular (customers increasing flex commitments intra contract). On the question of the durability of Flex adoption in the top quartile of CrowdStrike power users vs. the entire installed base (i.e. is the low hanging fruit being captured?), CrowdStrike reiterated its confidence that Flex will have a meaningful impact throughout the customer base.
- **On next gen SIEM,** CrowdStrike noted its architectural advantage centers on its ability to tag data with context, which would otherwise not be captured in a data lake (even if the volume of data in competitor data lakes is greater). CrowdStrike also has a speed and cost advantage with querying in part because of its index enhancements (e.g. LogScale is a temporal database with a probabilistic data filter). CrowdStrike believes this advantage will hold even vs. new entrants like **Databricks**.

Please see our takeaways from our 4/8 conversation with Adam Gellar, SVP of Counter Adversary, [here](#).

Decagon AI

We hosted Decagon AI’s RVP of Strategic Accounts Katie Coutermarsh. Decagon AI is an enterprise platform that enables companies to build, deploy, and scale AI agents for automating customer support across voice, chat, email, and SMS with centralized control and orchestration.

- **While customers look to Decagon for cost savings, the more compelling proposition is centered around increasing revenue.** Call center interactions can cost ~\$18 per call, whereas Decagon can deliver higher-quality experiences for under \$1, thus the value proposition has been clear for customers. However, rather than focusing on just cost savings, customers are more focused on the ROI framed around revenue uplift (higher LTV, upsell, and conversion).
- **Speed of iteration structurally favors AI-native platforms over incumbents.** Decagon’s AI-native architecture and absence of legacy tech debt enable materially shorter iteration cycles (measured in weeks rather than quarters). More importantly, customers can self-serve changes through the platform via natural-language AOPs (Agent Operating Procedures), without needing to rely on Decagon’s engineering team or SIs. By contrast, incumbents remain constrained by services-heavy delivery models, IT gating, and fragmented stacks, where incremental changes often require

bespoke workstreams and extended production loops, slowing time-to-value and limiting responsiveness as use cases evolve.

- **Build vs. buy dynamics are shifting toward buy.** While many enterprises initially explored building in-house during the early days of AI which still remained the case a year and half ago, momentum has since moved toward “buy” as customers recognize that success requires far more than a point-in-time agent. Operating production-grade conversational AI agents demands ongoing investment in testing, safety, model upgrades, and orchestration, which requires engineering teams to work on the project indefinitely. Customers are now increasingly questioning whether bespoke agents are truly core IP and whether allocating top engineering talent is justified, as opposed to partnering with platforms like Decagon that have already solved this complexity at scale.
- **Agents are alleviating labor constraints in call centers.** Call centers face persistent challenges staffing for volatile demand peaks (e.g., Black Friday) and managing ~50% attrition rates, making it difficult to scale headcount efficiently. Agents are increasingly used to absorb routine tasks and volume surges, reducing the need to staff for peak capacity times while improving resiliency. As AI handles repeatable tasks, human agents are redeployed toward higher-value, relationship-driven interactions, repositioning CX teams as consultative partners.

Fivetran

订阅研报添加微信 LCD123321Cad 04-16

We hosted Fivetran’s CEO George Fraser. Fivetran is a fully managed data integration platform that automates the ingestion of structured and semi-structured data from source systems into cloud data warehouses, enabling reliable, governed analytics at scale with minimal ongoing maintenance.

- **The company views customer freedom to move between warehouses and analytics engines as structurally positive, arguing that attempts to trap data ultimately erode trust and retention.** Instead, Fivetran aims to be the priority ingestion and curation layer regardless of where workloads ultimately run. The need to reliably aggregate, organize, and continuously update internal data has become mission-critical. Fivetran increasingly sits at a critical point in the modern data stack, supplying curated enterprise data that AI systems require to reason and take action.
- **Fivetran does not view foundation model providers as credible infrastructure competitors.** Management was direct in stating that frontier model companies are unlikely to recreate deeply embedded data infrastructure categories like ingestion, transformation, and curation. While models excel at reasoning and code generation, the company sees a clear separation between intelligence and the operational systems required to reliably move and maintain enterprise data in production. As such, Fivetran sees AI as a demand tailwind rather than a displacement risk.
- **The company believes offense, not defense, is the right response as software barriers collapse.** With the cost of building software falling sharply, management argued that traditional defensive moats in data infrastructure have weakened.

Fivetran's strategy is to stay on the offensive—continuously expanding coverage across sources, formats, and use cases—rather than relying on switching costs or legacy entrenchment. While AI is increasingly capable of extracting value from messy data, companies still need leadership-driven mandates to centralize and curate data before it can be meaningfully operationalized at scale.

Glean

We hosted Glean's VP of Finance and Business Operations Michael Miao. Glean is a Work AI platform that connects and understands all of enterprise data, enabling employees to find answers, automate tasks, and generate content grounded in company knowledge. Glean also empowers businesses to securely build multistep prompts and custom AI agents that autonomously plan and execute workflows with natural language.

- **Enterprise AI is constrained by context, not model capability, and Glean is positioned at that bottleneck.** While frontier models continue to improve, their primary constraint in enterprise settings is secure, governed access to proprietary context across fragmented systems. Enterprises intentionally prevent frontier model labs from ingesting internal data, creating a need for an independent context and orchestration layer. Glean's value proposition is solving this problem by unifying enterprise context across SORs, enforcing permissions and governance, and safely delivering relevant inputs to models.
- **Differentiation is driven by depth of retrieval, platform neutrality, and cost efficiency as use cases become agentic.** Glean's core product enterprise search and its investment in hundreds of secure connectors enable high-precision retrieval, which management views as superior to relying on any SORs. This cross-application context becomes increasingly important as customers move from simple Q&A toward multi-step, agentic workflows that both retrieve information and take action. Platform and model neutrality were highlighted as key buying criteria for large enterprises, allowing customers to avoid lock-in as the "best" model changes. Importantly, better retrieval reduces unnecessary tokens, lowering inference costs and mitigating margin pressure even as workflows become more sophisticated.

Grafana

We hosted Grafana's CFO Wailun Chan, VP of Engineering Manoj Acharya, VP of Emerging Product Marc Chipouras, and VP of Finance Robert Ni. Grafana is a leading open source observability provider offering a unified platform for monitoring and visualizing metrics, logs, and traces.

- **Management stated that open source has a structural advantage vs. closed-source vendors like Datadog by enabling customers to ingest, query, and visualize data from any source without forcing consolidation onto a single proprietary stack.** For example, for customers to fully leverage the value of Datadog's platform, it would require migrating all of a customer's disparate

observability data into its platform, creating vendor lock-in and duplicate spend for customers running multiple tools. Grafana emphasizes interoperability and cost efficiency, allowing customers to unify diverse observability environments without paying twice.

- **Grafana aims to deliver TCO reduction by materially shrinking data volume.** For example, Adaptive Traces reduces 90% of ingested traces while Adaptive Profiles enables ~85% reduction in profiles, directly addressing the “Q” side of the TCO cost equation. On the other hand, agent-driven workflows are driving 5–20x higher query loads vs. humans as agents consume broader datasets more frequently, making cost efficiency critical. Grafana argues its knowledge graph and adaptive controls allow agents to extract more value from observability data at lower ingest rates, increasing platform stickiness as customers rely on more of the stack without proportional cost inflation.
- **Focus shifting from TCO to open standard.** Grafana notes a clear shift in buyer priorities over the last 18 months, moving from a primary focus on TCO toward open standards and interoperability. While earlier conversations centered on avoiding runaway spend as observability bills scaled, cost has become more secondary as customers plan for an agentic world and push back on vendor lock in. Across the company’s ~9K customers, none operates on a single vendor, and consolidation has proven difficult. This elevates the importance of open standards that allow agents and users to reason across fragmented observability stacks.
- **Grafana sees Datadog as its clear primary competitor, with New Relic, Dynatrace, and Splunk representing a second tier of competitors. ClickHouse, by contrast, is largely complementary.** Management noted that Grafana overlaps with Datadog >50% of the time, while Clickhouse has low single-digit competitive overlap and customers commonly deploying both platforms side by side for distinct workloads. Grafana remains differentiated in time-series and real-time observability use cases, while ClickHouse is more often used for longer-term, historical, and columnar analytics, limiting direct competition today. Looking ahead, management views the market as evolving toward broader IT observability with early signs of convergence between observability and business analytics.

Guidewire

We hosted Guidewire’s CFO Jeff Cooper.

AI is driving more urgency for cloud conversion activity, and making it easier for insurers to move off of mainframe systems. The largest benefit to Guidewire from AI this year is accelerating cloud conversion activity, both through existing customer migrations and new customer wins. Recall that well over half of total P&C industry premium is still supported by on-premise or mainframe core systems, and customers are increasingly citing AI as a driver for accelerating cloud conversion plans. This is showing up, too, with larger deal sizes, longer duration deals, and bundling Claims with Policy in the initial sale at greater rates. While the migration itself is still difficult given pulling data from mainframe into the cloud requires more than just making code ready for the cloud,

management noted that it is deploying Frontier Model capabilities to its product development and professional services teams to reduce cost/time to implement. Management is targeting a ~50% cost reduction for migrations over time (they are not there yet).

AI framed as purely TAM expansive; competition is forming outside of Guidewire's existing revenue base. AI-natives to this point haven't shown signs of challenging Guidewire's system of record. There continues to be experimentation with insurers building agentic functionality or working with forward deployed engineer businesses for certain use cases, but that activity currently exists in areas that are purely TAM expansive for Guidewire like Underwriting (rather than in areas that could be value destructive). Federato is one business that has shown signs of trying to connect Underwriting with Policy Administration, but there haven't been proof points that the Policy side has gotten traction. As for the debate around whether there could be a unified data layer that forms on top of system of record components where agentic use cases take place, such a data layer would not solve underlying infrastructure issues for customers not fully on Guidewire cloud (i.e. if a customer used Guidewire for Claims/Policy and had a mainframe billing system), and ultimately could be viewed as duplicative of Guidewire's functionality (which Guidewire mentioned in one case had already happened).

Despite tougher comps in 2H, strong visibility into dense pipeline (which includes some large deals). Management framed the pipeline as strong with high visibility into commitments and increasing win-rates/close-rates among "best case" prospects (i.e. portion with real interest but not quite at commitment). The company typically leans into this visibility to a degree in guidance but does not assume the high win rates of recent quarters continue, which underscores the upside risk to numbers near-term should momentum continue at rates we've seen in recent history. Management reiterated Guidewire's commitment to grow Fully Ramped ARR faster than ARR for FY26, which in our view is encouraging given the tough comp from the Large Tier 1 insurer win from F4Q25.

Hightouch

We hosted Hightouch's Co-Founder & Co-CEO, Kashish Gupta, and CFO, Cherry Miao. Hightouch is a composable customer data platform (CDP) that syncs warehouse-native data directly into downstream business tools, enabling teams to activate customer data across marketing, sales, and product without duplicating or moving data out of the warehouse.

- **Hightouch is positioning itself around "on-brand" AI rather than generic LLM output.** Management was explicit that many CMOs were burned by early LLM experimentation in 2024, citing off-brand, unshippable outputs that could not be deployed into production workflows. Hightouch believes its differentiation lies in serving proprietary customer, brand, and asset context to foundation models, allowing LLM-generated content to follow brand guidelines and business logic. The company increasingly views defensibility through verticalized use cases and exclusive access to first-party context rather than model ownership.

- **AI is driving a step-function increase in marketing leverage, not just incremental productivity gains.** Management framed the platform's impact as enabling small, strategic marketing teams to do the work of materially larger organizations, with agentic workflows increasingly handling campaign creation, personalization, and execution. In this model, marketers spend less time writing briefs and stitching tools together, and more time directing AI "coworkers," with some customers already seeing materially higher campaign performance versus non-AI workflows.
- **The company sees consolidation at the workflow layer, not the data layer.** Management sees a shift toward fewer, strategic platforms that sit above data systems and orchestrate work via agents, while warehouse-native architectures become mission-critical, as marketing moves from UI-centric tools to infrastructure. CMOs increasingly want a single agentic layer spanning campaigns, channels, and decision-making, rather than maintaining dozens of specialized point solutions. Hightouch aims to be that orchestration layer, integrating with (rather than replacing) downstream tools.
- **Pricing and economics are deliberately aligned with customer ROI, not near-term gross margin optimization.** Management emphasized that Hightouch's pricing has always scaled with personalization volume (campaigns), enabling revenue to grow alongside customer success. While AI introduces incremental cost, leadership is prioritizing long-term value creation over short-term GM maximization. Today Hightouch has a 9-figure ARR business with 120%+ NRR, >90% GRR, and a burn multiple below 0.5.

← x b¥mûR_@Oá LCD123321Cad 04-16

Huntress

We hosted Huntress's CEO, Kyle Hanslovan. Huntress provides a managed security platform focused on SMB and mid-market customers, combining endpoint, identity, and SIEM capabilities with a 24/7 human-led SOC. The company emphasizes delivering outcomes through a fully managed service, helping customers detect, investigate, and respond to threats without requiring in-house security expertise.

- AI adoption is showing up very differently across the market. Large enterprises are actively experimenting and pushing usage, while the SMB base is still much earlier and largely looking for someone to manage outcomes for them.
- On the product side, they were direct that a lot of the "AI replaces the SOC" narrative is overstated. Automation is improving, but real investigation and response still require human judgment, especially as attack paths get more complex.
- The bigger shift is around how security is measured, less about blocking everything upfront and more about how quickly you can detect and respond, though near term they do think attackers are benefiting more from AI than defenders.

← x b¥mûR_@Oá LCD123321Cad 04-16

We hosted Intercom's CFO, Dan Griggs. Intercom is an AI native customer service platform that pairs an AI powered help desk with its integrated support agent, Fin, to automate and manage customer support across email, phone, live chat, SMS, and other channels.

- **Intercom highlighted a dramatic difference in cost savings from shifting away from 3P models and toward proprietary models built in house.** The company has moved from relying on 12+ 3P models to a mix that is now primarily custom-built, which has driven meaningful gross margin expansion. While Intercom continues to use select 3P models for areas that have yet to be rebuilt or where performance remains superior, management expects the model stack to migrate fully in house over time. This transition has created a structural cost advantage vs. competitors tied to 3P token pricing, giving Intercom flexibility to expand margins and reinvest in performance as usage scales. Importantly, custom-built proprietary models have also delivered meaningful performance improvements, with recent examples including a custom rebuilt reranker that delivers higher answer quality, an estimated 80% reduction in reranking costs, and greater product flexibility.
- **The application and software layers surrounding models and agentic functionality are becoming cheaper and easier to build, pushing durable differentiation toward the underlying model technology.** Intercom views it as inevitable that competitors will also need to develop proprietary models, both to remain competitive on cost and to avoid long term margin pressure from 3P, token based pricing. As agentic tooling becomes increasingly commoditized, sustainable advantage is shifting to domain specific training data and the ability to deliver superior accuracy, latency, and trust in regulated, multi step workflows. Beyond cost, management believes that as inference costs fall, pricing power will increasingly be driven by performance, with proprietary models enabling greater monetization per unit of compute and supporting durable competitive advantage.
- **Intercom prices Fin directly against customer service labor economics, driving a clear ROI led budget mix shift toward outcome based spend.** Customer service teams typically spend ~\$5–6 / conversation / ticket on average, with a human agent handling ~300–400 conversations / month and ~\$49–\$149 per seat in traditional software costs. By comparison, resolving ~300 conversations via Fin costs roughly \$300 for \$0.99 / resolution while effectively eliminating the associated labor costs. While management notes that all “seat mover” spend is theoretically at risk over time, direct cannibalization of help desk seats has been limited to date. Instead, customers are using Fin to absorb incremental growth in conversation volume and redeploying labor capacity toward higher value activities such as customer success, upsell, and cross sell, though management acknowledges that longer term help center economics are likely to evolve as agentic AI adoption deepens.

We hosted Island Technology President Steve Tchejeyan. Island provides an enterprise browser platform that embeds security controls directly into the browser, enabling organizations to manage application access, protect data, and enforce policy at the point where users interact with corporate systems. Island believes its approach to SASE solves for latency, performance and security challenges in a more sophisticated way than traditional SASE vendors, and that it is well positioned to take share from SASE incumbents.

- Island is positioning the browser as the primary control point for modern security, arguing that since the browser is where users access applications and data, it is the most logical place to enforce policy, visibility, and data protection across both SaaS and legacy environments.
- The core problem is shifting to enabling AI safely - enterprises are under pressure to adopt AI tools quickly, but lack a clear way to control how sensitive data is accessed and shared across models, driving demand for a control layer that allows “safe enablement” rather than outright blocking.
- The platform is expanding beyond browser security into a broader control plane for application access and AI usage, including enabling multiple LLMs within the enterprise, enforcing data boundaries, and supporting agentic workflows, while increasingly displacing legacy infrastructure like VDI, VPN, and parts of the SASE stack.

KnowBe4

We hosted KnowBe4's President & CEO, Bryan Palma. KnowBe4 is a human risk management platform focused on security awareness, phishing simulation, and email security, helping organizations train employees and reduce social engineering risk. The company is expanding into adjacent areas, including agent risk management, as it evolves into a broader platform addressing both human and AI-driven attack surfaces.

- KnowBe4 is repositioning from security awareness training into a broader “human risk management” platform, expanding into phishing simulation, email security (via its Egress acquisition), and now agent risk management, with the goal of addressing both human and non-human attack surfaces.
- The company's differentiation is anchored in its data and scale, with 15+ years of telemetry on how users interact with threats and a large installed base that enables more effective training, simulation, and increasingly agent-driven personalization of security workflows.
- AI is being used both to enhance the platform (agent-led training, content generation, risk scoring) and as a new attack vector (agents as a new vulnerable surface), with management leaning into this shift by building capabilities around agent governance while continuing to drive toward a multi-product platform and re-accelerating growth back toward ~20%.

We hosted Kong's CTO & Co-Founder, Marco Palladino. Kong is an API and AI connectivity software company providing a unified platform to manage, secure, and govern traffic across APIs, services, and AI workloads, enabling enterprises to operationalize AI at scale with consistent policy enforcement, observability, and cost control across models and agents.

- **Kong's evolution reflects a pivot from an early API marketplace to critical enterprise infrastructure as APIs proved foundational.** Founded in 2017 after experimenting with an API marketplace model, Kong realized the market was too early and open sourced its core gateway technology. That decision catalyzed adoption as APIs moved from experimentation to production infrastructure, allowing Kong to scale into a platform supporting APIs, microservices, and now agent connectivity across global enterprises.
- **Enterprises are moving from AI experimentation to operationalization, driving demand for governed API and agent infrastructure.** Kong works with highly regulated enterprises that need AI, APIs, identity, security, and access control managed centrally. While many organizations are spending on models and agents today, usage remains peripheral rather than core due to trust and non-determinism. Kong positions itself as the control plane to move AI into mission-critical workflows with outcome guardrails, visibility, and the ability to shut agents down when needed.
- **Core API demand and AI use cases are deeply intertwined, with AI now a major growth driver.** Every agent built at scale requires APIs at scale, and Kong's pipeline has doubled as customers deploy AI platforms. Today, ~50% of pipeline is AI-related, net revenue retention is 120%+, and hundreds of organizations are using Kong to manage connectivity for agents, models, and tools. Customers with mature API investments are seeing materially more value from AI than those without.
- **Kong is positioning for the agentic economy by owning the connectivity, governance, and monetization layer.** Kong supports MCP and A2A as core protocols, enabling agents to discover tools, stream data, and interact with other agents. Beyond governance, Kong is extending into monetization and memory optimization, leveraging its position between every API and AI call. Its platform can run fully managed, hybrid, or on-prem, serving risk-averse enterprises while enabling new agent-to-agent economic models where APIs become products and revenue streams.

订阅研报添加微信 LCD123321Cad 04-16

We hosted Okta's President & COO, Eric Kelleher.

- **Building an "identity fabric" to secure AI agents end-to-end.** Okta continues to invest in its platform with the goal of building a unified control plane where customers can monitor all identities, including AI agents, through a single console. The company outlined core priorities for securing agents including visibility into

where agents exist, control over what systems and data they can access, and governance over the action they are permitted to take. A key milestone has been driving standards around how agents interact across applications, which has been attained through cross agent access – an important step toward reducing friction and increasing adoption of identity solutions. Okta already supports AI agent use cases for customer identity through Auth0 for AI agents, with workforce capabilities for AI agents set to go live on April 30.

- **Okta noted its cloud-first identity stack has minimal overlap with the types of code security vulnerabilities found on premise or in endpoint technology (we heard similar thoughts from Zscaler).** That said, the broader focus on securing identities continues to drive conversations, supporting ongoing focus on progress in its ITP and PAM solutions. The company also discussed applying LLM capabilities internally, aggregating signals across its own platform and third-party sources to contextualize signals across an organization.
- **Near-term financial impact from agents limited; GTM and competitive backdrop stable.** Management expects more meaningful impact in FY27. The company continues to invest in a more specialized sales force to better serve both developers and enterprise buyers, while also highlighting ongoing improvements in sales productivity following steady ramp over the past several quarters. From a competitive standpoint, management noted no meaningful change in dynamics, including limited impact from Palo Alto Networks' acquisition of CyberArk. Similarly, it does not view Microsoft's bundling of agentic security into E7 as a significant risk, citing comparatively narrower capabilities. Okta's two priorities remain to not get hacked, and to grow faster.

Outreach

We hosted Outreach's CFO, Drew Laxton. Outreach is a sales execution and engagement platform designed to orchestrate and automate end-to-end revenue workflows, combining signal-driven analytics with emerging agentic capabilities to help sellers move opportunities through the funnel more efficiently.

- **Outreach is moving beyond copilots toward fully agentic, workflow-driven sales execution.** Management framed the next phase of product evolution as a shift from assistive copilots to task- and workflow-level agents that can autonomously move prospects through the funnel, informed by first- and third-party data, prior interaction history, and internal CRM signals. Rather than optimizing for surface-level productivity gains, Outreach is aiming to become the system of action at the moment an opportunity emerges, recommending — and in high-volume use cases executing — next best actions on an AE's behalf.
- **AI is structurally reshaping sales org design, not just rep productivity.** Outreach highlighted that customers are largely trying to hold sales headcount flat while driving incremental output through AI, with agents already delivering a significant productivity lift per rep. Over time, this enables higher manager-to-rep ratios, reduced reliance on RevOps analysts, and a shift in sales management toward data

interpretation rather than deal intervention. Management expects traditional concepts like rigid sales stages to erode as agentic systems operate continuously across signals instead of discrete pipeline checkpoint.

- **Pricing and ROI are migrating away from seats and tokens toward action-based outcomes.** While Outreach still prices primarily per user, management underscored a growing mix of credit- and consumption-based packaging tied to actions rather than raw token usage. Importantly, Outreach is focusing on customer value, grounding ROI in whether AI improves top- or bottom-line outcomes (e.g., deal velocity, win rates, coverage). Pilots remain central to both adoption and pricing confidence, allowing customers to tune how much decision-making authority agents are granted, particularly in higher-touch outbound motions.

Physical Intelligence

We hosted Physical Intelligence's Co-Founder & CEO, Lachy Groom. Physical Intelligence is an AI research company building foundation models for the physical world, aiming to enable general purpose robots that can learn, reason, and act across a wide range of real world tasks. The company focuses on unifying perception, control, and learning to move beyond narrow robotics systems toward more adaptable embodied intelligence.

- **Generalization key to specialization.** Management challenged the conventional bias toward specialized robotics, arguing that generalization is a prerequisite to downstream specialization. Drawing explicit parallels to the evolution of language models, the company argued that training a single model across highly diverse tasks leads to superior performance even on narrow use cases, while task-specific robots ultimately cap learning and economic relevance. Management noted that this stance underpins Physical Intelligence's rejection of warehouse-only or humanoid-only approaches and reinforces its belief that broad embodiment across robots, tools, and environments is the only viable path to scalable labor replacement.
- **Long-term moat is proprietary, deployment driven data.** The company emphasized that locomotion and navigation are largely solved through simulation-driven reinforcement learning, while real-world manipulation (e.g., deformable objects, food preparation, tool use) remains fundamentally unsimulatable and therefore data-constrained. As a result, Physical Intelligence is orienting its research and capital toward tasks where real-world interaction is unavoidable and where labor displacement potential is highest. Management detailed a closed-loop system in which human-operated robots generate video and action data for specific tasks, models are retrained and validated on live robots, and subsequent deployments generate incremental proprietary data. Critically, this proprietary dataset improves with each deployment, reinforcing management's view that data accumulation is the primary competitive advantage.
- **Capturing large labor opportunity.** Management discussed prioritizing the structurally large labor opportunity rather than near-term monetization. In order to capture this opportunity, Physical Intelligence emphasized that competitiveness will

require foundation-model-level scale in both compute and data – in fact, management argued that internet-scale data will become increasingly insufficient to train models, highlighting that more real-world data is generated every minute relative to the entire internet corpus. That progress over the next several years will require operating at similar compute scale to large foundation models. Against this backdrop, management framed physical intelligence as a long-duration, capital-intensive effort aimed at unlocking an addressable labor market opportunity of significant magnitude, with monetization viewed as a downstream outcome once capability thresholds are reached.

订阅研报添加微信 LCD123321Cad 04-16

Salesforce

We hosted Salesforce's EVP & GM of Agentforce, Madhav Thattai.

- **Model choice should be task-specific, not frontier-first.** While large reasoning models are useful for deep discovery, coding, and long-form reasoning, they are inefficient for deterministic workflows such as a “checklist-like conversation” as they add cost with little incremental value. Older or smaller models deliver better latency and cost for many customer-facing tasks. As a result, Salesforce takes an “opinionated harness” approach to its reasoning chain, selectively combining LLMs, smaller fine-tuned models, and deterministic logic to optimize speed, cost, and efficiency, measuring success through tokens per Agentic Work Units (AWUs).
- **Salesforce is monetizing Agentforce at the run-time layer, focusing on capturing value from agentic execution rather than consumption.** Management emphasized its approach to price at the highest level of abstraction, aligning monetization with customer outcomes instead of inputs like tokens, while retaining flexibility to meet customers at different stages of maturity. In practice, this shows up across multiple pricing motions: 1) per-seat licensing with unlimited agent access for employee-facing productivity use cases; 2) consumption or action-based pricing for externally facing and operational workflows such as candidate qualification for Adecco; and 3) AELE contracts for customers pursuing multimillion-dollar transformations where flexibility matters most. Salesforce continues to iterate upon pricing to refine higher-order metrics like conversions or leads generated. Management further noted that what Salesforce monetizes is not just task execution, but the Agentforce harness including context, metadata, governance, trust, and observability that underpins enterprise-grade outcomes at scale.
- **API access is gated rather than unfettered, which creates a monetization lever through premium SKUs.** Salesforce has not seen a material increase in API usage or data egress to date, and even when customers do port data out, the primary friction is the need to reconstruct Salesforce's metadata layer to make the raw data usable. That added complexity raises switching costs and reduces the practicality of rebuilding agentic workflows outside the Salesforce ecosystem.
- **Salesforce positions Agentforce as a scalable, enterprise-grade platform**

订阅研报添加微信 LCD123321Cad 04-16

relative to point-solution competitors and closed AI ecosystems. Management noted that it has won head-to-head deals against Sierra and reclaimed customers that initially deployed single-point POC agents but returned to Salesforce to develop agentic decisions at scale. Salesforce is able to implement agents and go-live within a similar timeframe as that of Sierra (~8 weeks). Relative to competitors such as Sierra and Decagon that rely on bespoke agents per customer, Salesforce emphasized its ability to scale horizontally across 20K+ Agentforce customers. In some cases, Salesforce's full-modality support via SMS has enabled share gains in overlapping accounts with Sierra. Management also cautioned that outcome-based pricing can be problematic given inconsistent definitions between vendor and customer. On the question of building bespoke with frontier labs or hyperscalers, Salesforce underlined that customers will need to recreate metadata, workflows, and governance at scale, coupled with higher risk of vendor lock-in. Customers who choose to build in frontier model ecosystems need to weigh the trade-offs of being bounded to every part of their respective proprietary harnesses.

- **Data readiness, control, and change management are the main inhibitors to adoption.** Many customers are held back by legacy data debt and incomplete architectures, limiting what agents can do effectively. Beyond data, consistent agentic control and accuracy are critical, with enterprises focused on reproducibility, trust, and governance over time. ~95% of the effort comes after go-live, centered on observability, KPIs, and continuous tuning to ensure agents perform as intended and deliver business outcomes. As customers reach production scale, change management and run-time operations, not build speed, have become the primary inhibitors to adoption.
- **Opening agent development to customer-preferred environments can accelerate adoption and deployment.** For enterprise customers, Salesforce noted that deployments have been FDE-heavy, which has constrained velocity and limited how quickly customers can expand use cases. Management thus sees enabling customers to build agents in tools like ones offered by frontier AI labs expands the developer base and speeds execution. This shift allows small admin teams to support high-volume use cases while unlocking far greater scale as more internal engineers participate. For SMB and mid-market customers, out-of-the-box capabilities already enable faster time-to-value, and the addition of hybrid reasoning has driven deployments toward a roughly 50/50 mix of custom and out-of-box agents.
- **Slack has emerged as the breakout interface driving Agentforce adoption and usage.** Management highlighted Slackbot as the fastest-growing Agentforce deployment by AWUs, attributing to its strong product-market fit. Salesforce views Slack as the most natural UI for agentic work alongside knowledge workers, positioning it as the primary surface where agents are managed, decisions are made, and core work gets done.

We hosted Dominic Phillips (CFO) and Mike Chang (Head of IR) of Samsara.

- **Fuel inflation is sharpening ROI-driven customer engagement and creating a repeatable sales motion.** Management noted they are not seeing demand pressure from higher oil prices, highlighting that many customers have operated through volatile fuel environments for decades. Instead, rising fuel costs are pushing customers to scrutinize operating inputs and efficiency more closely. In response, Samsara has rolled out and packaged capabilities to help customers better manage fuel spend, including temperature-based idling reduction, routing and commercial navigation to optimize fueling locations, and broader fuel optimization tools. These offerings are being used to anchor both new and expansion conversations around near-term payback, turning macro volatility into a catalyst for engagement.
- **Management framed AI differentiation around the difficulty of replicating insights without network-scale, physically sourced data.** In addressing concerns that value could accrue primarily to LLM platforms, Management pointed to Samsara's ability to generate insights by learning across trillions of data points from roughly 12,000 customers and millions of connected assets accumulated over more than a decade. They noted that certain use cases, such as drowsiness detection and advanced safety insights, are difficult to produce using siloed, single-customer datasets. Importantly, these insights require physically instrumenting vehicles and assets and pulling real-world data into the cloud at scale, which Management views as a meaningful barrier to software-only or DIY approaches. They also emphasized that their customers typically lack large IT budgets and are seeking turnkey solutions from trusted vendors rather than building proprietary agents, while API access remains tiered and monetized without being positioned as a defensive control.
- **Management views autonomy as reinforcing, not displacing, Samsara's role across physical operations.** While adoption of autonomous vehicles remains early and fragmented, Management emphasized that customers want Samsara to serve as a single pane of glass across autonomous and non-autonomous vehicles, as well as tools and equipment in the field. They highlighted that roughly 85% of customers operate vans and field-service vehicles in urban and job-site environments, where utilization patterns and economics differ materially from long-haul autonomy use cases. More broadly, Management noted that about 50% of AI detections occur outside the vehicle cab today, supported by more than 40 computer vision models deployed across the platform. Examples include waste management customers using out-of-cab detections to validate proof of service, identify hazardous materials, and support billing and compliance workflows. Management suggested this supports continued platform relevance as vehicle operating models evolve, with customers focused on orchestrating and managing heterogeneous asset environments rather than optimizing solely around the driver.
- **Growth drivers are broadening while internal AI efficiency supports disciplined profitability and operating leverage.** Emerging products are becoming a more meaningful contributor to growth, driving roughly 23% of NNACV, as customers

increasingly land with broader product sets and more non-vehicle applications, including Asset Tags. Asset Tags expand the addressable asset base beyond vehicles and enable use cases around utilization, theft prevention, geofencing, and workflow automation, with adoption paced more by customers identifying assets and integrating data into workflows than by product capability. International NNACV remains roughly 15% of mix and is growing in-line with the US business, contributing to recent growth acceleration. At the same time, internal use of AI is driving productivity gains across R&D, marketing, and back-office functions, allowing slower headcount growth outside GTM while maintaining a fast release cadence, with incremental AI spend evaluated through a clear ROI lens and an emphasis on balanced growth and profitability.

Serval

We hosted Serval's Co-Founder & CEO, Jake Stauch. Serval is an AI-native IT automation company that builds autonomous agents to resolve IT issues by generating and executing code-based workflows across enterprise systems. It focuses on complex, real-world IT environments, positioning its agents to operate alongside existing platforms like ServiceNow.

- **Core differentiation is code-generated automation as the system of record.** Serval's thesis is that no-code ITSM tooling breaks at scale because it relies on rigid abstractions (inputs, triggers, field mappings) that introduce friction and require scarce IT time to maintain. Instead, Serval uses AI to generate production code as the source of truth, with workflows visualized only after the fact. Management emphasized that AI is well-suited to writing and reviewing code, fixing errors at runtime, and handling durability, permissions, approvals, and long-lived execution. Over time, generated code becomes reusable, enabling tens of thousands of automations to run without incremental inference cost. This architecture also provides an opportunity for CMDB data quality: currently CMDBs data hygiene is low driven by low incentives to keep ITSM data (relative to CRM or HR systems which are more consistently updated). Serval's workflow engine actually populates the CMDB within Serval. That is, it generates code to reach out to the appropriate systems and keep the CMDB in sync rather than relying on manual hygiene.
- **Prioritizing adoption over rip-and-replace.** Serval does not position itself as an overnight ServiceNow replacement - initial deployments often automate simple help-desk or onboarding workflows that start and end inside ServiceNow, while Serval mirrors the full ticket object and maintains its own copy of the data. This creates functional redundancy over time, with the goal of prompting customers to question why two platforms are required. Management framed displacement as a multi-year outcome (3 to 10 years), with the primary near-term constraint being organizational decision-making and process redesign rather than technical implementation.
- **Unit economics SaaS-like or better.** Today, Serval prices on a platform fee based on organization size and use cases, while tracking token utilization internally (but not

explicitly charging for it). Management cited two reasons for this: 1) avoiding friction in expansion beyond IT into HR, finance, legal, and sales; and 2) maintaining predictable IT budgets. Importantly, they noted that not all workflows require LLM usage, as reusable, pre-generated code allows many automations to execute with zero token consumption. As model costs decline and reuse increases, management sees a path to long-term gross margins 80%+, with a cost profile closer to SaaS than usage-based AI.

We hosted Shield AI's CEO, Gary Steele. Shield AI is a defense technology company that builds AI-powered autonomy software and drones used by military and national security forces. Its core platform enables aircraft and robots to operate and navigate in complex, GPS-denied environments without human control.

- **Technical differentiation centered on edge autonomy, coordination, and assurance.** Hivemind, the autonomous decision-making stack, is designed to operate fully air-gapped, reflecting real-world defense requirements. Management stressed the importance of decision-making at the edge, particularly in environments with degraded or denied communications, where platforms must continue to make autonomous, mission-appropriate decisions. Simulation was described as a core and required component of the autonomy lifecycle, especially for designing and validating coordinated behaviors across multiple autonomous systems (e.g., swarms), which cannot be realistically tested in live environments. Assurance, safety, and guardrails are built into the software stack to meet airworthiness and mission-level requirements. Customers can typically achieve initial autonomous flight within ~6 weeks, followed by continued refinement and progression to production deployment over ~6 months. This process is supported by Shield providing FDEs alongside customer developers, a key enabler of faster customer time-to-value relative to traditional autonomy programs (which typically require multi-year integration cycles).
- **Competitive Environment.** Shield both competes with and partners alongside incumbent defense primes, including cases where primes serve as subcontractors while Shield acts as prime on programs. Management emphasized the use of private capital to fund product development ahead of formal U.S. programs, allowing the company to avoid being constrained by slow procurement cycles or obsolete specifications. This strategy also enables Shield to sell directly to international customers and foreign militaries (~2/3 of revenue) while remaining compliant with U.S. government contractor requirements and export controls.
- **Revenue model and unit-economics.** Shield described a dual revenue stream across hardware and software, with the current business roughly 50/50 split between the two. On the software side, Hivemind monetization is driven by two components: 1) developer seat licensing for the Hivemind; and 2) an "AI pilot" fee, which management framed as ~5–15% of the value of the system or platform being flown autonomously. Management stated that leading indicators for scaling software

economics are the number of developer seats deployed and the number of AI pilots operating in production, pointing to an emerging recurring revenue base.

Sierra Technologies

We hosted Sierra's Co-Founder & CEO, Bret Taylor. Sierra is a conversational AI platform that enables companies to deploy AI agents to communicate with customers using natural language, solve problems, and take action.

- **Applied AI is fundamentally changing how companies are built and scaled.** Sierra scaled to \$100M ARR in 7 quarters and \$150M ARR in 8 quarters, with momentum continuing at a comparable pace. Management noted that this outcome exceeded expectations because AI agents ultimately prove more valuable than traditional software itself. Specifically, agents are digitizing labor across core business functions such as software engineering, customer service, and legal work, rather than delivering incremental productivity gains.
- **Speed and adaptability have become the defining sources of competitive advantage.** The pace of technological change is so extreme today that state-of-the-art capabilities are quickly commoditized, compressing the window in which any single innovation remains differentiated. In this environment, management believes that debating what constitutes as durable software IP to build matters less than executing relentlessly and delivering real solutions to customers. Management said Sierra's growth reflects a willingness to operate amid uncertainty, reject ivory-tower planning, and focus on rapid deployment, continuous iteration, and tangible business outcomes.
- **Moats in AI will be built at the solution and workflow layer, not at the model layer.** While the cost of creating software and accessing frontier models is falling rapidly, most enterprises will not build every software themselves. Durable value accrues to companies that productize complex workflows, integrate across fragmented enterprise systems, and align closely with business buyers. Over time, Management expects that the moat for traditional SaaS and IaaS will increasingly become distribution and delivering value by being close to the business problem as opposed to the underlying technology.
- **AI-native companies can strengthen margins by anchoring value to business outcomes.** Management believes that AI-native companies should not train their frontier models, and finetuning is sufficient. Running large models on every transaction is economically unviable, so leading teams increasingly rely on a mix of open-source, fine-tuned, and distilled models that are "good enough" for high-volume workflows. As frontier models have matured, advances now show up more through distillation and efficiency gains than step-function improvement in model capability, enabling lower-cost deployment. Critically, because applied AI is sold on measurable improvements like sales performance or customer satisfaction rather than tokens or inference volume, costs can be optimized independently of customer value, thus gross margins could strengthen.

We hosted Superhuman's CEO, Shishir Mehrotra. Superhuman is a productivity-focused AI suite and email-native platform designed to accelerate work by managing, drafting, and organizing emails using artificial intelligence. It features "Auto Labels" for automatic inbox categorization, "Autodrafts" for tone-matched replies, and "Go," a proactive assistant that integrates with calendars and external apps (like Slack, Jira, Salesforce) to handle tasks, schedule meetings, and provide insights.

- **Distribution is the primary adoption driver.** Management disclosed high engagement, with ~40mn DAUs, ~125bn LLM calls per week, and ~3–4k calls per user per day. The key takeaway is that ambient, in-context delivery inside applications materially outperforms destination-based chat or task interfaces, suggesting distribution at the work surface increases AI productivity.
- **"Go" reframes the platform from a single agent to an orchestration layer.** The transition from Grammarly as a standalone agent to a system of specialized agents (calendar, sales, support, fact-finding, etc.) was a structural shift. Superhuman aims to be the primary AI surface for users by orders of magnitude, driven by ambient deployment rather than explicit user prompts. The most tangible datapoint was internal Glean usage at Superhuman: standalone Glean averaged <1 query per user per day, but when embedded via Go, usage increased by 2 orders of magnitude as it reduced change management and brought Glean's capability where people work. This underscores Superhuman's role as the "last mile of AI" that materially increases utilization of existing enterprise tools.
- **The infrastructure is proprietary**, with ~97% of inference volumes handled on private, hosted, fine-tuned models, and only ~3% routed to third-party LLMs. This architecture is driven by both capabilities (with specialized models more effective at certain tasks) and latency requirements (as Superhuman's agents operate in milliseconds as a user types vs. multi-second responses from general-purpose models.) Moreover, management emphasized the depth of technical investment required to achieve application-level integration (e.g., emulating native clients to Google Docs in order to access full screen context), noting that while replication is possible, recreating a footprint spanning ~1mn unique app surfaces per day would be meaningfully difficult.
- **Assist vs. chat vs. do.** Management framed AI tools into chat (conversation-led), do (task delegation), and assist (contextual intervention), arguing that market attention has skewed toward chat and do. Superhuman is positioning itself within 'assist', inserting guidance at the point of work without interface switching. The key insight is that assist interactions become a high-leverage entry point for downstream chat or task execution, yet remain underserved by competitors.

We hosted Tekion Founder & CEO Jay Vijayan, and CFO Jacob Shulman. Tekion is an

AI-first dealer management system (DMS) built to run the full dealership tech stack on a single platform, spanning dealer management, CRM, digital retail, services, and payments. Tekion exited 2025 with ~\$250mn in ARR after >60% YoY growth, and expects to grow subscription revenue >50% YoY in 2026. Tekion's differentiation is its end-to-end platform (opposed to bolt-on tools) where AI can safely execute workflows due to its integration in the underlying system of record.

- **AI as labor arbitrage + throughput unlock.** AI impact on dealerships is centered on 1) increased sales efficiency through lower labor costs (dealerships are certain use cases, but that activity currently exists in areas that both on total volume on expansive for Guidewire like Underwriting (rather than in areas vehicles as well as vehicles sold per salesperson). Tekion is explicitly aiming at the most expensive constraint in the P&L - labor - highlighting that dealers spend >\$90B on labor and positioning AI agents (including an early technician-focused agent) as the mechanism to eliminate or redeploy headcount and grow with fewer people. If a salesperson has been stuck around ~10 cars/month for decades, Tekion believes agentic workflows can plausibly push that materially higher (they cite a world where it could be ~20 per person).
- **Measurable ROI outweighing price sensitivity, driving increased adoption of AI.** While dealerships are generally highly price sensitive, Tekion notes this customer base will adopt and pay for AI when the value is measurable such as increased profit per vehicle / repair, F&I, and ultimately hard ROI tied to headcount reduction or cheaper labor substitution. Tekion frames AI monetization as value-based, with AI SKUs in the \$15–25K ARR range, with hybrid pricing based on tiered usage or flat models, increasingly shifting toward outcome-based pricing (e.g., pay per appointment) because it aligns with how dealers think about spend.
- **Trust & compliance requirements as the gating moat for incumbents.** The auto retail end market is highly regulated and litigation-prone, which drives lower risk tolerance and discourages experimentation with generic LLM tools in production workflows. That said, Tekion positions this friction as an advantage as they own the integrated system and emphasize security/compliance. Through this, Tekion can offer a trusted AI interface that can securely execute tasks and workflows across the dealership stack. This trust posture is positioned as a key reason dealers, despite being price sensitive, are willing to adopt AI through Tekion when the alternative (uncontrolled LLM usage) creates increased compliance and operational risk.

Traversal

订阅研报添加微信 LCD123321Cad 04-16

We spoke with Anish Agarwal (CEO) and Ryan Powers (SVP Marketing) of Traversal.

Traversal is an AI-native site reliability engineering platform that autonomously investigates, diagnoses, and ultimately remediates complex production incidents by building a continuously updated model of enterprise systems and applying causal reasoning across logs, metrics, traces, and code.

- **Causality, not correlation, is the product.** Traversal's core claim is that most incident response today stalls because existing observability tooling surfaces correlations and symptoms, not causes. Traversal instead reconstructs the full causal

chain of an incident, traversing telemetry, dependency graphs, and recent system changes to isolate the true failure point. Management framed this as the difference between humans searching dashboards versus AI agents running thousands of targeted investigations in parallel, collapsing hours of manual triage into minutes. The result is a fundamentally different MTTR curve, not an incremental workflow improvement.

- **Built for agentic investigation at enterprise scale.** Traversal was explicit that its architecture is designed for AI agents, not humans. The platform re-indexes telemetry into a causal search layer and maintains a continuously updated “world model” of production environments, enabling autonomous agents to query, reason, and validate hypotheses at a scale that existing observability APIs cannot support. This is why management believes AI incident response cannot be bolted onto legacy tools without running into latency, cost, and query-limit ceilings.
- **From investigation layer to system of record over time.** While Traversal enters through incident response, the long-term ambition is broader: moving down-stack into the data layer and becoming a system of record for operational workflows. Management tied this directly to agentic change - as AI agents gain permission to write code, deploy infrastructure, and orchestrate workflows, failure modes become faster, larger, and harder to reason about. Traversal views autonomous diagnosis and, eventually, autonomous remediation as baseline requirements for operating modern production systems at scale, particularly in downtime-sensitive and regulated environments.

订阅研报添加微信 LCD123321Cad 04-16

Turing

We hosted Turing’s Founder & CEO, Jonathan Siddharth. Turing is a research accelerator for frontier AI labs and a proprietary-intelligence partner for Enterprises. Turing sits in the middle of the model supply chain, helping push frontier capabilities forward (especially in coding, reasoning, agents, and multimodality) while also translating those gains into production systems inside large companies. In practice, Turing supports a compounding loop - generate high-quality data/evals (often via reinforcement-learning-style environments), learn where models break in real enterprise workflows, then feed those failure modes back into the training and evaluation pipeline to improve the next iteration.

- **Frontier model compression will occur, though models will not become commoditized.** Although model performance continues to improve, the right-sized model matters more than raw scale once you move from demos to workflows. Management’s view is that you often don’t need a trillion-parameter generalist to automate a specific business process; you may need a much smaller, fine-tuned model that is cheaper, faster, and easier to operationalize - implying compression, specialization, and a more barbell-shaped stack of models in production. Mr. Siddharth’s argument against simple commoditization is that enterprise value accrues to models that win in narrow vertical contexts (finance, healthcare, etc.), so the ecosystem becomes plural rather than monoculture - different models, different

edges, different jobs.

- **AI will not eliminate the need for engineers, though the role must evolve.** As coding is increasingly handled by AI-technologies (a theme that was echoed by a number of participants at our conference), coding will move downstream. The role of the engineer will increasingly shift to defining the problem, decomposing it into phases, translating intent into sprints, then monitoring and iterating post-deploy. As models become better at coding, that doesn't delete the need for engineers so much as it drags the center of gravity upstream - from typing code to specifying systems, judging tradeoffs, and running tighter build-measure-learn loops. Per management, the uncomfortable corollary is transition risk: fewer employees in some functions, a relocation of labor, and a period where org charts shrink faster than new roles harden.
- **Operational data will become the real differentiator.** Across the four types of data Mr. Siddharth discussed (public, synthetic, hybrid-synthetic, and operational), the real differentiator becomes operational data (how real humans interact with models in real workflows), with rubrics that can verify quality. Public data is largely consumed at this stage, synthetic data helps but is better at teaching "student" models than pushing the "teacher" frontier. This benefits Turing as the company builds RL-like environments and enterprise solutions, watch where they fail, and convert those failures into differentiated training/eval data.

Rubrik

We hosted Arvind Nithrakashyap (CTO, Co-Founder) and Kiran Choudary (CFO) of Rubrik.

- **Agent Cloud represents a potentially important new control layer for agentic AI.** Management framed the core problem as governance and dynamic authorization, not just authentication: once an agent is authenticated into systems like Salesforce, Jira, GitHub, or Outlook, the harder question is whether a given action should actually be allowed in context. Rubrik's pitch is that this cannot be solved by static identity rules alone, since agents can behave unpredictably inside approved systems; instead, customers need policy-aware monitoring, model-level inspection of prompts and tool calls, and the ability to remediate or "rewind" actions when something goes wrong. That is where Rubrik believes small language models can act as judges against natural-language policy, which in our view is a differentiated and more operationally grounded framing of agent security than many of the early point-solution pitches in market.
- **Rubrik's competitive positioning rests on combining data, metadata, and identity context in a way peers may struggle to replicate.** Management argued the architectural advantage is not simply knowing who a user or agent is, but understanding what data exists, where it lives, how it is connected, what permissions have changed over time, and which actions are being taken against it across cloud, SaaS, and on-prem environments. That is the basis for Rubrik's argument that identity-centric vendors can define broad access boundaries, but are less

well-positioned to make the fine-grained, dynamic decisions required once an agent is operating inside an application. The same logic also underpins Rubrik's view versus legacy data protection vendors like Commvault, where Management argued older architectures are not as well-positioned for modern, metadata-rich, application-level resilience. Separately, on the hardware side, Management suggested memory pricing has had some impact at the margin (in context of F4Q results), but not in a way that is materially changing demand, given the business mix has shifted meaningfully toward pure software / cloud offerings and customers still have virtual and alternate deployment options for hybrid workloads.

- **Rubrik highlighted both expanding protection needs across newer workloads and a measured approach to scaling newer products like Identity and Agent Cloud.** Management suggested AI is increasing the number of workloads customers now view as business-critical, including analytical data, large S3 estates, and modern data lake environments that historically may not have sat at the top of the protection stack. They also pointed to more customer discussion around sovereignty, geographic redundancy, and vaulting in the context of geopolitical risk, including questions around whether recovery environments can sit in regions better insulated from physical disruption. Separately, Management described a more specialized incubation motion for newer products like Identity and Agent Cloud, where a smaller team helps refine the buyer, messaging, and product motion before broader scale-out through the core field organization. That approach seems sensible given those products extend Rubrik beyond traditional IT / infrastructure conversations and further into the CISO organization and, in Agent Cloud's case, into AI governance workflows.

Writer

We hosted Matan-Paul Shetrit, Head of Product. Writer offers a fully integrated solution to deploy secure and reliable AI applications and agents that automate mission critical work. These development tools are supported by Palmyra, Writer's family of LLMs, alongside graph-based RAG and customizable AI guardrails.

- **Competitive differentiation vs. model providers and incumbents.** On competition, Writer positioned itself between model providers and incumbent application vendors; where frontier model companies remain API-centric and are still building enterprise distribution, governance, and deployment capabilities, while incumbents benefit from entrenched data and workflows but are constrained by legacy architectures and permissive automation layers. Writer's differentiation sits in controlling the orchestration layer end-to-end (coordinating agents, models, and enterprise systems) without forcing customers to rebuild their stack or surrender control of sensitive workflows. For example, in consumer packaged goods (CPG) GTM use cases, management highlighted material reductions in launch timelines; from roughly four months historically, to two months, to now largely automated. This is despite fragmentation across retailers, formats, and geographies. The key takeaway is not just speed but repeatability: once a workflow is automated and

productized, it can be reused and deployed broadly with less incremental effort.

- **Model strategy focused on cost-efficient, domain-specific systems.** Writer laid out a framework centered on smaller, specialized models that improve through usage within a customer's isolated environment, using 20–30B parameter models that are cheap, fast, and trained continuously on customer-specific usage within an isolated environment. Management noted that after ~5 runs, these domain-specific models can match leading LLM performance for targeted tasks, while avoiding high costs associated with frontier models. This was reinforced by Palmyra Fin being the first model to pass the CFA exam. The implication is lower inference costs, better domain performance over time, and a more sustainable enterprise economics profile compared with reliance on large, expensive general-purpose models.
- **Internal product development.** Writer is rethinking product development in response to AI-driven software creation; management highlighted significant experimentation around product specifications and development process, noting that fully removing specs led to poor outcomes and engineering friction, while overly rigid documentation slowed iteration. The key takeaway was a shift toward lighter-weight specs and a renewed emphasis on engineers as system designers rather than code producers, reinforcing the view that AI accelerates execution but does not eliminate the need for structured product discipline.

Workday

We hosted Workday's President of Product & Technology, Gerrit Kazmaier.

Workday described its API strategy as primarily designed to align usage with cost.

Historically, application-level APIs were broadly free and used for general integrations, but Workday has seen a sharp uptick yoy concentrated in specific customer cohorts, driving materially higher infrastructure spend. Every customer still receives a default API entitlement tied to size and historical usage, but usage beyond design intent now drives overage charges. For example, some customers effectively use Workday as a document database in addition to its core functionality, and it didn't make sense for this to be free. Management emphasized the discussion around monetization is largely noise rather than signal: for most of the base, incremental cost is a non-event, simply formalizing a rational value discussion. The shift mirrors practices at other scaled software companies and reflects a focus on capturing value where usage is non-uniform rather than raising prices across the board.

More structurally, Workday is separating traditional application APIs from a premium class of AI-oriented APIs, sometimes likened to MCP servers, but exposed selectively through Extend Pro at the highest platform tier. These premium APIs are designed for agentic workflows, accrue purely on consumption, and are intentionally limited by use case rather than opened wholesale. Monetization is abstracted through a credit system decoupled from underlying token economics: customers buy flex credits mapped internally to "units of work," while Workday manages token efficiency behind the scenes. Credit consumption varies by agent and task complexity, with simpler self-service actions costing less. Management was explicit that pricing is for value, not

cost, even as internally they optimize aggressively for token efficiency, similar to how they historically optimized storage and compute COGS.

On growth, management acknowledged that it will take time to drive acceleration at its current run rate even with its initial success with organic and inorganic AI-first revenue. That said, momentum is building—Workday’s March agent release already has more than 2,000 customers live, and AI-driven win conversations now dominate new deals around recruiting, contracts, and financial automation. Workday expects to have new product releases in May. Payroll remains <60% penetrated, but management views net-new customer growth—particularly outside North America in Asia and Europe—as more attractive than retrofit conversions. Longer term, Workday sees first-party agents, premium APIs, and outcome-based monetization replacing seats in many workflows, with value defined at the enterprise level (e.g., aggregate HR case volume at a Walmart-scale customer) rather than by individual power users.

Zscaler

We hosted Zscaler’s CFO, Kevin Rubin and EVP of Agentic AI, Swamy Kocherlakota.

- **Zscaler noted ongoing engagement with LLM providers, including for latest gen models.** It sees a distinction between the types of vulnerabilities typically found in hardware (such as firewalls, that have complex on premise code) and end user devices (such as endpoint), such that these types of vulnerabilities can be more easily hacked using AI. In contrast, Zscaler’s code installed in customer environments is a more simple packet forwarder, which has minimal attack surface to hack. Zscaler will benefit nonetheless from being able to harden its internal infrastructure.
- **On demand, Zscaler is seeing momentum building for its AI products, including AI Security Suite launched in January, and expects material contribution next year.** These products sometimes serve as the first land for Zscaler, independent of its core ZIA and ZPA offering. For data security, Zscaler reiterated that these products are growing faster than the core. There is also the potential next wave of SASE deployments if customers see the risk of LLM exploited vulnerabilities as a catalyst to move to cloud-hosted SASE. For agents, Zscaler noted the advantages of a zero trust architecture with in line enforcement, and its future plans to release in line proxy for agents. Zscaler is seeing an uptick of agentic traffic on its network and has a component of usage based pricing in its existing solutions (e.g. guest wifi which is not priced by number of users); over time, Zscaler has an opportunity to monetize a step change in traffic along with any changes in capacity.